

OpenAI Agents SDK

Abdul Rehman (PIAIC & GIAIC)
MCQS

OpenAI Agents SDK - Introductory Section

1. What kind of AI apps does the OpenAI Agents SDK enable you to build?

- ☐ a) Data analysis apps
 - ☒ b) **Agentic AI apps**
 - ☐ c) Static website generators
 - ☐ d) Mobile gaming apps
-

2. The Agents SDK is described as a production-ready upgrade of what previous experimentation for agents?

- ☐ a) Hive
 - ☐ b) Colony
 - ☒ c) **Swarm**
 - ☐ d) Nexus
-

3. Which of the following is NOT listed as a primitive of the Agents SDK?

- ☐ a) Agents
- ☐ b) Handoffs
- ☒ c) **Guardrails** (*Trick question! Guardrails **are** a primitive*)
- ☐ d) Workflows

(Note: All are actual primitives, so this is a questionable phrasing — likely intended to test awareness.)

4. What are "Agents" defined as in the SDK?

- ☐ a) Simple Python functions
 - ☒ b) **LLMs equipped with instructions and tools**
 - ☐ c) Data validation modules
 - ☐ d) User interface components
-

5. What do "Handoffs" allow agents to do?

- ☐ a) Log their activities
- ☐ b) Self-correct errors

- ☒ c) **Delegate to other agents for specific tasks**
 - ☐ d) Interact with external APIs
-

6. "Guardrails" enable what functionality in relation to agent inputs?

- ☐ a) Sending results to the LLM
 - ☐ b) Calling external tools
 - ☒ c) **Validating the inputs to agents**
 - ☐ d) Orchestrating multiple agents
-

7. In combination with what programming language are the SDK's primitives powerful enough to express complex relationships?

- ☐ a) Java
 - ☐ b) JavaScript
 - ☒ c) **Python**
 - ☐ d) C++
-

8. What is a key benefit of the SDK having "very few abstractions"?

- ☐ a) It makes the SDK more complex.
 - ☒ b) **It makes it quick to learn.**
 - ☐ c) It limits customization options.
 - ☐ d) It requires more boilerplate code.
-

9. Which design principle states that the SDK "Works great out of the box, but you can customize exactly what happens"?

- ☐ a) Python-first
 - ☐ b) Few primitives
 - ☐ c) Agent loop
 - ☒ d) **Works great out of the box, but you can customize exactly what happens**
-

10. What does the "Agent loop" feature primarily handle?

- ☐ a) Defining agent instructions
 - ☐ b) Generating automatic responses
 - ☒ c) **Calling tools, sending results to the LLM, and looping until the LLM is done**
 - ☐ d) Visualizing agent traces
-

11. The "Python-first" approach means using what for orchestration and chaining agents?

- ☐ a) A new, proprietary domain-specific language
- ☐ b) A visual programming interface

- ✓c) **Built-in language features**
 - ✗d) External YAML configuration files
-

12. What specific capability do "Handoffs" provide for multiple agents?

- ✗a) Error handling
 - ✗b) Parallel execution of single tasks
 - ✓c) **Coordination and delegation**
 - ✗d) Automatic fine-tuning
-

13. How do "Guardrails" typically behave if checks fail?

- ✗a) They log a warning and continue
 - ✗b) They prompt the user for a new input
 - ✓c) **They break early.**
 - ✗d) They attempt to fix the input automatically
-

14. What do "Function tools" allow you to turn any Python function into?

- ✗a) A data validation model
 - ✗b) A new agent
 - ✓c) **A tool**
 - ✗d) A guardrail
-

15. "Function tools" come with automatic schema generation and validation powered by what library?

- ✗a) NumPy
 - ✗b) Pandas
 - ✓c) **Pydantic**
 - ✗d) SciPy
-

16. What is the benefit of the built-in "Tracing" feature?

- ✗a) It automates code deployment
 - ✗b) It encrypts sensitive data
 - ✓c) **It lets you visualize, debug and monitor workflows, and use OpenAI tools for evaluation, fine-tuning, and distillation.**
 - ✗d) It generates user documentation automatically
-

17. What is the command to install the OpenAI Agents SDK?

- ✗a) npm install openai-agents
- ✓b) **pip install openai-agents**
- ✗c) conda install openai-agents
- ✗d) go get openai-agents

18. What is the name given to the Agent in the "Hello world" example?

- ☐ a) "Haiku Generator"
 - ☒ b) "Assistant"
 - ☐ c) "Recursion Expert"
 - ☐ d) "Programming Helper"
-

19. What instruction is given to the Agent in the "Hello world" example?

- ☐ a) "Write a poem about nature."
 - ☐ b) "Solve a math problem."
 - ☒ c) "You are a helpful assistant"
 - ☐ d) "Explain AI concepts."
-

20. What specific task is the Agent asked to perform in the "Hello world" example?

- ☐ a) Generate a story about a robot
 - ☐ b) Explain quantum physics
 - ☒ c) Write a haiku about recursion in programming.
 - ☐ d) Translate a sentence into French
-

21. What is the expected final output of the "Hello world" example (the haiku)?

- ☐ a) "AI, a new dawn, / Machines learn, systems grow, / Future now unfolds."
 - ☒ b) "Code within the code, / Functions calling themselves, / Infinite loop's dance."
 - ☐ c) "Program's gentle flow, / Logic weaves, commands obey, / Digital new form."
 - ☐ d) "Recursion deep, / Calls itself, repeats the task, / Stack grows, then unwinds."
-

22. What environment variable needs to be set to run the "Hello world" example successfully?

- ☐ a) API_KEY
 - ☐ b) OPENAI_API_ENDPOINT
 - ☒ c) OPENAI_API_KEY
 - ☐ d) AGENT_SDK_KEY
-

23. The SDK is described as enabling you to build agentic AI apps in a package that is:

- ☐ a) Complex and heavy
 - ☒ b) Lightweight and easy-to-use
 - ☐ c) Experimental and feature-limited
 - ☐ d) Dependent on many external abstractions
-

24. What previous project is the Agents SDK considered a "production-ready upgrade" of?

- ☐ a) Core Agents
 - ☐ b) AgentLab
 - ☒ c) **Swarm**
 - ☐ d) AI Forge
-

25. The primitives of the Agents SDK (Agents, Handoffs, Guardrails) in combination with Python are powerful enough to express what?

- ☐ a) Only simple linear processes
- ☐ b) Basic data transformations
- ☒ c) **Complex relationships between tools and agents**
- ☐ d) Standalone data visualization

Quickstart

OpenAI Agents SDK - Quickstart Specifics

MCQs

1. What is the recommended first step for setting up a new project to use the Agents SDK?

- ☐ a) Install the SDK directly.
 - ☐ b) Set the OpenAI API key.
 - ☒ c) **Create a project directory and a Python virtual environment.**
 - ☐ d) Define your first agent immediately.
-

2. After creating the project directory and virtual environment, what is the next action advised before installing the SDK?

- ☐ a) Create an OpenAI API key.
 - ☒ b) **Activate the virtual environment.**
 - ☐ c) Define the main function.
 - ☐ d) Run a test command.
-

3. The Agent constructor allows for optional configuration. Which specific optional config parameter is mentioned when creating your first agent?

- ☐ a) api_version
- ☐ b) debug_mode
- ☒ c) **model_config**
- ☐ d) trace_level

4. When adding "a few more agents" like History Tutor and Math Tutor, what specific parameter is added to provide context for handoff routing?

- ☐ a) routing_instructions
 - ☐ b) context_description
 - ☒ c) **handoff_description**
 - ☐ d) specialty_area
-

5. What is the explicit instruction given to the Triage Agent regarding its role?

- ☐ a) "You will answer all user questions directly."
 - ☐ b) "You will provide help with math or history problems."
 - ☒ c) **"You determine which agent to use based on the user's homework question."**
 - ☐ d) "You validate user inputs for safety."
-

6. How does the Triage Agent know which specialist agents (e.g., History Tutor, Math Tutor) it can delegate to?

- ☐ a) It infers them from their instructions.
 - ☒ b) **They are provided in a list to the handoffs parameter of the Triage Agent.**
 - ☐ c) They are automatically discovered by the SDK.
 - ☐ d) It must manually import them during runtime.
-

7. In the "Run the agent orchestration" example, what specific method of the `Runner` class is used to initiate the agent workflow asynchronously?

- ☐ a) `Runner.run_sync()`
 - ☐ b) `Runner.execute()`
 - ☒ c) **`Runner.run()`**
 - ☐ d) `Runner.start()`
-

8. The initial test query for the agent orchestration is "What is the capital of France?". Which agent is this query first sent to?

- ☐ a) `history_tutor_agent`
 - ☐ b) `math_tutor_agent`
 - ☒ c) **`triage_agent`**
 - ☐ d) `guardrail_agent`
-

9. When defining a custom guardrail, what is the purpose of the `HomeworkOutput` class inheriting from `pydantic.BaseModel`?

- ☐ a) To store the guardrail's internal state.
- ☐ b) To define the agent's instructions.
- ☒ c) **To provide a structured output type for the `guardrail_agent`'s check result.**

- ☒ d) To activate the virtual environment.
-

10. The `guardrail_agent` is specifically instructed to check for what?

- ☒ a) Malicious content in the input.
 - ☒ b) Grammatical errors in the user's question.
 - ☒ c) **If the user is asking about homework.**
 - ☒ d) The length of the user's input.
-

11. What is the role of `ctx.context` when running the `guardrail_agent` inside the `homework_guardrail` function?

- ☒ a) It's unused in this context.
 - ☒ b) **It passes the run's context object down to the guardrail's agent sub-run.**
 - ☒ c) It specifies the model configuration for the guardrail.
 - ☒ d) It stores the guardrail's output.
-

12. The `homework_guardrail` function returns `GuardrailFunctionOutput`. What value determines if the "tripwire" is triggered?

- ☒ a) `final_output.is_homework`
 - ☒ b) `final_output.reasoning`
 - ☒ c) **not `final_output.is_homework`**
 - ☒ d) `result.final_output`
-

13. How is the `homework_guardrail` attached to the `triage_agent` in the full workflow example?

- ☒ a) As an `output_guardrail`.
 - ☒ b) Within the `handoffs` list.
 - ☒ c) **As an `InputGuardrail` in the `input_guardrails` list.**
 - ☒ d) As a `function_tool`.
-

14. What standard Python module is used to run the `async def main()` function in the full "Put it all together" example?

- ☒ a) `threading`
 - ☒ b) `concurrent`
 - ☒ c) **`asyncio`**
 - ☒ d) `multiprocessing`
-

15. In the combined example, when the `triage_agent` is given "what is life" as an input, which part of the system is expected to trigger?

- ☒ a) The History Tutor agent.
- ☒ b) The Math Tutor agent.

- ✓c) **The `homework_guardrail`.**
 - ✗d) The handoff mechanism.
-

16. To review what happened during an agent run, where are users instructed to navigate?

- ✗a) The project's local log files.
 - ✗b) The `Runner` object's internal state.
 - ✓c) **The Trace viewer in the OpenAI Dashboard.**
 - ✗d) The Python interpreter's history.
-

17. The "Next steps" suggest learning how to build "more complex agentic flows." This implies the quickstart example primarily showcases what?

- ✗a) Only the most advanced features.
 - ✗b) Only theoretical concepts.
 - ✓c) **Foundational elements and a basic, yet functional, agent orchestration.**
 - ✗d) Production-ready deployment strategies.
-

18. What is the explicit purpose of `handoff_descriptions` for the History Tutor and Math Tutor agents?

- ✗a) To provide alternative names.
 - ✓b) **To help the Triage Agent determine routing based on context.**
 - ✗c) To summarize their instructions.
 - ✗d) To indicate if they are ready for production.
-

19. The SDK is described as a "lightweight, easy-to-use package with very few abstractions." How does the Quickstart example demonstrate this?

- ✗a) By showing complex inheritance hierarchies.
 - ✓b) **By defining agents and their interactions with clear, concise Python code.**
 - ✗c) By requiring extensive configuration files.
 - ✗d) By relying on a graphical user interface.
-

20. The `is_homework` field in `HomeworkOutput` is a boolean. What is the type of the reasoning field?

- ✗a) `list[str]`
 - ✗b) `int`
 - ✓c) **`str`**
 - ✗d) `bool`
-

Examples

OpenAI Agents SDK - Examples Section MCQs

1. The examples in the SDK repository are primarily intended to demonstrate what?

- ☐ a) Optimal performance benchmarks.
 - ☐ b) The SDK's internal architecture.
 - ☒ c) **Different agent design patterns and capabilities.**
 - ☐ d) Compatibility with all Python versions.
-

2. If a developer wants to understand how to ensure an agent workflow follows a predictable sequence of steps, which example category would be most helpful?

- ☐ a) basic
 - ☒ b) **agent_patterns**
 - ☐ c) handoffs
 - ☐ d) tool examples
-

3. The `agent_patterns` category mentions "Agents as tools." What does this imply about the design flexibility of the SDK?

- ☐ a) Agents can only be used as primary responders.
 - ☒ b) **Agents themselves can be modular components utilized by other agents or systems.**
 - ☐ c) Tools must always invoke an agent.
 - ☐ d) The SDK only supports simple agent-tool interactions.
-

4. Which category would be essential for a developer trying to implement real-time interaction feedback or partial responses from the LLM?

- ☐ a) handoffs
 - ☐ b) tool examples
 - ☒ c) **basic (due to "Streaming outputs")**
 - ☐ d) model providers
-

5. What distinguishes the "Dynamic system prompts" examples from typical agent instructions?

- ☐ a) They are always very short.
 - ☐ b) They are defined as separate files.
 - ☒ c) **They allow the agent's guiding instructions to change based on runtime conditions or context.**
 - ☐ d) They are specific to non-OpenAI models.
-

6. For a user looking to integrate a feature like web search into their agent, which example category would guide them on how to do this using OAI-hosted functionalities?

- ☐ a) agent_patterns
 - ☒ b) **tool examples**
 - ☐ c) basic
 - ☐ d) research_bot
-

7. The `model_providers` category is crucial for developers who are concerned with what aspect of their agent application?

- ☐ a) Only using OpenAI's latest models.
 - ☒ b) **Utilizing LLMs from different vendors or custom sources.**
 - ☐ c) Optimizing model inference speed.
 - ☐ d) Ensuring model interpretability.
-

8. The `customer_service` example illustrates a system for an airline. What is the broader implication of such "built-out examples"?

- ☐ a) They are purely theoretical demonstrations.
 - ☐ b) They only show basic agent setup.
 - ☒ c) **They showcase how to construct more complete, real-world, industry-specific applications.**
 - ☐ d) They primarily focus on debugging agent traces.
-

9. If a user wants to build an agent that can perform extensive information gathering similar to a specialized research assistant, which example would serve as a direct reference?

- ☐ a) customer_service
 - ☐ b) handoffs
 - ☒ c) **research_bot**
 - ☐ d) agent_patterns
-

10. What specific type of AI interaction is highlighted by the `voice` examples category?

- ☐ a) Visual recognition.
 - ☐ b) Text summarization.
 - ☒ c) **Conversational agents that interact via spoken language.**
 - ☐ d) Code generation.
-

11. The mention of TTS and STT models in the voice examples indicates the SDK's capability to bridge agents with what kind of interfaces?

- ☐ a) Database interfaces.
- ☐ b) Command-line interfaces.
- ☒ c) **Audio-based interfaces.**
- ☐ d) Graphical user interfaces.

12. What does the inclusion of `Parallel agent execution in agent_patterns` suggest about the SDK's runtime capabilities?

- ☒ a) Agents must always run sequentially.
 - ☒ b) It supports basic multi-threading only.
 - ☒ c) **It can orchestrate multiple agents to work concurrently on tasks.**
 - ☒ d) It is limited to single-agent operations.
-

13. The organization of examples into categories like `handoffs`, `agent_patterns`, and `basic` is designed to help users primarily with what?

- ☒ a) Identifying bugs in the SDK.
 - ☒ b) **Learning specific architectural approaches and foundational features.**
 - ☒ c) Benchmarking different LLM providers.
 - ☒ d) Understanding the SDK's installation process.
-

14. What could a user learn from the `Lifecycle events` examples within the `basic` category?

- ☒ a) How to define an agent's name.
 - ☒ b) How to force an agent to use a tool.
 - ☒ c) **How to observe and react to different stages of an agent's operation.**
 - ☒ d) How to set up a virtual environment.
-

15. The phrase "Simple deep research clone" for `research_bot` implies it's designed to mimic what kind of functionality?

- ☒ a) A search engine's indexing.
- ☒ b) A human's ability to thoroughly investigate a topic.
- ☒ c) **An automated system capable of in-depth information gathering and synthesis.**
- ☒ d) A tool for quick factual lookups.

Documentation

Agents

OpenAI Agents SDK - Agents Section MCQs (Nuanced)

1. Beyond temperature and top_p, what broader category of parameters can be configured via model_settings for an LLM?

- ☐ a) Only specific API keys.
 - ☒ b) **Model tuning parameters.**
 - ☐ c) Networking configurations.
 - ☐ d) Agent name and instructions.
-

2. The get_weather tool is decorated with @function_tool. What does this decorator implicitly handle for the Python function?

- ☐ a) Automatic caching of results.
 - ☒ b) **Automatic schema generation for the LLM to understand and use the tool.**
 - ☐ c) Dynamic scaling of the function.
 - ☐ d) Error handling within the function's logic.
-

3. The Context mechanism is described as "dependency-injection." What does this term imply?

- ☐ a) They must explicitly import all external modules.
- ☒ b) **Dependencies are provided to them externally, rather than being created internally.**

- ☒ c) They can only access global variables.
 - ☒ d) They can only inject other agents as dependencies.
-

4. In the `UserContext` example, what is `async def fetch_purchases()` designed to represent?

- ☒ a) A static utility function.
 - ☒ b) **Asynchronous operations that can fetch data relevant to the user during an agent run.**
 - ☒ c) A method for storing user credentials.
 - ☒ d) A function for generating new users.
-

5. What does Pydantic's `TypeAdapter` flexibility suggest about `output_type` in tools/agents?

- ☒ a) Only `BaseModel` subclasses are supported.
 - ☒ b) Only simple Python primitives are allowed.
 - ☒ c) **It allows for a broad range of Python types (dataclasses, lists, `TypedDict`, etc.).**
 - ☒ d) It implies a performance overhead.
-

6. What is the main benefit of combining instruction: `Extract calendar events from text with output_type=CalendarEvent`?

- ☒ a) It prevents the agent from making any mistakes.
 - ☒ b) It allows the agent to generate plain text and structured data simultaneously.
 - ☒ c) **It guides the LLM to format identified events into structured, machine-readable format.**
 - ☒ d) It enables the agent to directly add events to a calendar.
-

7. What problem do modular handoffs between agents aim to solve?

- ☒ a) Overfitting of a single large LLM.
 - ☒ b) **The challenge of one LLM needing to handle too many diverse tasks.**
 - ☒ c) Network latency between agents.
 - ☒ d) The cost of running multiple small LLMs.
-

8. Why is context passed to the `dynamic_instructions` function?

- ☒ a) To allow the agent to modify its own name.
 - ☒ b) **To enable instructions to be personalized or conditional based on runtime data.**
 - ☒ c) To force the agent to use a specific tool.
 - ☒ d) To manage the agent's internal memory.
-

9. Why would a developer use "lifecycle events" (hooks)?

- ☒ a) To dynamically change the agent's model.
- ☒ b) **To perform actions like logging, monitoring, or data pre-fetching at specific stages.**
- ☒ c) To bypass guardrail checks.

- ☒ d) To automatically regenerate instructions.
-

10. What is the advantage of guardrails running "in parallel to the agent"?

- ☒ a) It ensures the agent runs faster.
 - ☒ b) It allows for real-time model fine-tuning.
 - ☒ c) **It enables early breaking if input validation fails, saving resources.**
 - ☒ d) It simplifies the agent's instruction set.
-

11. What is the use of the `clone()` method on agents?

- ☒ a) Running agents synchronously.
 - ☒ b) **Creating specialized versions of a base agent with minor changes.**
 - ☒ c) Forcing tool use behavior.
 - ☒ d) Defining dynamic instructions.
-

12. What does the `tool_choice` setting mean for the LLM's control?

- ☒ a) The content of the tool's output.
 - ☒ b) **The selection of the most appropriate tool from the available list.**
 - ☒ c) Whether to run indefinitely.
 - ☒ d) Its own internal reasoning process.
-

13. Why does `tool_choice` reset to "auto" after each tool call?

- ☒ a) Slow API responses.
 - ☒ b) **To avoid infinite loops where the LLM keeps calling tools unnecessarily.**
 - ☒ c) Incorrect tool output formats.
 - ☒ d) Overuse of `ModelSettings`.
-

14. What happens when `tool_use_behavior="stop_on_first_tool"` is set?

- ☒ a) It continues to process with `tool_choice="auto"`.
 - ☒ b) **It stops LLM processing after using the first tool and uses that result as final output.**
 - ☒ c) It re-runs the previous turn.
 - ☒ d) It triggers an input guardrail.
-

15. What does the trio of `instructions`, `model_settings`, and `tools` provide the LLM?

- ☒ a) Only raw text input.
- ☒ b) A fixed, unchangeable persona.
- ☒ c) **A defined persona, configuration, and tool access for rich interaction.**
- ☒ d) Solely conversational abilities.

Running Agents

OpenAI Agents SDK - Running Agents Section MCQs

✓Part 1: Running Methods & The Agent Loop

✓Part 1: Running Methods & The Agent Loop

1. Which Runner method is asynchronous and returns a RunResult?

- ✗a) Runner.run_sync()
- ✗b) Runner.run_streamed()
- ✓c) **Runner.run()**
- ✗d) Runner.execute()

2. Runner.run_sync() internally uses:

- ✗a) Runner.stream_events()
- ✓b) **Runner.run()**
- ✗c) Runner.execute_sync()
- ✗d) Runner.async_run()

3. What does Runner.run_streamed() return?

- ✗a) RunResult
- ✓b) **RunResultStreaming**
- ✗c) RunStream
- ✗d) StreamResult

4. Initial input types to Runner methods?

- ✗a) An integer or a boolean
- ✗b) A string or a dictionary
- ✓c) **A string (user message) or a list of input items**
- ✗d) A tuple or a set

5. If the LLM produces final_output:

- ✗a) The LLM does a handoff
- ✗b) The LLM produces tool calls
- ✓c) **The loop ends and the result is returned**
- ✗d) A MaxTurnsExceeded exception is raised

6. **If the LLM does a handoff:**
☒ a) The loop ends immediately
☒ b) **The current agent and input are updated, and the loop re-runs**
☒ c) Tool calls are executed
☒ d) The process waits for user intervention
7. **If the LLM produces tool calls:**
☒ a) The loop ends with an error
☒ b) The final_output is immediately returned
☒ c) **Those tool calls are run, their results are appended, and the loop re-runs**
☒ d) A handoff is automatically triggered
8. **If max_turns is exceeded:**
☒ a) The agent simply stops without an error
☒ b) **A MaxTurnsExceeded exception is raised**
☒ c) The agent automatically restarts
☒ d) The final_output is forced
9. **For output to be "final":**
☒ a) It must have done at least one handoff
☒ b) **There must be no tool calls**
☒ c) It must exceed max_turns
☒ d) It must be streamed
-

✓Part 2: Streaming & Run Config

10. **What does streaming provide?**
☒ a) Debug logs only.
☒ b) **Streaming events**
☒ c) Performance metrics.
☒ d) Model weights.
11. **Purpose of run_config?**
☒ a) To configure the Agent class directly.
☒ b) To define agent instructions.
☒ c) **To configure some global settings for the entire agent run.**
☒ d) To manage the virtual environment.
12. **Override model via run_config?**
☒ a) default_model
☒ b) global_llm
☒ c) **model**
☒ d) override_model

13. **Default model_provider?**

- ☐ a) Google
- ☐ b) Anthropic
- ☒ c) **OpenAI**
- ☐ d) Custom

14. **run_config.model_settings can be used to:**

- ☒ a) **Override agent-specific settings like temperature or top_p globally.**
- ☐ b) Define new tools for the agent.
- ☐ c) Specify handoff behavior.
- ☐ d) Disable tracing.

15. **run_config allows what guardrails?**

- ☐ a) Only output guardrails.
- ☐ b) Only input guardrails.
- ☒ c) **Both input and output guardrails.**
- ☐ d) Only global guardrails.

16. **handoff_input_filter allows you to:**

- ☐ a) Disable specific handoffs.
- ☒ b) **Edit the inputs that are sent to the new agent.**
- ☐ c) Prioritize certain handoffs.
- ☐ d) Redirect handoffs to external APIs.

17. **Disable tracing via:**

- ☐ a) disable_trace
- ☐ b) tracing_off
- ☒ c) **tracing_disabled**
- ☐ d) no_tracing

18. **trace_include_sensitive_data enables:**

- ☐ a) Whether traces are enabled or disabled.
- ☒ b) **Whether traces will include LLM and tool call inputs/outputs.**
- ☐ c) The location where trace data is stored.
- ☐ d) The format of the trace metadata.

19. **Recommended tracing parameter:**

- ☐ a) trace_id
- ☐ b) group_id
- ☒ c) **workflow_name**
- ☐ d) trace_metadata

✓Part 3: Conversations & Exceptions

20. A **Runner call** = what in chat?

- ✗a) A complete conversation.
- ✗b) A single independent query.
- ✓c) A **single logical turn**
- ✗d) A debugging session.

21. Get next turn input via:

- ✗a) `to_output_list()`
- ✗b) `get_next_input()`
- ✓c) `to_input_list()`
- ✗d) `prepare_for_next_turn()`

22. What links traces across multiple runs?

- ✗a) `workflow_name`
- ✗b) `trace_id`
- ✓c) **`group_id` (via `thread_id`)**
- ✗d) `trace_metadata`

23. What is **AgentsException**?

- ✗a) An exception specific to tool calls.
- ✓b) **The base class for all exceptions raised in the SDK.**
- ✗c) An exception related to model behavior.
- ✗d) An exception for user-made errors.

24. **MaxTurnsExceeded** is raised when:

- ✗a) When a guardrail is tripped.
- ✗b) When the model produces invalid outputs.
- ✓c) **When the run exceeds the `max_turns` passed to the run methods.**
- ✗d) When the user makes an error in the SDK.

25. **ModelBehaviorError** is raised for:

- ✗a) Exceeding `max_turns`.
- ✓b) **When the model produces invalid outputs (e.g., malformed JSON)**
- ✗c) User input validation failure.
- ✗d) Disabling tracing.

26. **UserError** is raised when:

- ✗a) Internal SDK bugs.
- ✗b) When the LLM generates an incorrect answer.

- ✓c) When the person writing code using the SDK makes an error.
- ✗d) Network connectivity issues.

27. **Guardrail exceptions:**

- ✗a) GuardrailError and TripwireError
- ✗b) InputGuardrailFailed and OutputGuardrailFailed
- ✓c) **InputGuardrailTripwireTriggered and OutputGuardrailTripwireTriggered**
- ✗d) ValidationError and CheckFailed

28. **Second turn input is created by:**

- ✗a) Only by the user's new question.
- ✗b) By the previous result.final_output.
- ✓c) **By combining result.to_input_list() with the new user message.**
- ✗d) By resetting the agent state.

29. **Initial Assistant instruction in example:**

- ✗a) "Write a haiku about recursion."
- ✓b) **"Reply very concisely."**
- ✗c) "What city is the Golden Gate Bridge in?"
- ✗d) "Explain your reasoning at each step."

30. **A Runner run represents what in a chat?**

- ✗a) A complete conversation history.
- ✗b) A single independent query.
- ✓c) **A single logical turn.**
- ✗d) A new conversation thread.

Results

OpenAI Agents SDK - Results Section MCQs

✓Part: Results & Run Items

1. **What type of object is returned when you call `Runner.run()` or `Runner.run_sync()`?**

- a) RunResultStreaming
- ✓b) **RunResult**
- c) RunResultBase
- d) RunnerResult

2. **If you call `Runner.run_streamed()`, what type of object do you receive?**

- ✓a) **RunResultStreaming**
- b) RunResult

- c) RunResultBase
 - d) StreamingResult
-

3. **Both RunResult and RunResultStreaming inherit from which base class where most useful information is present?**
- a) RunBase
 - ✓b) **RunResultBase**
 - c) ResultBase
 - d) RunnerBase
-

4. **What does the final_output property contain?**
- a) All intermediate steps of the agent run.
 - ✓b) **The final output of the last agent that ran.**
 - c) The original input provided to the runner.
 - d) A list of all tool calls made.
-

5. **If the last agent that ran did not have an output_type defined, what will be the type of final_output?**
- a) An object of last_agent.output_type
 - b) A list
 - ✓c) **A str**
 - d) None
-

6. **Why is the final_output property of type Any in the SDK?**
- a) To allow for easier debugging.
 - b) Because its type is statically known only when no handoffs occur.
 - ✓c) **Due to handoffs, as any Agent might be the last agent, so the output type is not statically known.**
 - d) It is always a generic object, regardless of output type.
-

7. **Which method allows you to turn the result into an input list suitable for concatenating with original input for subsequent runs?**
- a) result.get_input_list()
 - b) result.get_next_input()
 - ✓c) **result.to_input_list()**
 - d) result.input_history()
-

8. **The to_input_list() method is convenient for what purpose?**
- a) To reset the agent's state.
 - b) To explicitly define new handoffs.
 - ✓c) **To take outputs of one agent run and pass them into another run or loop with new user inputs.**
 - d) To disable tracing for the next run.
-

9. **What does the `last_agent` property contain?**

- a) The initial agent that started the run.
 - ✓b) **The last agent that ran.**
 - c) A list of all agents involved in the run.
 - d) The Triage Agent by default.
-

10. **When might storing and re-using the `last_agent` be useful?**

- a) When you want to force a specific tool use.
 - ✓b) **If you have a frontline triage agent that hands off to a specialized agent, to re-use the specialized agent for follow-up.**
 - c) To disable guardrails for the next run.
 - d) To always reset the conversation.
-

11. **What does the `new_items` property contain?**

- a) Only the final response.
 - ✓b) **The new items (`RunItems`) generated during the run.**
 - c) A list of all model responses.
 - d) The original input and final output combined.
-

12. **A `MessageOutputItem` indicates what from the LLM?**

- a) That a tool was invoked.
 - b) That a handoff occurred.
 - ✓c) **A message from the LLM.**
 - d) A reasoning item.
-

13. **What does a `HandoffCallItem` indicate?**

- ✓a) **That the LLM called the handoff tool.**
 - b) That a tool call output was received.
 - c) That a guardrail was tripped.
 - d) That the final output was generated.
-

14. **For a `HandoffOutputItem`, in addition to the raw item (tool response), what else can you access from it?**

- a) The LLM's full conversational history.
 - b) The `model_config` of the handoff agent.
 - ✓c) **The source and target agents.**
 - d) The `trace_id` of the handoff.
-

15. **Which `RunItem` indicates that the LLM invoked a tool?**

- a) `ToolResponseItem`
 - b) `ToolResultItem`
 - ✓c) **`ToolCallItem`**
 - d) `ToolOutputItem`
-

16. What information can you access from a `ToolCallOutputItem`?

- a) Only the raw tool call.
 - b) Only the LLM's reasoning.
 - ✓ c) **The tool output.**
 - d) The `model_settings` of the tool.
-

17. What does a `ReasoningItem` indicate?

- a) A message from a tool.
 - b) A call to a handoff.
 - ✓ c) **A reasoning item from the LLM.**
 - d) A final output from the agent.
-

18. Which properties contain the results of the guardrails, if any?

- a) `guardrail_status`
 - ✓ b) **`input_guardrail_results` and `output_guardrail_results`**
 - c) `check_results`
 - d) `validation_outcomes`
-

19. The `raw_responses` property contains what?

- a) Only user messages.
 - ✓ b) **The `ModelResponses` generated by the LLM.**
 - c) The final summarized output.
 - d) Intermediate tool call arguments.
-

20. What does the `input` property of a `RunResultBase` object contain?

- a) The processed input after guardrails.
 - b) The input generated by the LLM.
 - ✓ c) **The original input you provided to the run method.**
 - d) The inputs for the next turn.
-

Streaming

OpenAI Agents SDK - Streaming Section MCQs

Streaming & Stream Events – OpenAI Agents SDK

1. What is the primary benefit of using streaming in agent runs?

- a) To reduce the overall execution time of the agent.
- b) To enable offline processing of results.

- ✓c) To show the end-user progress updates and partial responses as the run proceeds.
 - d) To store all run data in a database.
-

2. Which Runner method must you call to enable streaming for an agent run?

- a) Runner.run()
 - b) Runner.run_sync()
 - ✓c) **Runner.run_streamed()**
 - d) Runner.stream_events()
-

3. When you call `Runner.run_streamed()`, what type of object is returned?

- a) RunResult
 - ✓b) **RunResultStreaming**
 - c) StreamResult
 - d) StreamingEvents
-

4. How do you access the asynchronous stream of `StreamEvent` objects from a `RunResultStreaming` object?

- a) By calling `result.get_stream()`.
 - b) By iterating directly over `result`.
 - ✓c) **By calling `result.stream_events()`.**
 - d) By subscribing to a global event bus.
-

5. What kind of events are `RawResponsesStreamEvent`?

- a) High-level events indicating tool calls.
 - ✓b) **Raw events passed directly from the LLM, in OpenAI Responses API format.**
 - c) Events indicating agent handoffs.
 - d) Debugging events for internal SDK processes.
-

6. `RawResponsesStreamEvent` objects contain data with a type and data. Which specific data type is used for streaming LLM text token-by-token?

- a) ResponseOutputEvent
 - b) LLMTextEvent
 - ✓c) **ResponseTextDeltaEvent**
 - d) TokenDeltaEvent
-

7. What is the primary use case for `RawResponsesStreamEvent`?

- a) To identify when a tool has completed its execution.
 - b) To determine which agent is currently active.
 - ✓c) **To stream response messages to the user as soon as they are generated (e.g., token-by-token).**
 - d) To capture the final output of the entire agent run.
-

8. `RunItemStreamEvents` are described as "higher level events." What specific kind of updates do they provide?

- a) Token-by-token generation.
 - ✓b) **Updates at the level of "message generated" or "tool ran".**
 - c) Changes in the LLM model used.
 - d) Debugging information about internal loop iterations.
-

9. **What information does an AgentUpdatedStreamEvent specifically provide?**

- a) The current LLM model's temperature settings.
 - b) The total number of turns completed by the agent.
 - ✓c) **Updates when the current agent changes (e.g., as a result of a handoff).**
 - d) The amount of time the agent has been running.
-

10. **In the example for RawResponsesStreamEvent, what is the `end=""` and `flush=True` in the print statement used for?**

- a) To append a newline character after each token.
 - b) To prevent any output from being printed until the end.
 - ✓c) **To print each token without a new line and immediately display it to the console.**
 - d) To store the tokens in a buffer for later processing.
-

11. **In the RunItemStreamEvents example, what is the instruction given to the Joker agent?**

- a) "Tell me exactly 5 jokes."
 - b) "Reply very concisely."
 - ✓c) **"First call the `how_many_jokes` tool, then tell that many jokes."**
 - d) "Generate a haiku about a clown."
-

12. **The `how_many_jokes` tool in the example returns a random integer. What Python module is imported for this functionality?**

- a) `math`
 - ✓b) **`random`**
 - c) `numpy`
 - d) `statistics`
-

13. **In the RunItemStreamEvents example, why is the `if event.type == "raw_response_event": continue` line used?**

- a) To explicitly process raw events first.
 - b) To throw an error if raw events are received.
 - ✓c) **To ignore token-level `raw_response_event` deltas and focus on higher-level events.**
 - d) To log raw response events to a file.
-

14. **When `event.type == "agent_updated_stream_event"` is true, what specific information is printed to the console?**

- a) The total number of agents.
- b) The previous agent's name.
- ✓c) **The name of the new agent.**
- d) The reason for the agent update.

15. What helper utility is used in the `RunItemStreamEvents` example to extract the text content from a `message_output_item`?

- a) `event.item.get_text()`
 - b) `event.item.message_content`
 - ✓ c) `ItemHelpers.text_message_output(event.item)`
 - d) `event.item.to_string()`
-

16. If `event.item.type == "tool_call_output_item"`, what information is specifically printed?

- a) Only that a tool was called.
 - ✓ b) The output received from the tool.
 - c) The name of the tool called.
 - d) The arguments passed to the tool.
-

17. What is the fundamental difference in granularity between `RawResponsesStreamEvent` and `RunItemStreamEvents`?

- a) One is for input, the other for output.
 - b) One is for synchronous runs, the other for asynchronous.
 - ✓ c) `RawResponsesStreamEvent` provides token-level updates, while `RunItemStreamEvents` provide updates for fully generated logical items (messages, tool outputs).
 - d) One is for debugging, the other for production.
-

18. What would happen if `asyncio.run(main())` were not used in the examples?

- a) The `Runner.run_streamed` call would become synchronous.
 - ✓ b) The asynchronous main function would not execute properly, as it needs an event loop.
 - c) The print statements would not display output.
 - d) The agent would automatically switch to synchronous mode.
-

19. The "Joker" agent's instructions ensure that a tool call (to `how_many_jokes`) occurs before generating jokes. This demonstrates streaming events in a scenario involving what?

- a) Only direct LLM responses.
 - ✓ b) Both LLM responses and tool interactions.
 - c) Only handoffs between agents.
 - d) Purely input validation.
-

20. After `result.stream_events()` is called and iterated through, what does the `RunResultStreaming` object still contain?

- a) Only the raw token stream.
- b) Only the last `StreamEvent`.
- ✓ c) The complete information about the run, including all new outputs produced.
- d) Only the final output of the agent.

REPL Utility

OpenAI Agents SDK - REPL Utility Section MCQs

✓REPL Utility — `run_demo_loop()`

1. What is the name of the utility provided by the SDK for quick interactive testing?
 - a) `Runner.run_interactive()`
 - ✓b) `run_demo_loop`
 - c) `Agent.repl()`
 - d) `interactive_shell`

2. What is the primary function of `run_demo_loop`?
 - a) To deploy agents to production.
 - b) To generate unit tests for agents.
 - ✓c) To prompt for user input in a loop, maintaining conversation history.
 - d) To evaluate agent performance metrics.

3. By default, how does `run_demo_loop` handle model output?
 - a) It stores all output in a file.
 - b) It prints the final output only after the loop exits.
 - ✓c) It streams model output as it is produced.
 - d) It sends output to an external logging service.

4. Which of the following commands or actions will *not* allow you to exit the `run_demo_loop`?
 - a) Typing `quit`.
 - b) Typing `exit`.
 - c) Pressing `Ctrl-D`.
 - ✓d) Pressing `Ctrl-C`.

Explanation: `quit`, `exit`, and `Ctrl+D` are explicitly handled by the REPL. `Ctrl+C` may raise a `KeyboardInterrupt` but is not the intended exit method per SDK docs.

5. In the provided example code, what Agent instruction is set for the "Assistant"?
 - a) "Always respond in haiku form."
 - b) "You are a helpful chatbot."
 - ✓c) "You are a helpful assistant."
 - d) "Provide help with math problems."

6. The `main` function in the REPL utility example is defined as `async def main()`. What standard Python module is used to run this asynchronous function?

- a) threading
 - b) concurrent
 - ✓ c) **asyncio**
 - d) multiprocessing
-

7. What is the main benefit of `run_demo_loop` “keeping the conversation history between turns”?
- a) It makes the agent run faster.
 - b) It disables external tool calls.
 - ✓ c) **It allows for continuous, context-aware interactions with the agent.**
 - d) It only stores the last user input.

Tools

OpenAI Agents SDK - Tools Section MCQs

Part 1: Overview and Hosted Tools

1. What is the fundamental role of "Tools" in the Agent SDK?
 - ✓ b) **To let agents take actions like fetching data, running code, or calling external APIs.**
 - ✗ a) To define the agent's instructions.
 - ✗ c) To manage conversation history.
 - ✗ d) To configure global run settings.

2. How many distinct classes of tools are there in the Agent SDK?
 - ✓ c) **Three**
 - ✗ a) One
 - ✗ b) Two
 - ✗ d) Four

3. Which class of tools runs on LLM servers alongside the AI models?
 - ✓ c) **Hosted tools**
 - ✗ a) Function calling tools
 - ✗ b) Agents as tools
 - ✗ d) Custom tools

4. Which of the following is NOT listed as a hosted tool offered by OpenAI?
 - ✓ d) **DatabaseQueryTool**
 - ✗ a) WebSearchTool
 - ✗ b) FileSearchTool
 - ✗ c) ComputerTool

5. **The FileSearchTool allows retrieving information from what specific OpenAI service?**
✓c) **OpenAI Vector Stores**
✗a) OpenAI API Gateway
✗b) OpenAI Embeddings Service
✗d) OpenAI Image Recognition API
6. **What is the purpose of the CodeInterpreterTool?**
✓c) **To let the LLM execute code in a sandboxed environment.**
✗a) To translate code
✗b) To debug production code
✗d) To generate comments
7. **In the Hosted tools example, which parameter specifies the file source?**
✓c) **vector_store_ids**
✗a) file_paths
✗b) document_ids
✗d) storage_buckets
8. **The ComputerTool enables agents to perform what kind of tasks?**
✓c) **Automating computer use tasks.**
✗a) Automating image generation
✗b) Spreadsheet automation
✗d) Web scraping
-

Part 2: Function Tools - Definition and Automation

9. **Which decorator turns a Python function into a tool?**
✓b) **@function_tool**
✗a) @agent_tool
✗c) @tool_callable
✗d) @make_tool
10. **How is a function tool's name determined by default?**
✓b) **It uses the name of the Python function.**
✗a) Explicitly provided
✗c) Unique ID
✗d) From the docstring
11. **Where is the tool description extracted from?**
✓b) **From the docstring of the function.**
✗a) Function arguments

- ☒ c) Config file
- ☒ d) name_override

12. Which Python module helps extract function signatures?

- ☒ b) inspect
- ☒ a) sys
- ☒ c) types
- ☒ d) os

13. Which library parses docstrings for descriptions?

- ☒ b) griffe
- ☒ a) Sphinx
- ☒ c) NumpyDocs
- ☒ d) Pydoc

14. What library is used for schema creation of inputs?

- ☒ c) pydantic
- ☒ a) jsonschema
- ☒ b) dataclasses
- ☒ d) typing

15. Which docstring formats are supported by griffe?

- ☒ b) google, sphinx, numpy
- ☒ a) JSON, YAML, TOML
- ☒ c) Markdown, reStructuredText
- ☒ d) PEP 257, Epytext

16. How to set a function tool's name explicitly?

- ☒ b) name_override="my_tool"
- ☒ a) tool_name="my_tool"
- ☒ c) Change variable name
- ☒ d) Set `tool.name`

17. How to prevent docstring parsing?

- ☒ c) use_docstring_info = False
- ☒ a) parse_docstring
- ☒ b) enable_docstring
- ☒ d) auto_docstring

18. In `fetch_weather`, what is Location?

- ☒ b) typing_extensions.TypedDict
- ☒ a) Pydantic BaseModel

- ☒ c) dict
 - ☒ d) dataclass
-

Part 3: Custom Function Tools & Agents as Tools

19. Required parameter for manual FunctionTool creation?

- ☒ c) `on_invoke_tool`
- ☒ a) `auto_schema`
- ☒ b) `docstring_format`
- ☒ d) `tool_priority`

20. Expected return type of `on_invoke_tool`?

- ☒ c) `str`
- ☒ a) Any
- ☒ b) dict
- ☒ d) None

21. What method parses JSON into `FunctionArgs`?

- ☒ c) `FunctionArgs.model_validate_json()`
- ☒ a) `parse_raw()`
- ☒ b) `from_json_string()`
- ☒ d) `load_json()`

22. Why model "agents as tools"?

- ☒ b) To let a central agent orchestrate other agents without a full handoff.
- ☒ a) Sequential execution
- ☒ c) Debugging
- ☒ d) Fewer LLM calls

23. Method to convert Agent into a tool?

- ☒ c) `agent.as_tool()`
- ☒ a) `agent.to_tool()`
- ☒ b) `agent.make_tool()`
- ☒ d) `agent.create_tool()`

24. What parameters define agent-as-tool presentation?

- ☒ c) `tool_name` and `tool_description`
- ☒ a) `input_type`, `output_type`
- ☒ b) `model_name`, `model_settings`
- ☒ d) `max_turns`, `run_config`

25. What is the limitation of `agent.as_tool()`?
- ✓b) It doesn't support full config like `max_turns`.
 - ✗a) Cannot use `FunctionTool`
 - ✗c) Only works with `run_sync`
 - ✗d) Prevents context access
26. What to do if `agent.as_tool()` config isn't enough?
- ✓c) Use `Runner.run` inside a custom function tool.
 - ✗a) Avoid agents-as-tools
 - ✗b) Use Handoff
 - ✗d) Override `as_tool`
-

Part 4: Output Extraction and Error Handling

27. What is `custom_output_extractor` for?
- ✓c) To reformat or modify the sub-agent's output.
 - ✗a) Summarize
 - ✗b) Translate
 - ✗d) Filter sensitive input
28. Why reverse `run_result.new_items`?
- ✓c) To find the latest JSON-like message efficiently.
 - ✗a) Chronological order
 - ✗b) Oldest first
 - ✗d) Memory efficiency
29. What `RunItem` is `custom_output_extractor` looking for?
- ✓b) `ToolCallOutputItem`
 - ✗a) `MessageOutputItem`
 - ✗c) `HandoffOutputItem`
 - ✗d) `ReasoningItem`
30. What happens if `failure_error_function` isn't defined?
- ✓b) Runs `default_tool_error_function` to inform LLM.
 - ✗a) Silently ignores
 - ✗c) Re-raises
 - ✗d) Retries
31. What happens if you pass `None` to `failure_error_function`?
- ✓c) Tool call errors are re-raised for manual handling.
 - ✗a) Silently ignored
 - ✗b) Returns empty
 - ✗d) Notifies LLM

32. What errors may be re-raised?
- ✓b) **ModelBehaviorError or UserError**
 - ✗a) NetworkError
 - ✗c) MemoryError
 - ✗d) AuthError
33. Where must you handle errors in manual FunctionTool?
- ✓c) **Inside `on_invoke_tool`**
 - ✗a) Global handler
 - ✗b) Outside
 - ✗d) Using `failure_error_function`
34. What does `ctx: RunContextWrapper[Any]` enable?
- ✓c) **Access to context-specific data or dependencies**
 - ✗a) Must be async
 - ✗b) Only orchestrator use
 - ✗d) Forces structured output
35. Role of `TypedDict` & `BaseModel`?
- ✓c) **Type hinting + schema generation/validation**
 - ✗a) Execution order
 - ✗b) DB integration
 - ✗d) API endpoint config

Model Context Protocol (MCP)

OpenAI Agents SDK - Model Context Protocol (MCP) Section MCQs

1. What is the primary function of the Model Context Protocol (MCP)?
- ✓c) **To standardize how applications provide context and tools to LLMs.**
 - ✗a) To manage the LLM's internal memory.
 - ✗b) To define a new type of LLM.
 - ✗d) To encrypt communication between agents.
2. MCP is compared to which technology for its standardization role?
- ✓c) **USB-C**
 - ✗a) Wi-Fi
 - ✗b) Bluetooth
 - ✗d) Ethernet

3. **How many kinds of MCP servers are defined based on transport mechanism?**

✓b) **Three**

✗a) Two

✗c) Four

✗d) Five

4. **Which type runs "locally" as a subprocess of your app?**

✓a) **stdio servers**

✗b) HTTP over SSE servers

✗c) Streamable HTTP servers

✗d) Remote servers

5. **Which server class connects using Streamable HTTP transport?**

✓c) **MCPServerStreamableHttp**

✗a) MCPServerStdio

✗b) MCPServerSse

✗d) MCPServerHttp

6. **Which command starts the official MCP filesystem server?**

✓c) **@modelcontextprotocol/server-fileSystem**

✗a) @modelcontextprotocol/server-http

✗b) @modelcontextprotocol/server-local

✗d) @modelcontextprotocol/server-console

7. **How are MCP servers integrated in an Agent instance?**

✓b) **Via the mcp_servers parameter.**

✗a) tools

✗c) model_settings

✗d) run_config

8. **What method is called to list available tools for the LLM?**

✓b) **list_tools()**

✗a) get_tools_list()

✗c) fetch_tools()

✗d) register_tools()

9. **What method is called when an MCP tool is invoked?**

✓b) **call_tool()**

✗a) execute_tool()

✗c) run_tool()

✗d) invoke_tool()

10. **Why implement caching for MCP servers?**
✓b) **To mitigate latency hits due to repeated list_tools() calls.**
✗a) Reduce memory use
✗c) Keep tools up-to-date
✗d) Simplify server code
11. **What parameter enables automatic tool list caching?**
✓b) **cache_tools_list**
✗a) enable_cache
✗c) auto_cache_tools
✗d) tool_cache_enabled
12. **When should you use automatic caching?**
✓c) **If you are certain the tool list will not change.**
✗a) For local servers
✗b) If tools change frequently
✗d) When debugging
13. **How to manually invalidate cached tools list?**
✓c) **invalidate_tools_cache()**
✗a) Restart the Agent
✗b) clear_cache()
✗d) set cache_tools_list=False
14. **What kind of server connects via a remote URL using SSE?**
✓b) **HTTP over SSE server**
✗a) stdio server
✗c) Streamable HTTP server
✗d) Localhost server
15. **What are command and args used for in MCPServerStdio?**
✓c) **To define the subprocess command/args that run the stdio server.**
✗a) Network address
✗b) Environment vars
✗d) Logging level
16. **Which MCP operations are captured in SDK tracing?**
✓b) **Calls to list_tools()**
✗a) Tool execution time only
✗c) Token generation
✗d) Input validation

17. What other operations include MCP tracing info?
- ✓b) Function calls
 - ✗a) Agent handoffs
 - ✗c) run_streamed events
 - ✗d) Agent initialization
18. Primary benefit of MCP for Agent developers?
- ✓c) Access a broad range of tools via standardized servers.
 - ✗a) Limits tools
 - ✗b) Simplifies FunctionTool creation
 - ✗d) Supports OpenAI tools only
19. Where to find complete examples of MCP usage?
- ✓c) examples/mcp
 - ✗a) examples/tools
 - ✗b) examples/handsoffs
 - ✗d) examples/basic
20. What challenge does MCP address?
- ✓b) Integrating LLMs with diverse external tools and data.
 - ✗a) Model accuracy
 - ✗c) LLM latency
 - ✗d) Prompt design

Handoffs

OpenAI Agents SDK - Handoffs Section MCQs

1. What is the primary purpose of "Handoffs"?
- ✓c) To allow an agent to delegate tasks to another specialized agent.
 - ✗a) Run code in sandbox
 - ✗b) Enable web search
 - ✗d) Generate images
2. How are handoffs represented to the LLM?
- ✓b) As tools
 - ✗a) Text instructions
 - ✗c) API calls
 - ✗d) State changes

3. **Default tool name for a Refund Agent handoff?**
✓c) **transfer_to_refund_agent**
✗a) refund_agent_tool
✗b) initiate_refund
✗d) handle_refunds
4. **Which Agent class parameter configures handoffs?**
✓c) **handoffs**
✗a) tools
✗b) delegations
✗d) sub_agents
5. **What else can `handoffs` accept besides an Agent?**
✓b) **A Handoff object**
✗a) FunctionTool
✗c) RunResult
✗d) RunContextWrapper
6. **Function to create a custom Handoff object?**
✓c) **handoff()**
✗a) create_handoff()
✗b) delegate_to()
✗d) define_handoff()
7. **Default tool name from `Handoff.default_tool_name()`?**
✓c) **transfer_to_<agent_name>**
✗a) handoff_to_agent_name
✗b) delegate_agent_name
✗d) handoff_agent_<name>
8. **What does `tool_name_override` do in `handoff()`?**
✓c) **Customize the tool name shown to LLM**
✗a) Change agent name
✗b) Specify return agent
✗d) Override tool description
9. **Purpose of `on_handoff` in `handoff()`?**
✓b) **Callback when handoff invoked, e.g., data fetch**
✗a) Prevent handoff
✗c) Redefine agent's instructions
✗d) Log output

10. `on_handoff` can optionally receive?
- ✓c) LLM-generated input (via `input_type`)
 - ✗a) Full conversation history
 - ✗b) Previous agent name
 - ✗d) `RunResult`
11. Purpose of `input_type` in `handoff()`?
- ✓b) Specify type of LLM-generated input to `on_handoff`
 - ✗a) Output type of target agent
 - ✗c) Filter user messages
 - ✗d) Async toggle
12. Class used in `EscalationData` example for input?
- ✓b) `BaseModel`
 - ✗a) `TypedDict`
 - ✗c) `dataclass`
 - ✗d) `NamedTuple`
13. How does the new agent see the conversation by default?
- ✓b) Sees the full previous history
 - ✗a) Starts fresh
 - ✗c) Only last user msg
 - ✗d) Only previous output
14. Purpose of `input_filter` in `handoff()`?
- ✓b) Modify input or filter history for new agent
 - ✗a) Validate output
 - ✗c) Define allowed handoffs
 - ✗d) Set target LLM
15. What argument does `input_filter` receive?
- ✓b) `HandoffInputData`
 - ✗a) `RunResult`
 - ✗c) `AgentContext`
 - ✗d) `List[dict]`
16. Where are common input filters like `remove_all_tools` found?
- ✓b) `agents.extensions.handoff_filters`
 - ✗a) `agents.core.filters`
 - ✗c) `agents.utils.input_filters`
 - ✗d) `agents.config.filters`

17. **Recommended practice for handoff clarity?**
- ✓c) **Include handoff info in agent instructions using recommended prompts**
 - ✗a) input_filter=None
 - ✗b) Always define tool_name_override
 - ✗d) Limit handoffs
18. **What is RECOMMENDED_PROMPT_PREFIX?**
- ✓c) **A suggested prefix for better LLM handoff understanding**
 - ✗a) Dynamic prompt generator
 - ✗b) Parameter to disable prompts
 - ✗d) Default handoff name
19. **Function to auto-include handoff info in prompts?**
- ✓c) **prompt_with_handoff_instructions()**
 - ✗a) add_handoff_instructions()
 - ✗b) generate_handoff_prompt()
 - ✗d) configure_handoff_prompt()
20. **Main benefit of specialized agents via handoffs?**
- ✓c) **Better specialization & modularity**
 - ✗a) Reduce LLM calls
 - ✗b) Make agents less modular
 - ✗d) Simpler debugging
21. **In Triage agent example – which agent uses which handoff method?**
- ✓b) **billing_agent uses Agent directly, refund_agent uses handoff()**
 - ✗a) refund_agent uses Agent directly
 - ✗c) both use handoff()
 - ✗d) both use Agent
22. **Purpose of tool_description_override?**
- ✓c) **Customize description seen by LLM**
 - ✗a) Shorten description
 - ✗b) Auto generate
 - ✗d) Describe inputs
23. **What does EscalationData's on_handoff show?**
- ✓c) **Trigger structured actions/logs from LLM input**
 - ✗a) Auto retraining
 - ✗b) Passive observation
 - ✗d) Disable further tool calls

24. **Why might an input_filter remove past tool calls?**

- ✓c) **Focus next agent only on new input/convo**
- ✗a) Block tool usage
- ✗b) Reduce tokens
- ✗d) Speed handoff

25. **What does modeling handoffs as tools enable?**

- ✓c) **LLM decides when/whom to delegate to**
- ✗a) Modify internal logic
- ✗b) Only run pre-defined functions
- ✗d) Bypass guardrails

Tracing

OpenAI Agents SDK - Tracing Section MCQs

📄 Part 1: Overview & Trace/Span Fundamentals

1. **Primary purpose of built-in tracing:**

- ✓b) **To collect a comprehensive record of events for debugging, visualization, and monitoring.**
- ✗a) Optimize performance
- ✗c) Version control
- ✗d) Encrypt data

2. **NOT collected by tracing:**

- ✓d) **Agent instruction changes**
- ✗a) LLM generations
- ✗b) Tool calls
- ✗c) Handoffs

3. **Is tracing enabled by default?**

- ✓b) **Enabled by default.**
- ✗a) Disabled
- ✗c) OS-dependent
- ✗d) Production-only

4. **Globally disable tracing?**

- ✓b) **Set `OPENAI_AGENTS_DISABLE_TRACING=1`.**
- ✗a) RunConfig
- ✗c) `trace.disable_all()`
- ✗d) Not possible

5. **Tracing under Zero Data Retention (ZDR)?**
☒ c) **Tracing is unavailable.**
☐ a) Auto-configured
☐ b) Redacted traces
☐ d) Custom processors required
6. **What is a Trace?**
☒ c) **A single end-to-end operation or "workflow".**
☐ a) LLM call
☐ b) Agent action
☐ d) Error log
7. **Required Trace property:**
☒ c) **workflow_name**
☐ a) started_at
☐ b) parent_id
☐ d) disabled
8. **Trace ID format:**
☒ b) **trace_<32_alphanumeric>**
☐ a) trace-<UUID>
☐ c) trace_YYYYMMDD_HHMMSS
☐ d) Any string
9. **group_id purpose:**
☒ c) **Link multiple traces from same conversation/flow.**
☐ a) Agent type
☐ b) Severity
☐ d) Duration limit
10. **What are spans?**
☒ b) **Operations with start and end times.**
☐ a) Workflow duration
☐ c) Trace metadata
☐ d) Final output
11. **Property pointing to parent span:**
☒ b) **parent_id**
☐ a) trace_id
☐ c) group_id
☐ d) span_data

12. What does `GenerationSpanData` contain?

- ✓c) LLM generation details.
 - ✗a) Agent init
 - ✗b) Tool I/O
 - ✗d) Guardrails
-

📖 Part 2: Default Tracing & Creating Traces/Spans

13. How are `Runner.run` methods traced?

- ✓b) Wrapped in a `trace()`
- ✗a) `custom_span`
- ✗c) Not traced
- ✗d) `agent_span`

14. Span type wrapping agent run:

- ✓c) `agent_span()`
- ✗a) `run_span`
- ✗b) `agent_exec_span`
- ✗d) `llm_call_span`

15. Default `workflow_name` if not set:

- ✓c) Agent trace
- ✗a) Default workflow
- ✗b) SDK Trace
- ✗d) Unnamed trace

16. Recommended way to create trace:

- ✓b) Use `trace()` as a context manager.
- ✗a) Manual start/finish
- ✗c) `RunConfig`
- ✗d) `asyncio.run()`

17. Update trace contextvars manually:

- ✓b) `mark_as_current` and `reset_current`
- ✗a) `set_current`
- ✗c) `start_current`
- ✗d) `activate`

18. Need to manually create spans?

- ✓b) No, most common ops are traced automatically.
- ✗a) Yes

- ☒ c) Only LLM
- ☒ d) Only handoffs

19. Track custom span info:

- ☒ b) `custom_span()`
- ☒ a) `log_span`
- ☒ c) `new_span`
- ☒ d) `record_span`

20. How are custom spans nested?

- ☒ c) Using Python contextvars for current trace/span.
- ☒ a) Explicit `parent_id`
- ☒ b) Global var
- ☒ d) RunConfig flag

Part 3: Sensitive Data & Custom Processors

21. Span storing LLM I/O (may be sensitive):

- ☒ c) `generation_span()`
- ☒ a) `agent_span`
- ☒ b) `tool_call_span`
- ☒ d) `handoff_span`

22. Disable sensitive data capture:

- ☒ c) `RunConfig.trace_include_sensitive_data = False`
- ☒ a) Env variable
- ☒ b) `metadata.sensitive`
- ☒ d) Use `custom_span`

23. Audio spans store (by default):

- ☒ c) Base64-encoded PCM data.
- ☒ a) File paths
- ☒ b) Transcripts
- ☒ d) Metadata

24. Disable sensitive audio capture:

- ☒ c) `VoicePipelineConfig.trace_include_sensitive_audio_data = False`
- ☒ a) RunConfig
- ☒ b) `trace_include_sensitive_data`
- ☒ d) Not possible

25. **Component that creates traces:**

- ☒ **b) TraceProvider**
- ☐ a) TraceManager
- ☐ c) SpanFactory
- ☐ d) TraceGenerator

26. **BatchTraceProcessor sends traces to:**

- ☒ **c) BackendSpanExporter**
- ☐ a) Log file
- ☐ b) Remote DB
- ☐ d) StreamEvent

27. **BackendSpanExporter purpose:**

- ☒ **b) Export spans/traces to OpenAI backend.**
- ☐ a) Process sensitive data
- ☐ c) Human-readable output
- ☐ d) Manage span lifecycle

28. **Add additional processor:**

- ☒ **b) add_trace_processor()**
- ☐ a) set_trace_processors
- ☐ c) register_trace_processor
- ☐ d) configure_trace_processor

29. **Consequence of set_trace_processors () use:**

- ☒ **c) You must include OpenAI exporter explicitly.**
- ☐ a) Auto includes both
- ☐ b) Custom must forward
- ☐ d) Local-only effect

30. **NOT a supported tracing backend:**

- ☒ **d) Prometheus**
- ☐ a) Weights & Biases
- ☐ b) LangSmith
- ☐ c) MLflow

31. **Why are multiple runs inside trace("Joke workflow")?**

- ☒ **c) To group multiple Runner.run () calls into one trace.**
- ☐ a) Trace final only
- ☐ b) Disable tracing
- ☐ d) Prevent concurrency

32. **Span for audio input (speech-to-text):**

- ✓c) **transcription_span()**
- ✗a) audio_input_span
- ✗b) speech_to_text_span
- ✗d) input_audio_span

33. **Span for audio output (text-to-speech):**

- ✓b) **speech_span()**
- ✗a) audio_output_span
- ✗c) text_to_speech_span
- ✗d) output_audio_span

34. **Parent span for related audio spans:**

- ✓c) **speech_group_span()**
- ✗a) audio_workflow_span
- ✗b) voice_interaction_span
- ✗d) composite_audio_span

35. **Purpose of Traces dashboard:**

- ✓c) **Debugging, visualizing, monitoring workflows.**
- ✗a) Raw data collection
- ✗b) Deployment management
- ✗d) Model optimization

Context Management

OpenAI Agents SDK - Context Management Section MCQs

📁 Part 1: Local Context

1. **Two main classes of context?**

- ✓c) **Context available locally to your code and Context available to LLMs**
- ✗a) Input/Output
- ✗b) Static/Dynamic
- ✗d) User/System Context

2. **Class representing local context?**

- ✓c) **RunContextWrapper**
- ✗a) AgentContext
- ✗b) LLMContext
- ✗d) LocalContext

3. **How to pass custom context to an agent run?**
✓c) As a context keyword argument to Runner.run() methods
✗a) Agent constructor
✗b) Env variable
✗d) Global config
4. **Accessing your context object from a wrapper?**
✓b) wrapper.context
✗a) get_context()
✗c) wrapper.data
✗d) wrapper.T
5. **Most important thing to know about context objects?**
✓b) Every agent, tool, and hook must use the same type
✗a) Must be immutable
✗c) Auto-sent to LLM
✗d) Only primitives allowed
6. **Common use case for local context?**
✓c) Carrying contextual data like UID, logger, etc.
✗a) Instructions
✗b) LLM history
✗d) Grounding responses
7. **Is local context visible to the LLM?**
✓c) No, purely for local use
✗a) Yes, always
✗b) If configured
✗d) If Pydantic
8. **UserInfo structure in example?**
✓c) dataclasses.dataclass
✗a) NamedTuple
✗b) BaseModel
✗d) TypedDict
9. **First argument of function tool needing context?**
✓c) RunContextWrapper[UserInfo] (or RunContextWrapper[T])
✗a) RunContext
✗b) Context
✗d) Any

10. **What does the LLM see for generation?**

- ✓c) **The conversation history**
- ✗a) Local context
- ✗b) RunConfig
- ✗d) Agent name/tools

11. **To make data visible to LLM?**

- ✓b) **Include in conversation history**
- ✗a) Direct function call
- ✗c) API call
- ✗d) Encrypted metadata

12. **System prompt = developer message = ?**

- ✓c) **Adding it to Agent instructions**
- ✗a) Input to Runner.run
- ✗b) Function tools
- ✗d) Retrieval/web search

13. **Instructions can be:**

- ✓c) **Static or dynamic functions returning strings**
- ✗a) Static only
- ✗b) Dynamic only
- ✗d) Pydantic only

14. **Common data for instructions:**

- ✓b) **Always useful info like name/date**
- ✗a) On-demand info
- ✗c) Computation results
- ✗d) Confidential data

15. **Runner.run input vs Agent instructions:**

- ✓b) **Input appears lower in prompt hierarchy**
- ✗a) Preferred for static info
- ✗c) Invisible to LLM
- ✗d) For sensitive data only

16. **Function tools are for what context type?**

- ✓c) **On-demand context**
- ✗a) Static

- ☒ b) Pre-loaded
- ☒ d) Internal logging

17. When LLM *might* need data? Use:

- ☒ c) **Function tools**
- ☒ a) Instructions
- ☒ b) context=...
- ☒ d) input_filter

18. Retrieval / web search is for:

- ☒ c) **Fetching relevant data to ground responses**
- ☒ a) Local context
- ☒ b) Agent state
- ☒ d) History encryption

19. Grounding supported by:

- ☒ c) **Retrieval or web search**
- ☒ a) run_config.metadata
- ☒ b) RunContextWrapper
- ☒ d) on_handoff

20. Tell LLM current datetime always:

- ☒ b) **Include it dynamically in Agent instructions**
- ☒ a) Use a function tool
- ☒ c) context=...
- ☒ d) Add to final_output

21. Which is true about RunContextWrapper?

- ☒ b) **Can pass dependencies like loggers**
- ☒ a) Prevents tool calls
- ☒ c) Auto-added to LLM
- ☒ d) Always visible in Traces

22. Main takeaway from local vs LLM context?

- ☒ c) **Clear separation between internal and LLM data**
- ☒ a) All data must be visible to LLM
- ☒ b) LLM accesses Python objects
- ☒ d) Local = debug, LLM = prod

23. Agent[UserInfo] + context=ProductInfo → ?

- ☒ c) **Error or unexpected behavior due to type mismatch**
- ☒ a) Adapts automatically

- ✗b) Ignores context
- ✗d) ProductInfo overrides

24. **Dynamic Agent instructions allow for:**

- ✓c) **Personalized system prompts using local context**
- ✗a) Real-time API calls
- ✗b) Change LLM model
- ✗d) Disable tracing

25. **Best way to enhance LLM reasoning w/ external info:**

- ✓c) **Inject into conversation history or expose via tools**
- ✗a) Modify model weights
- ✗b) Use hosted tools only
- ✗d) Rely on pre-trained knowledge

Guardrails

OpenAI Agents SDK - Guardrails Section MCQs

📖 Part 1: General Concepts and Types

1. **How do guardrails run in relation to agents?**

- ✓b) **In parallel to your agents**
- ✗a) Sequentially
- ✗c) Only after
- ✗d) As sub-process

2. **Key benefit of guardrails (slow/expensive model example)?**

- ✓b) **Immediately halt agent execution and save time/money if misuse is detected**
- ✗a) Enhance reasoning
- ✗c) Detailed tracing
- ✗d) Rephrase input

3. **How many guardrail types are described?**

- ✓b) **Two**

- ☒ a) One
- ☒ c) Three
- ☒ d) Four

4. Which runs on initial user input?

- ☒ c) **Input guardrails**
- ☒ a) Output
- ☒ b) Pre-processing
- ☒ d) User

5. Which runs on final agent output?

- ☒ c) **Output guardrails**
- ☒ a) Input
- ☒ b) Post-processing
- ☒ d) Final

Part 2: Input and Output Guardrail Mechanics

6. What does an Input Guardrail receive first?

- ☒ a) **The same input passed to the agent**
- ☒ b) Agent instructions
- ☒ c) LLM output
- ☒ d) History

7. Input guardrail result is wrapped in?

- ☒ b) **InputGuardrailResult**
- ☒ a) GuardrailResult
- ☒ c) GuardrailOutput
- ☒ d) InputCheckResult

8. If `.tripwire_triggered` is true → ?

- ☒ c) **Raises InputGuardrailTripwireTriggered exception**
- ☒ a) Logs and continues
- ☒ b) Rephrases input
- ☒ d) Reroutes agent

9. When do input guardrails run?

- ☒ c) **Only if the agent is the first in the flow**
- ☒ a) If tools enabled
- ☒ b) If streaming
- ☒ d) If input has sensitive data

10. **Why are guardrails attached to Agent object?**
✓c) **Because they relate to the specific agent (better readability)**
✗a) Enforce global policy
✗b) Simplify run()
✗d) Modify state
11. **Key difference in input vs output guardrail execution?**
✓c) **Input: on user input; Output: on final agent output**
✗a) Tool timing
✗b) Sync/async
✗d) Fast vs slow model
12. **When do output guardrails run?**
✓c) **Only if the agent is the last agent in the execution flow**
✗a) JSON output
✗b) Handoff
✗d) LLM message
-

📖 Part 3: Tripwires and Implementation

13. **What does a tripwire signal?**
✓c) **Input/output failed guardrail check**
✗a) Success
✗b) Ready to rerun
✗d) New type detected
14. **Effect of triggered tripwire?**
✓c) **Halts execution, raises exception**
✗a) Self-corrects
✗b) Logs warning
✗d) Re-prompts user
15. **Required return type for guardrail functions?**
✓d) **GuardrailFunctionOutput**
✗a) bool
✗b) str
✗c) dict
16. **How is math homework checked in example?**
✓c) **Runs another Agent (guardrail agent)**
✗a) Keyword search
✗b) External API
✗d) Forbidden phrase DB

17. **How is `.tripwire_triggered` determined?**
✓c) From `final_output`'s `is_math_homework` / `is_math` boolean
✗a) Hardcoded
✗b) Always False
✗d) Context state
18. **Python construct to handle guardrail trip?**
✓c) `try...except` for `TripwireTriggered` exception
✗a) `if/else`
✗b) `while` loop
✗d) `assert`
19. **Input types accepted by `math_guardrail`?**
✓c) `str` or `list[TResponseInputItem]`
✗a) `str` only
✗b) `list` only
✗d) Any
20. **Decorator for input guardrail?**
✓c) `@input_guardrail`
✗a) `@guardrail`
✗b) `@agent_guardrail`
✗d) `@check_input`
21. **Decorator for output guardrail?**
✓d) `@output_guardrail`
✗a) `@guardrail`
✗b) `@agent_guardrail`
✗c) `@check_output`
22. **Why is `OutputGuardrailTripwireTriggered` raised?**
✓c) To prevent finalizing/sending non-compliant output
✗a) Retry
✗b) Alt destination
✗d) Help LLM understand output
23. **Why use a separate agent for guardrails?**
✓b) To leverage LLM for complex logic
✗a) Must use agents
✗c) Fewer deps
✗d) Sync execution

24. What's in `output_info` of `GuardrailFunctionOutput`?

- ✓c) **final_output** from internal guardrail agent
- ✗a) Boolean
- ✗b) Error msg
- ✗d) Raw input

25. Why run guardrails “in parallel”?

- ✓b) **Efficiency** – avoid costly operations if early issues
- ✗a) Code complexity
- ✗c) Model choice
- ✗d) UI responsiveness

Orchestrating Multiple Agents

OpenAI Agents SDK - Orchestrating Multiple Agents Section MCQs

🔍 Part 1: Overview and Orchestrating via LLM

1. What does "Orchestration" refer to in agent apps?

- ✓b) **The flow of agents, including their order and decision-making for what happens next.**
 - ✗a) Deployment process
 - ✗c) LLM memory management
 - ✗d) Training agent models
-

2. Two main ways to orchestrate agents:

- ✓c) **Allowing the LLM to make decisions and Orchestrating via code.**
 - ✗a) Synchronous and Asynchronous
 - ✗b) Local and Remote
 - ✗d) Sequential and Parallel
-

3. What defines an LLM-based agent?

- ✓c) **Instructions, Tools, and Handoffs.**
 - ✗a) Memory, Sensors, Actuators
 - ✗b) Models, Data, Logging
 - ✗d) Prompts, Outputs, Traces
-

4. How does LLM tackle open-ended tasks?

- ✓b) **Autonomously plans, reasons, and uses tools/handoffs.**

- ☒ a) Follows script
 - ☒ c) Always asks for human help
 - ☒ d) Randomly chooses
-

5. Which is NOT a tool for LLM agents?

- ☒ d) **Model Fine-tuning.**
 - ☒ a) Web search
 - ☒ b) File search
 - ☒ c) Computer use
-

6. When is LLM orchestration ideal?

- ☒ b) **Open-ended tasks where LLM reasoning is needed.**
 - ☒ a) Deterministic tasks
 - ☒ c) Low-resource tasks
 - ☒ d) Rule-based logic
-

7. Important tactic in LLM orchestration:

- ☒ c) **Use good prompts to define tools and parameters.**
 - ☒ a) Fewer agents
 - ☒ b) Avoid tools
 - ☒ d) Disable monitoring
-

8. Meaning of "introspect and improve":

- ☒ b) **Run in a loop, self-critique, and learn from errors.**
 - ☒ a) Manually update code
 - ☒ c) Restrict access
 - ☒ d) Give only positive feedback
-

9. Why use specialized agents?

- ☒ b) **So each agent excels in one thing.**
 - ☒ a) Makes debugging harder
 - ☒ c) Reduces agent count
 - ☒ d) Forces handoffs
-

10. Why "Invest in evals"?

- ☒ b) **Helps train and improve agent performance.**
- ☒ a) Reduces LLM cost
- ☒ c) Eases deployment
- ☒ d) Real-time error fixing

□ □ Part 2: Orchestrating via Code

11. Benefit of code orchestration:

- ✓b) More deterministic and predictable (speed, cost).
 - ✗a) Flexibility
 - ✗c) Fewer prompts needed
 - ✗d) Easier LLM integration
-

12. Structured output use case:

- ✓c) Classify task, then choose agent.
 - ✗a) Write full blog
 - ✗b) Parallel tasks
 - ✗d) Self critique
-

13. What is "chaining agents"?

- ✓c) Output of one → Input of next agent.
 - ✗a) Run simultaneously
 - ✗b) One agent, many jobs
 - ✗d) Recursive calls
-

14. Chaining example:

- ✓c) Writing a blog post (research → write → critique).
 - ✗a) Customer service
 - ✗b) Math problems
 - ✗d) Translation
-

15. Pattern using a while loop:

- ✓c) Iterative feedback until criteria met.
 - ✗a) Parallel tasks
 - ✗b) Task classification
 - ✗d) Structured data creation
-

16. Python method for parallel agents:

- ✓c) `asyncio.gather`
- ✗a) `threading.Thread`
- ✗b) `multiprocessing.Process`

- ☒ d) ThreadPoolExecutor
-

17. When to use parallel agents:

- ☒ b) When tasks are independent and need speed.
 - ☒ a) Tasks depend on each other
 - ☒ c) Let LLM decide
 - ☒ d) Debugging
-

18. Why mix LLM and code orchestration?

- ☒ c) Combines LLM reasoning with code precision.
 - ☒ a) You must pick one
 - ☒ b) Makes design harder
 - ☒ d) Only works for simple tasks
-

19. Where are orchestration examples?

- ☒ c) In `examples/agent_patterns` folder.
 - ☒ a) SDK README
 - ☒ b) Blog post
 - ☒ d) API docs
-

20. Best for precise, cost-controlled tasks:

- ☒ b) Orchestrating via code.
- ☒ a) LLM-based
- ☒ c) LLM-heavy hybrid
- ☒ d) Human-in-loop

Models

Models

OpenAI Agents SDK - Models Section MCQs

🔖 Part 1: OpenAI Models and Non-OpenAI Integration

1. What is the recommended model class to use with OpenAI in the SDK?
 - a) OpenAIChatCompletionsModel
 - ✓b) **OpenAIResponsesModel**
 - c) OpenAITextCompletionModel
 - d) OpenAILegacyModel
2. Which API does OpenAIChatCompletionsModel use?
 - a) Responses API
 - ✓b) **Chat Completions API**
 - c) Embeddings API
 - d) Moderations API
3. What dependency is needed to use most non-OpenAI models?
 - a) openai-agents[all]
 - b) openai-agents[models]
 - ✓c) **openai-agents[litellm]**
 - d) openai-agents[external]
4. What prefix is used when integrating with LiteLLM models?
 - a) ollama/
 - b) external/
 - ✓c) **litellm/**
 - d) custom/
5. How can you integrate other LLM providers (besides LiteLLM)?
 - a) Modify Runner constructor
 - b) Edit SDK source code
 - ✓c) **Use `set_default_openai_client`**
 - d) Tool-level ModelFactory
6. `set_default_openai_client()` is useful when the endpoint is:
 - a) Proprietary API
 - ✓b) **OpenAI-compatible API endpoint**
 - c) XML-based
 - d) GraphQL-based
7. Where is ModelProvider applied to cover all agents in one run?
 - a) Tool function level
 - ✓b) **Runner.run level**
 - c) Global app lifecycle
 - d) Inside Agent constructor

8. How do you assign different LLM providers to different agents?
- a) `set_default_openai_client`
 - b) `model_provider` in `Runner`
 - ✓ c) **Set `Agent.model` per agent**
 - d) Use only LiteLLM
9. What to do if you don't have an OpenAI API key?
- a) Use dummy key
 - ✓ b) **Disable tracing via `set_tracing_disabled()`**
 - c) Nothing needed
 - d) Tracing adjusts automatically
10. Which API is often used with non-OpenAI models in examples?
- a) Embeddings API
 - ✓ b) **Chat Completions API/model**
 - c) Completions API
 - d) Fine-tuning API
-

📖 Part 2: Mixing & Matching Models

11. Why use different models per agent?
- a) Increase latency
 - b) Add security complexity
 - ✓ c) **Use smaller model for triage, larger for complex tasks**
 - d) Limit agents
12. Which is NOT a valid way to select a model?
- a) Provide model name like "gpt-3.5-turbo"
 - b) Use `ModelProvider` with name
 - c) Use `OpenAIChatCompletionsModel(...)`
 - ✓ d) **Write model in Agent instructions**
13. Why use a single model shape per workflow?
- a) Ensure pricing consistency
 - b) Simplify logging
 - ✓ c) **Model shapes support different tools/features**
 - d) Optimize disk
14. What to ensure if mixing model shapes?
- a) Don't mix
 - b) Only use Responses

- c) Use only Chat
- ✓d) **Features must exist in both shapes**

15. Purpose of `trriage_agent` in model mix example?

- a) Analyze data
- b) Multilingual generation
- ✓c) **Route to agent based on language**
- d) Evaluate other agents

16. Which param is used for settings like temperature?

- a) `model_config`
- b) `llm_settings`
- ✓c) **`model_settings`**
- d) `config_params`

17. Effect of `temperature=0.1`?

- a) More creative
- ✓b) **More deterministic, less random**
- c) Increase length
- d) Reduce tokens

🔗 Part 3: Issues with Other LLM Providers

18. Cause of “Tracing client error 401”?

- a) Model not compatible
- ✓b) **No OpenAI API key for tracing**
- c) Bad internet
- d) Temp too high

19. How to fix tracing error 401?

- a) Set empty key
- b) Use `OpenAIResponsesModel`
- ✓c) **Call `set_tracing_export_api_key(...)` with OpenAI key**
- d) Update agent instructions

20. Fix for 404 error due to Responses API unsupported?

- a) Disable tracing
- b) Change `litellm` version
- ✓c) **Use `set_default_openai_api("chat_completions")` or switch to `OpenAIChatCompletionsModel`**
- d) Set custom `ModelProvider`

21. Error message when json_schema not supported?
- a) Host not found
 - b) Invalid model name
 - ✓c) **BadRequestError: 'response_format.type'...**
 - d) Unexpected keyword
22. SDK suggestion when provider lacks JSON schema support?
- a) Manually parse text
 - b) Set temp to 0
 - ✓c) **Use provider that supports JSON schema**
 - d) Reduce input size
23. What to consider when mixing providers?
- a) All features same
 - b) Performance same
 - ✓c) **Feature differences like multimodal & structured output**
 - d) API keys interchangeable
24. Precaution with text-only models?
- a) Long inputs
 - b) JSON inputs
 - ✓c) **Filter out multimodal inputs**
 - d) One-word prompts
25. OpenAI API key for tracing must be from:
- a) Current provider
 - ✓b) **platform.openai.com**
 - c) Any cloud provider
 - d) Local only
26. SDK model interactions are mainly:
- a) Synchronous
 - ✓b) **Asynchronous**
 - c) Blocking
 - d) Multi-threaded
27. How does SDK get API key & URL for OpenAI?
- a) Passed manually
 - b) Hardcoded
 - ✓c) **From OPENAI_API_KEY & OPENAI_BASE_URL env vars**
 - d) Generates each time

28. Example parameter in ModelSettings:
- a) api_key
 - b) base_url
 - c) model_name
 - ✓d) **temperature**
29. Result if provider doesn't support json_schema?
- a) Output is converted
 - b) Agent uses OpenAI
 - ✓c) **App breaks with malformed JSON**
 - d) Model gives empty string
30. Where are example configs for non-OpenAI models?
- a) examples/models/
 - b) examples/providers/
 - ✓c) **examples/model_providers/**
 - d) examples/custom_models/

Using any model via LiteLLM

OpenAI Agents SDK - Using any Model via LiteLLM Section MCQs

📖 Using Any Model via LiteLLM – MCQs

1. What is the current status of the LiteLLM integration in the Agents SDK?
 - a) Stable and production-ready
 - ✓b) **In beta**
 - c) Deprecated
 - d) Under experimental development

2. If you encounter issues with the LiteLLM integration, especially with smaller model providers, what is the recommended action?
 - a) Revert to using OpenAI models only
 - b) Wait for a new SDK version automatically
 - ✓c) **Report the issues via GitHub issues**
 - d) Try a different LiteLLM version

3. What is LiteLLM described as?
 - a) A new type of LLM from OpenAI
 - ✓b) **A library that allows you to use 100+ models via a single interface**

- c) A debugging tool for agents
 - d) An SDK for creating custom tools
4. To enable the LiteLLM integration, which specific command should you run?
- a) `pip install litellm`
 - b) `pip install openai-agents`
 - ✓ c) **`pip install "openai-agents[litellm]"`**
 - d) `pip install openai-agents-models`
5. Once LiteLLM is installed, which class should you use within an Agent to specify a LiteLLM model?
- a) `OpenAIResponsesModel`
 - b) `LitellmClient`
 - ✓ c) **`LitellmModel`**
 - d) `ExternalModel`
6. In the provided example, when running the script, what information will you be prompted to enter?
- a) Only the model name
 - b) Only the API key
 - ✓ c) **A model name and an API key**
 - d) The agent's instructions
7. Which of the following is an example of a model name format for LiteLLM given in the documentation?
- a) `gpt-4.1`
 - ✓ b) **`openai/gpt-4.1`**
 - c) `litellm_gpt-4.1`
 - d) `gpt-4.1-openai`
8. Where can you find a full list of models supported by LiteLLM?
- a) In the Agents SDK documentation
 - b) On the OpenAI API reference page
 - ✓ c) **In the LiteLLM providers docs**
 - d) Within the `agents.extensions.models` module
9. In the example `main()` function, what are the `model` and `api_key` arguments passed to when creating the `LitellmModel` instance?
- a) `name` and `instructions`
 - b) `tools` and `handoffs`
 - ✓ c) **`model` and `api_key`**
 - d) `temperature` and `top_p`
10. What specific instruction is given to the Assistant agent in the example, which dictates its response style?
- a) `"You are a helpful assistant."`

- b) "You only speak English."
 - ✓c) **"You only respond in haikus."**
 - d) "You must use the get_weather tool."
11. The `get_weather` function in the example is decorated with `@function_tool`. This means it will be exposed to the LLM as what?
- a) An internal utility
 - b) A local context provider
 - ✓c) **A callable tool**
 - d) A guardrail
12. The example code includes `set_tracing_disabled()`. What is the likely reason for including this in a LiteLLM example?
- a) Tracing is not compatible with LiteLLM models
 - b) It's a standard practice for all agent examples
 - ✓c) **Users might not have an OpenAI API key for tracing, which is needed by default**
 - d) To improve the performance of LiteLLM models
13. How does the example script obtain the model and `api_key` if they are not provided via command-line arguments?
- a) It uses predefined environment variables
 - b) It fetches them from a configuration file
 - ✓c) **It prompts the user for input using `input()`**
 - d) It generates them dynamically
14. What library is used in the example to handle asynchronous execution of the main function?
- a) `threading`
 - b) `multiprocessing`
 - ✓c) **`asyncio`**
 - d) `concurrent.futures`
15. The LiteLLM integration allows you to use `anthropic/claude-3-5-sonnet-20240620`. This demonstrates the ability to use models from providers other than just OpenAI, such as:
- a) Google
 - b) Cohere
 - ✓c) **Anthropic**
 - d) Hugging Face

Configuring the SDK

OpenAI Agents SDK - Configuring the SDK Section MCQs

Part 1: API Keys and Clients

1. **By default, where does the Agents SDK look for the OpenAI API key for LLM requests and tracing?**
 - a) In a config.ini file
 - b) Directly from ~/.openai/credentials
 - ✓c) **In the OPENAI_API_KEY environment variable**
 - d) As a command-line argument

 2. **If you cannot set the OPENAI_API_KEY environment variable before your app starts, which function can you use to configure it?**
 - a) set_openai_key()
 - ✓b) **set_default_openai_key()**
 - c) configure_api_key()
 - d) set_env_variable()

 3. **By default, what type of OpenAI client instance does the SDK create?**
 - a) OpenAI (synchronous)
 - ✓b) **AsyncOpenAI**
 - c) SyncOpenAI
 - d) OpenAIClient

 4. **To use a custom AsyncOpenAI instance (e.g., with a different base_url), which function should you use?**
 - a) set_custom_client()
 - b) configure_openai_instance()
 - ✓c) **set_default_openai_client()**
 - d) override_openai_client()

 5. **What is the default OpenAI API that the Agents SDK uses?**
 - a) The Chat Completions API
 - ✓b) **The Responses API**
 - c) The Embeddings API
 - d) The Legacy Completions API

 6. **To override the default OpenAI API to use the Chat Completions API, which function is available?**
 - a) use_chat_completions_api()
 - b) set_openai_api_type("chat_completions")
 - ✓c) **set_default_openai_api("chat_completions")**
 - d) configure_api_version()
-

7. **What is the default state of tracing in the Agents SDK?**
 - ✓ a) **Enabled**
 - b) Disabled
 - c) It depends on the environment variable
 - d) It requires manual configuration for each run

8. **By default, where does tracing get its API key from?**
 - a) A separate OPENAI_TRACING_KEY environment variable
 - b) A dedicated tracing configuration file
 - ✓ c) **The same OpenAI API keys used for LLM requests (environment variable or `set_default_openai_key()`)**
 - d) It generates a temporary key for each trace

9. **Which function allows you to set a specific API key to be used only for tracing, separate from LLM requests?**
 - a) `set_tracing_key()`
 - b) `configure_trace_api()`
 - ✓ c) **`set_tracing_export_api_key()`**
 - d) `override_trace_key()`

10. **How can you disable tracing entirely in the SDK?**
 - a) By setting OPENAI_TRACING_ENABLED=False environment variable
 - ✓ b) **By calling `set_tracing_disabled(True)`**
 - c) By removing the `openai.agents.tracing` logger handler
 - d) Tracing cannot be disabled

11. **By default, without any handlers set, which logging levels from the SDK's Python loggers are sent to stdout?**
 - a) All logs (DEBUG, INFO, WARNING, ERROR, CRITICAL)
 - b) Only ERROR logs
 - ✓ c) **Only WARNING and ERROR logs**
 - d) INFO and above

12. **To enable verbose logging from the SDK to stdout, which convenience function can be used?**
 - a) `enable_debug_logging()`
 - b) `set_verbose_logging(True)`
 - ✓ c) **`enable_verbose_stdout_logging()`**
 - d) `configure_logging_level("VERBOSE")`

13. **If you want to customize logging beyond the convenience function, which Python module is recommended to use?**
 - a) `sys`

- ✓b) logging
- c) os
- d) configparser

14. What are the names of the two main Python loggers in the SDK that you can customize?

- a) "openai.sdk" and "openai.logs"
- ✓b) "openai.agents" and "openai.agents.tracing"
- c) "agents.main" and "agents.debug"
- d) "openai.llm" and "openai.tools"

15. To prevent logging of LLM inputs and outputs, which environment variable should be set?

- a) OPENAI_AGENTS_DISABLE_LLM_LOGS=1
- b) OPENAI_AGENTS_NO_LLM_DATA=1
- ✓c) OPENAI_AGENTS_DONT_LOG_MODEL_DATA=1
- d) OPENAI_AGENTS_HIDE_PROMPTS=1

16. To disable logging of tool inputs and outputs, which environment variable should be set?

- a) OPENAI_AGENTS_DISABLE_TOOL_LOGS=1
- b) OPENAI_AGENTS_NO_TOOL_DATA=1
- c) OPENAI_AGENTS_HIDE_TOOL_CALLS=1
- ✓d) OPENAI_AGENTS_DONT_LOG_TOOL_DATA=1

17. If you set a logger's level to logging.INFO, what types of logs will typically be displayed (assuming a handler is present)?

- a) Only INFO logs
- b) DEBUG and INFO logs
- ✓c) INFO, WARNING, ERROR, and CRITICAL logs
- d) Only WARNING and ERROR logs

18. What is the default output stream for logs when no handlers are explicitly added, as per the logging.StreamHandler example?

- ✓a) stdout
- b) stderr
- c) A file named agents.log
- d) A network socket

19. Why might you want to disable logging of sensitive data (like LLM or tool inputs/outputs)?

- a) To improve application performance significantly
- ✓b) To ensure user data privacy and compliance
- c) To reduce the number of API calls made to OpenAI
- d) To make debugging easier

20. The ability to set `base_url` when configuring a custom `AsyncOpenAI` client is particularly useful for what scenario?
- ✓ a) When you want to use a proxy server or an OpenAI-compatible local LLM
 - b) When you want to specifically connect to OpenAI's European data center
 - c) When you are using a different API version for OpenAI
 - d) When you want to enable real-time streaming of responses

Agent Visualization

OpenAI Agents SDK - Agent Visualization Section MCQs

📖 Part: Agent Visualization in the SDK

1. What is the primary purpose of Agent Visualization in the SDK?
 - a) To automatically optimize agent performance.
 - b) To generate executable code from agent definitions.
 - ✓ c) To create a structured graphical representation of agents and their relationships.
 - d) To simulate agent behavior without running them.

2. Which external library is used by the Agents SDK for generating these graphical representations?
 - a) Matplotlib
 - b) NetworkX
 - ✓ c) Graphviz
 - d) Mermaid.js

3. To install the necessary dependencies for agent visualization, which optional dependency group should you install?
 - a) "openai-agents[graph]"
 - ✓ b) "openai-agents[viz]"
 - c) "openai-agents[draw]"
 - d) "openai-agents[render]"

4. Which function is used to generate an agent visualization graph?
 - a) `generate_graph()`
 - b) `visualize_agents()`
 - c) `create_diagram()`
 - ✓ d) `draw_graph()`

5. In the generated graph, how are Agents represented visually?
 - a) Green ellipses.
 - ✓ b) Yellow boxes (rectangles).
 - c) Blue circles.
 - d) Red diamonds.

6. **How are Tools represented in the generated agent visualization graph?**
- a) Yellow boxes.
 - ✓b) **Green ellipses.**
 - c) Blue circles.
 - d) Red diamonds.
7. **What do solid arrows in the generated graph typically represent?**
- a) Tool invocations.
 - b) Data flow between agents.
 - ✓c) **Agent-to-agent handoffs.**
 - d) Conditional execution paths.
8. **What do dotted arrows in the generated graph indicate?**
- a) Agent-to-agent handoffs.
 - ✓b) **Tool invocations.**
 - c) Bidirectional communication.
 - d) Error pathways.
9. **What is the purpose of the `__start__` node in the visualization?**
- a) It represents the end of the execution flow.
 - b) It indicates a sub-agent.
 - ✓c) **It indicates the entry point of the agent system.**
 - d) It marks a debugging breakpoint.
10. **In the example usage, which agent is passed to `draw_graph` to visualize the entire structure?**
- a) `spanish_agent`
 - b) `english_agent`
 - ✓c) **`triage_agent`**
 - d) `get_weather`
11. **By default, how does `draw_graph` display the generated graph?**
- ✓a) **Inline.**
 - b) In a separate browser tab.
 - c) As a text-based ASCII diagram.
 - d) It automatically saves it as a PDF.
12. **To display the generated graph in a separate window, what method should be chained to the `draw_graph` function call?**
- a) `.display()`
 - b) `.show()`
 - ✓c) **`.view()`**
 - d) `.open()`

13. To save the generated graph as a file, which parameter should be used in the `draw_graph` function?
- a) `output_file`
 - b) `save_as`
 - ✓ c) **`filename`**
 - d) `export_path`
14. If you use `draw_graph(triage_agent, filename="agent_graph")`, what will be the name and default format of the generated file?
- a) `agent_graph.txt`
 - b) `agent_graph.pdf`
 - c) `agent_graph.svg`
 - ✓ d) **`agent_graph.png`**
15. The `triage_agent` in the example is configured with `handoffs=[spanish_agent, english_agent]`. How will these relationships be visually represented in the graph?
- a) With green ellipses connecting the agents.
 - b) With dotted arrows from `triage_agent` to `spanish_agent` and `english_agent`.
 - ✓ c) **With solid arrows from `triage_agent` to `spanish_agent` and `english_agent`.**
 - d) These relationships will not be shown in the graph.

Release Process

OpenAI Agents SDK - Release Process Section MCQs

📖 Part: Agents SDK Versioning

1. What versioning format does the OpenAI Agents SDK project follow?
- a) Major.Minor.Patch (X.Y.Z)
 - b) Year.Month.Day (YYYY.MM.DD)
 - ✓ c) **A modified semantic versioning using the form 0.Y.Z**
 - d) Incremental build numbers (0.0.0.X)
-
2. What does the leading 0 in the 0.y.z versioning scheme signify for the SDK?
- a) It's a stable, production-ready release.
 - b) It indicates a pre-alpha stage.
 - ✓ c) **The SDK is still evolving rapidly.**
 - d) It is a long-term support release.
-
3. When are Minor (Y) versions incremented (e.g., from 0.0.x to 0.1.x)?
- a) For any new feature.
 - b) For bug fixes only.

✓c) **For breaking changes to any public interfaces that are not marked as beta.**

d) For changes to private interfaces.

4. **If a user wants to avoid breaking changes in their project, what is the recommended approach for pinning the SDK version?**

a) Pinning to 0.Y.Z for specific Y values.

b) Pinning to 0.Y.X with X as a wildcard.

✓c) **Pinning to 0.0.x versions.**

d) Not pinning at all, and always using the latest.

5. **What kind of changes trigger an increment in Patch (Z) versions?**

a) Only major architectural overhauls.

b) Changes that require users to rewrite significant parts of their code.

✓c) **Non-breaking changes.**

d) Changes that introduce new models.

6. **Which of the following types of changes would result in a Patch (Z) version increment?**

a) Renaming a public class.

b) Changing the required arguments for a public function.

✓c) **A bug fix.**

d) Removing a publicly exposed method.

7. **If the SDK goes from version 0.0.5 to 0.1.0, what type of change has likely occurred?**

a) A non-breaking new feature.

b) A bug fix.

✓c) **A breaking change to a public interface.**

d) An update to a beta feature.

8. **If a new feature is added to the SDK, but it does not break any existing public interfaces, what version component will be incremented?**

a) The leading 0.

b) The Minor (Y) version.

✓c) **The Patch (Z) version.**

d) A new component will be added.

9. **Changes made to "private interfaces" typically result in an increment of which version component?**

a) Minor (Y) version, as they are still changes.

b) The leading 0, to indicate internal shifts.

✓c) **Patch (Z) version, as they are non-breaking for public users.**

d) No version increment, as they are internal.

10. If a beta feature receives an update, but it's not a bug fix, which version component would typically be incremented?
- a) Minor (Y) version.
 - b) The leading 0.
 - ✓c) **Patch (Z) version.**
 - d) A separate beta version number.

Voice agents

OpenAI Agents SDK - Quickstart (Voice) Section

📖 Part 1: Prerequisites & Core Concepts

1. What is the specific command to install the optional voice dependencies for the Agents SDK?
- a) `pip install 'openai-agents[audio]'`
 - b) `pip install 'openai-agents[speech]'`
 - ✓c) **`pip install 'openai-agents[voice]'`**
 - d) `pip install 'openai-agents[sound]'`
-

2. What is the main concept introduced for voice-enabled applications in the SDK?
- a) `AudioProcessor`
 - b) `SpeechWorkflow`
 - ✓c) **`VoicePipeline`**
 - d) `AgentListener`
-

3. How many steps are involved in the `VoicePipeline` process?
- a) 2
 - ✓b) **3**
 - c) 4
 - d) 5
-

4. Which of the following is NOT one of the three steps in the `VoicePipeline`?
- a) Run a speech-to-text model.
 - ✓b) **Run a sentiment analysis model.**
 - c) Run your code (usually an agentic workflow).
 - d) Run a text-to-speech model.
-

5. The `VoicePipeline` takes ☐ **Audio Input** and produces what as its final output?
- a) Text response
 - b) A transcribed file
 - ✓c) ☐ **Audio Output**
 - d) A video stream

🔗 Part 2: Agent Setup for Voice Pipeline

6. In the example, what is the purpose of the `get_weather` function?
- a) To transcribe audio input.
 - b) To convert text to speech.
 - ✓ c) To act as a `function_tool` for agents to get weather information.
 - d) To play audio output.
-
7. What is the `spanish_agent`'s primary instruction regarding language?
- a) It translates English to Spanish.
 - b) It can speak both English and Spanish.
 - ✓ c) It must speak in Spanish.
 - d) It only understands Spanish commands.
-
8. Both the Spanish agent and the Assistant agent in the example use which specific LLM model?
- a) gpt-4o
 - ✓ b) gpt-4o-mini
 - c) gpt-3.5-turbo
 - d) claude-3-5-sonnet-20240620
-
9. What is the Assistant agent instructed to do if the user speaks in Spanish?
- a) Attempt to respond in Spanish itself.
 - b) Ask the user to switch to English.
 - ✓ c) Handoff to the `spanish_agent`.
 - d) Call a special translate tool.
-
10. The Assistant agent is configured with `handoffs=[spanish_agent]`. What does this enable?
- a) The Assistant agent can directly call `spanish_agent`'s tools.
 - b) The Assistant agent can receive handoffs from the `spanish_agent`.
 - ✓ c) The Assistant agent can delegate tasks to the `spanish_agent`.
 - d) It creates a bidirectional communication link between them.
-
11. What function from `agents.extensions.handoff_prompt` is used to enhance the agents' instructions?
- a) `add_handoff_instructions`
 - ✓ b) `prompt_with_handoff_instructions`
 - c) `handoff_aware_prompt`
 - d) `enrich_instructions_for_handoff`
-

🔗 Part 3: Running the Voice Pipeline

12. In the `VoicePipeline` setup, which specific workflow class is used with the agent instance?
- a) `MultiAgentVoiceWorkflow`

- b) VoiceAgentWorkflow
 - ✓c) **SingleAgentVoiceWorkflow**
 - d) DefaultVoiceWorkflow
-

13. For simplicity in the quickstart example, what kind of audio input is generated using `np.zeros`?

- a) A short spoken phrase.
 - ✓b) **3 seconds of silence.**
 - c) A random noise signal.
 - d) A pre-recorded song.
-

14. What is the `samplerate` configured for the `sounddevice.OutputStream` player in the example?

- a) 8000 Hz
 - b) 16000 Hz
 - ✓c) **24000 Hz**
 - d) 44100 Hz
-

15. When streaming the result from `pipeline.run()`, which event type signifies that actual audio data is being received for playback?

- a) `voice_stream_event_text`
 - b) `voice_stream_event_start`
 - ✓c) **`voice_stream_event_audio`**
 - d) `voice_stream_event_end`
-

16. What library is used in the example to handle asynchronous operations like running the pipeline and streaming results?

- a) `threading`
 - b) `concurrent.futures`
 - ✓c) **`asyncio`**
 - d) `multiprocessing`
-

17. If you wanted to test this voice quickstart example with actual microphone input, where would you find a demo for that?

- a) In the `examples/voice/live` directory.
 - b) In the `examples/voice/microphone` directory.
 - ✓c) **In the `examples/voice/static` directory.**
 - d) In the `examples/voice/interactive` directory.
-

18. What is the purpose of `set_tracing_disabled()` being called in the "Put it all together" section?

- a) To improve audio quality.
- b) To reduce memory usage.
- ✓c) **To prevent potential errors if an OpenAI API key for tracing is not set.**
- d) To switch to a different LLM model.

-
19. The buffer for `AudioInput` is created with `dtype=np.int16`. What does this specify?
- a) The number of audio channels.
 - b) The duration of the audio.
 - ✓c) The data type of the audio samples (16-bit integers).
 - d) The sampling rate of the audio.
-
20. In a real-world scenario, what would the buffer for `AudioInput` typically contain instead of silence?
- a) Pre-recorded music.
 - ✓b) Data from a microphone.
 - c) Text from a user.
 - d) Video frames.

Pipelines and Workflows

OpenAI Agents SDK - Pipelines and Workflows Section MCQs

📖 Part 1: VoicePipeline Overview and Configuration

1. What is the primary function of the `VoicePipeline` class?
- a) To manage agent memory.
 - ✓b) To facilitate turning agentic workflows into a voice application.
 - c) To visualize agent interactions.
 - d) To handle LLM model configuration.
-
2. Which of the following is *not* a responsibility of the `VoicePipeline` as described?
- a) Transcribing input audio.
 - b) Detecting when audio ends.
 - ✓c) Providing pre-built agent instructions.
 - d) Turning workflow output back into audio.
-
3. When configuring a `VoicePipeline`, which of the following is a mandatory component to pass?
- a) model provider
 - b) tracing config
 - ✓c) workflow
 - d) prompt language
-

4. The `config` parameter of `VoicePipeline` allows configuration of various aspects. Which of these is an example given for config settings related to tracing?
- a) Agent name.
 - b) Number of handoffs.
 - ✓ c) **Whether audio files are uploaded.**
 - d) The tool definitions.
-

5. Which specific setting within the `VoicePipeline`'s config allows for mapping model names to actual model instances?
- a) `tts_model_settings`
 - b) `stt_model_settings`
 - ✓ c) **`model_provider`**
 - d) `workflow_config`
-

🔗 Part 2: Running the Pipeline and Handling Results

6. Which method is used to initiate a voice pipeline run?
- a) `start()`
 - b) `execute()`
 - ✓ c) **`run()`**
 - d) `process_audio()`
-

7. Which `AudioInput` type is suitable when you have a complete audio transcript and do not need to detect when the speaker is done speaking?
- ✓ a) **`AudioInput`**
 - b) `StreamedAudioInput`
 - c) `RecordedAudioInput`
 - d) `BatchAudioInput`
-

8. For what kind of applications is `AudioInput` particularly useful?
- a) Real-time conversational AI with interruptions.
 - ✓ b) **Push-to-talk apps where the user's speaking completion is clear.**
 - c) Systems requiring continuous microphone listening.
 - d) Applications focused on long-form dictation.
-

9. Which `AudioInput` type should be used if the pipeline needs to automatically detect when a user is done speaking?
- a) `AudioInput`
 - ✓ b) **`StreamedAudioInput`**
 - c) `ContinuousAudioInput`
 - d) `ActivityDetectedAudioInput`
-

10. What is the process called by which the voice pipeline automatically runs the agent workflow at the right time when using `StreamedAudioInput`?
- a) Voice activity detection (VAD).
 - b) Turn-taking.
 - c) Endpointing.
 - ✓d) **Activity detection.**
-

11. The result of a voice pipeline run (`pipeline.run(input)`) is an object of what type?
- a) `VoiceResult`
 - b) `AudioOutput`
 - c) `AgentResponse`
 - ✓d) **`StreamedAudioResult`**
-

12. Which type of `VoiceStreamEvent` contains a chunk of audio that can be played back?
- a) `VoiceStreamEventText`
 - ✓b) **`VoiceStreamEventAudio`**
 - c) `VoiceStreamEventData`
 - d) `VoiceStreamEventStream`
-

13. Which type of `VoiceStreamEvent` would inform you that a new turn in the conversation has started or ended?
- a) `VoiceStreamEventStatus`
 - b) `VoiceStreamEventControl`
 - ✓c) **`VoiceStreamEventLifecycle`**
 - d) `VoiceStreamEventInfo`
-

14. What should you do when handling a `VoiceStreamEventError`?
- a) Ignore it and continue playing audio.
 - b) Restart the entire pipeline.
 - ✓c) **Implement error handling logic (`elif event.type == "voice_stream_event_error"`).**
 - d) Log it only to a file.
-

🔹Part 3: Best Practices – Interruptions

15. Does the Agents SDK currently offer built-in interruption support for `StreamedAudioInput`?
- a) Yes, it's fully supported.
 - b) Only for specific models.
 - ✓c) **No, it does not.**
 - d) It's in alpha stage.
-

16. If you want to handle interruptions within your application, which type of events should you listen for?
- a) `VoiceStreamEventAudio`
 - b) `VoiceStreamEventError`

- ✓c) **VoiceStreamEventLifecycle**
 - d) Any VoiceStreamEvent type
-

17. Which specific `VoiceStreamEventLifecycle` event indicates that a new turn was transcribed and processing by the workflow is beginning?

- a) `turn_completed`
 - ✓b) **`turn_started`**
 - c) `audio_transcribed`
 - d) `workflow_beginning`
-

18. The `turn_ended` `VoiceStreamEventLifecycle` event triggers after what occurs?

- a) The user starts speaking again.
 - b) The workflow completes its processing.
 - ✓c) **All the audio was dispatched for a respective turn.**
 - d) The speech-to-text model finishes.
-

19. A suggested method to handle interruptions by muting/unmuting the microphone involves listening to which two `VoiceStreamEventLifecycle` events?

- a) `turn_started` and `audio_sent`
 - ✓b) **`turn_started` and `turn_ended`**
 - c) `transcription_complete` and `response_received`
 - d) `mic_on` and `mic_off`
-

20. What is the fundamental purpose of an "agentic workflow" that the `VoicePipeline` turns into a voice app?

- a) It's a predefined sequence of steps that always executes the same way.
 - ✓b) **It's an AI-driven process where autonomous AI agents make decisions, take actions, and coordinate tasks.**
 - c) It's a simple script to perform speech-to-text and text-to-speech.
 - d) It's a static set of rules for handling conversational turns.
-

Voice Pipeline Tracing

OpenAI Agents SDK - Voice Pipeline Tracing MCQs

1. How are voice pipelines traced in the Agents SDK?

- ✓b) **They are automatically traced, similar to agents.**
 - ✗a) They must be manually traced by the developer.
 - ✗c) Tracing is disabled by default for voice pipelines.
 - ✗d) Only errors in voice pipelines are traced.
-

2. **What class is specifically used to configure tracing for a VoicePipeline?**

- ✓c) **VoicePipelineConfig**
 - ✗a) TracingConfig
 - ✗b) VoiceTracingSettings
 - ✗d) PipelineConfig
-

3. **By default, what is the status of tracing for voice pipelines?**

- ✓c) **Enabled.**
 - ✗a) Disabled.
 - ✗b) It depends on the global tracing setting.
 - ✗d) It is only enabled in debug mode.
-

4. **Which field in VoicePipelineConfig controls whether tracing is enabled or disabled for the pipeline?**

- ✓c) **tracing_disabled**
 - ✗a) enable_tracing
 - ✗b) disable_tracing
 - ✗d) trace_status
-

5. **The trace_include_sensitive_data field in VoicePipelineConfig specifically controls sensitive data in traces for which part of the system?**

- ✓b) **Data specifically from the voice pipeline (e.g., audio transcripts).**
 - ✗a) Only data from within your agent workflow.
 - ✗c) Data from external tools only.
 - ✗d) All sensitive data across the entire SDK.
-

6. **Which field controls whether actual audio data is included in the voice pipeline traces?**

- ✓d) **trace_include_sensitive_audio_data**
 - ✗a) trace_audio_only
 - ✗b) include_audio_in_traces
 - ✗c) trace_sensitive_audio
-

7. **What is the purpose of the workflow_name field in VoicePipelineConfig?**

- ✓c) **It sets the name of the trace workflow.**
 - ✗a) It sets the name of the agent within the workflow.
 - ✗b) It defines the input prompt for the speech-to-text model.
 - ✗d) It specifies the audio file to be used.
-

8. **The group_id field in VoicePipelineConfig serves what function?**

- ✓b) **To link multiple traces (e.g., from the same conversation).**
- ✗a) To specify the group of agents involved in the pipeline.
- ✗c) To categorize pipelines based on their functionality.
- ✗d) To set permissions for trace access.

-
9. If you want to add arbitrary key-value pairs to a voice pipeline trace, which field in `VoicePipelineConfig` should you use?
- ✓c) `trace_metadata`
 - ✗a) `additional_data`
 - ✗b) `custom_props`
 - ✗d) `extra_info`
-
10. By default, will a voice pipeline trace include audio transcripts, assuming default `VoicePipelineConfig` settings?
- ✓c) Yes, as `trace_include_sensitive_data` defaults to `True`.
 - ✗a) No, sensitive data is disabled by default.
 - ✗b) Yes, but only for errors.
 - ✗d) It depends on an environment variable.

API Reference

Agents

OpenAI Agents SDK - Agents Module Functions MCQs

📖 OpenAI Agents SDK - Tracing & Key Management MCQs

1. What is the primary purpose of `set_default_openai_key()`?
- ✓c) To set the OpenAI API key for LLM requests (and optionally tracing).
 - ✗a) To configure a specific LLM model.
 - ✗b) To enable verbose logging.
 - ✗d) To disable tracing globally.
-
2. If the `OPENAI_API_KEY` environment variable is already set, what happens if `set_default_openai_key()` is called with a different key?
- ✓b) The key provided to `set_default_openai_key()` will be used instead.
 - ✗a) The environment variable takes precedence.
 - ✗c) An error will be raised.
 - ✗d) Both keys will be used interchangeably.
-

3. **By default, when `set_default_openai_key()` is called, is the provided key also used for tracing?**
✓b) **Yes, `use_for_tracing` defaults to `True`.**
✗a) No, tracing requires a separate configuration.
✗c) Only if the `OPENAI_API_KEY` environment variable is not set.
✗d) Only if verbose logging is enabled.
-
4. **Which function allows you to provide a custom `AsyncOpenAI` client instance to the SDK?**
✓c) **`set_default_openai_client()`**
✗a) `configure_openai_client()`
✗b) `set_custom_openai_client()`
✗d) `override_openai_client()`
-
5. **What is a common use case for providing a custom `AsyncOpenAI` client via `set_default_openai_client()`?**
✓c) **To connect to a different `base_url` for an OpenAI-compatible endpoint.**
✗a) To disable all tracing.
✗b) To change the default logging level.
✗d) To automatically retry failed LLM requests.
-
6. **By default, which OpenAI API does the Agents SDK use for LLM requests?**
✓b) **The Responses API.**
✗a) The Chat Completions API.
✗c) The Embeddings API.
✗d) The Moderation API.
-
7. **To switch the default OpenAI API for LLM requests to the Chat Completions API, which function would you use?**
✓c) **`set_default_openai_api("chat_completions")`**
✗a) `use_chat_completions_api()`
✗b) `set_openai_api_type("chat_completions")`
✗d) `configure_api_version("chat_completions")`
-
8. **If you want to use a specific API key only for sending traces to the backend, which function should you call?**
✓c) **`set_tracing_export_api_key()`**
✗a) `set_default_openai_key(key, use_for_tracing=False)`
✗b) `set_tracing_key_only()`
✗d) `configure_trace_destination()`
-
9. **What is the purpose of `set_tracing_disabled(disabled: bool)`?**
✓d) **To globally enable or disable tracing for the entire SDK.**
✗a) To disable specific types of traces.

- ☒ b) To only disable tracing for a single agent.
 - ☒ c) To disable tracing for the OpenAI API calls.
-

10. What argument should be passed to `set_tracing_disabled()` to turn off all tracing?

- ☒ c) **True**
 - ☒ a) "off"
 - ☒ b) 0
 - ☒ d) False
-

11. For advanced customization, if you want to replace the SDK's default trace processing mechanism entirely, which function is used?

- ☒ c) **set_trace_processors()**
 - ☒ a) `add_trace_processor()`
 - ☒ b) `register_trace_handler()`
 - ☒ d) `override_default_tracer()`
-

12. The `set_trace_processors()` function accepts a list of objects that implement which specific interface?

- ☒ c) **TracingProcessor**
 - ☒ a) `TraceHandler`
 - ☒ b) `LogProcessor`
 - ☒ d) `TraceExporter`
-

13. What is the effect of calling `enable_verbose_stdout_logging()`?

- ☒ c) **It enables more detailed SDK logs (e.g., INFO, DEBUG) to be printed to stdout.**
 - ☒ a) It changes the default OpenAI API to Chat Completions.
 - ☒ b) It disables all tracing.
 - ☒ d) It sets a custom OpenAI API key.
-

14. When would `enable_verbose_stdout_logging()` be most useful?

- ☒ c) **During development and debugging.**
 - ☒ a) In production environments to monitor system health.
 - ☒ b) When deploying the application to a server.
 - ☒ d) For generating final user reports.
-

15. If `set_default_openai_key()` is called with `use_for_tracing=False`, and no other tracing API key is set, what will happen to tracing?

- ☒ d) **Tracing will attempt to use the `OPENAI_API_KEY` environment variable if available; otherwise, it will likely not export traces.**
- ☒ a) Tracing will use a dummy key.

- ✗b) Tracing will automatically try to find a key from `OPENAI_API_KEY` environment variable.
- ✗c) Tracing will be entirely disabled.

Agents Module

OpenAI Agents SDK - Agents Module MCQs

✔Part 1: Core Agent Attributes and Behaviors

1. What is the `instructions` attribute of an Agent primarily used for?
 - ✔c) To act as the "system prompt" describing the agent's behavior and goal.
 - ✗a) To define the agent's name.
 - ✗b) To specify the tools the agent can use.
 - ✗d) To list the agents it can handoff to.

2. If an Agent's `instructions` is set to a function, what must that function return?
 - ✔c) A string.
 - ✗a) A Prompt object.
 - ✗b) An Agent instance.
 - ✗d) A list of tool names.

3. The `handoff_description` attribute is used when an agent is included in which specific part of another agent's configuration?
 - ✔b) Its `handoffs` list.
 - ✗a) Its tools list.
 - ✗c) Its `mcp_servers` list.
 - ✗d) Its `input_guardrails`.

4. If an Agent's `model` attribute is not explicitly set, what model will it default to using?
 - ✔b) The default model configured in `openai_provider.DEFAULT_MODEL` (currently "gpt-4o").
 - ✗a) gpt-3.5-turbo
 - ✗c) A local, CPU-only model.
 - ✗d) The last used model in the application.

5. What is the purpose of the `model_settings` attribute?
 - ✔b) To configure model-specific tuning parameters like temperature.
 - ✗a) To configure tracing options.
 - ✗c) To set the API key for the model.
 - ✗d) To define the model's supported output types.

6. Which attribute defines a list of other Agent instances that the current agent can delegate tasks to?

- ✓c) `handoffs`
 - ✗a) `delegates`
 - ✗b) `sub_agents`
 - ✗d) `collaborators`
-

7. What does `input_guardrails` specifically check, and under what condition do they run?

- ✓c) They check inputs before generating a response, running only if the agent is the first in the chain.
 - ✗a) They check the agent's final output, always running.
 - ✗b) They check tool call results, running only if a tool is called.
 - ✗d) They check agent instructions for compliance.
-

8. The `output_type` attribute of an Agent defaults to `str` if not provided. What is a common way to customize it for structured output?

- ✓c) Passing a regular Python type like a dataclass or Pydantic model.
 - ✗a) Passing a list of strings.
 - ✗b) Setting it to `Any`.
 - ✗d) Using a Callable that returns the output.
-

9. What is the default `tool_use_behavior` for an Agent?

- ✓c) `"run_llm_again"`
 - ✗a) `"stop_on_first_tool"`
 - ✗b) A list of specific tool names.
 - ✗d) A `ToolsToFinalOutputFunction`.
-

10. If `tool_use_behavior` is set to `"stop_on_first_tool"`, what happens after the first tool call?

- ✓c) The output of the first tool call is used as the final output, and the LLM does not process it.
 - ✗a) The LLM processes the tool output and generates a new response.
 - ✗b) The agent raises an error indicating immediate termination.
 - ✗d) The agent proceeds to call all other available tools.
-

✓Part 2: Advanced Agent Configuration & Methods

11. What is the primary purpose of MCP servers in the context of an Agent?

- ✓c) To dynamically provide tools to the agent.
 - ✗a) To manage agent authentication.
 - ✗b) To provide a distributed computing environment.
 - ✗d) To store conversation history.
-

12. What crucial action must the user perform when using `mcp_servers` with an Agent?

- ✓c) Call `server.connect()` before passing them and `server.cleanup()` when done.
- ✗a) Ensure they are all running on the same port.

- ☒ b) Disable tracing for those servers.
 - ☒ d) Provide a `handoff_description` for each server.
-

13. What does the `reset_tool_choice: bool = True` attribute aim to prevent?

- ☒ c) The agent from entering an infinite loop of tool usage.
 - ☒ a) The agent from using too many tools in a single turn.
 - ☒ b) The LLM from generating overly long responses.
 - ☒ d) Errors related to tool schema mismatches.
-

14. How does the `clone()` method facilitate agent configuration?

- ☒ c) It creates a copy of the agent with specified arguments changed.
 - ☒ a) It creates an exact duplicate of the agent that runs in parallel.
 - ☒ b) It allows deep copying all mutable attributes to prevent side effects.
 - ☒ d) It saves the agent's current state to disk.
-

15. When transforming an Agent into a Tool using `as_tool()`, what is a key difference in how the new agent receives input compared to a handoff?

- ☒ c) It receives generated input (e.g., a specific query or data).
 - ☒ a) It receives the full conversation history.
 - ☒ b) It receives a summary of the conversation.
 - ☒ d) It receives direct access to the calling agent's internal state.
-

16. What is a key difference in control flow when an Agent is called via `as_tool()` versus being used as a handoff?

- ☒ b) With `as_tool()`, the new agent returns control to the original calling agent.
 - ☒ a) With `as_tool()`, the new agent takes over the conversation.
 - ☒ c) Handoffs are synchronous, while `as_tool()` calls are asynchronous.
 - ☒ d) `as_tool()` calls are always faster than handoffs.
-

17. If an Agent has its `prompt` attribute set to a `DynamicPromptFunction`, what method would you call to get the resolved prompt during a run?

- ☒ b) `get_prompt()`
 - ☒ a) `get_resolved_prompt()`
 - ☒ c) `resolve_dynamic_prompt()`
 - ☒ d) `get_configured_prompt()`
-

18. The `get_mcp_tools()` method allows an agent to access tools from what source?

- ☒ b) From configured Model Context Protocol servers.
 - ☒ a) Only from its directly configured tools list.
 - ☒ c) From a global tool registry.
 - ☒ d) From external API endpoints directly.
-

19. Which attribute allows a class to receive callbacks on various lifecycle events of an agent?
- ✓c) **hooks**
 - ✗a) `event_listeners`
 - ✗b) `callback_handlers`
 - ✗d) `lifecycle_managers`
-
20. If `output_type` is set to `AgentOutputSchema(MyClass, strict_json_schema=False)`, what does `strict_json_schema=False` enable?
- ✓c) **Non-strict schema validation, allowing for more flexible JSON outputs.**
 - ✗a) The agent to output any data type, ignoring `MyClass`.
 - ✗b) A faster schema conversion process.
 - ✗d) Automatic schema generation based on the LLM's response.

OpenAI Agents SDK - Runner MCQs

✓Part 1: Running Workflows (`run`, `run_sync`, `run_streamed`)

1. What is the primary function of the Runner class?
- ✓b) **To execute agent workflows.**
 - ✗a) To configure OpenAI API keys.
 - ✗c) To visualize agent graphs.
 - ✗d) To manage trace processors.
-
2. Which describes the loop mechanism of an agent workflow run?
- ✓a) **Agent invoked → Tool calls → Handoff → Final output.**
 - ✗b) Agent invoked → Final output → Handoff → Tool calls.
 - ✗c) Final output → Agent invoked → Tool calls → Handoff.
 - ✗d) Handoff → Tool calls → Agent invoked → Final output.
-
3. When does the agent workflow loop terminate?
- ✓c) **When the agent generates a final output (of its `agent.output_type`).**
 - ✗a) When `max_turns` is reached.
 - ✗b) When all available tools have been called.
 - ✗d) When a handoff occurs.
-
4. Which exception is raised if the `max_turns` limit is exceeded?
- ✓d) **MaxTurnsExceeded**
 - ✗a) `TurnLimitExceeded`
 - ✗b) `MaxIterationsReached`
 - ✗c) `LoopLimitError`

5. What defines a “turn” in the context of `max_turns`?

- ✓c) One AI invocation (including any tool calls that might occur).
 - ✗a) Every time a tool is called.
 - ✗b) Every time an agent receives new input.
 - ✗d) One successful agent response.
-

6. Which `run` method is asynchronous?

- ✓a) `run`
 - ✗b) `run_sync`
 - ✗c) `run_streamed`
 - ✗d) `run_async_only`
-

7. When should `run_sync()` be avoided?

- ✓c) If there’s already an event loop (e.g., in an async function, Jupyter notebook, or FastAPI).
 - ✗a) When using many tools.
 - ✗b) When tracing is enabled.
 - ✗d) When the agent has handoffs.
-

8. Return type of `run_streamed()` method?

- ✓c) `RunResultStreaming`
 - ✗a) `RunResult`
 - ✗b) `StreamedResult`
 - ✗d) `EventStream`
-

9. Purpose of `previous_response_id` argument in `run` methods?

- ✓d) To skip passing in input from the previous turn when using OpenAI models via the Responses API.
 - ✗a) To link traces to previous runs.
 - ✗b) To enable continuous conversation history management.
 - ✗c) To specify the previous agent in a handoff chain.
-

10. What is unique about input guardrails in a workflow run?

- ✓c) Only the first agent’s input guardrails are run.
 - ✗a) All agents in the chain run their input guardrails.
 - ✗b) Input guardrails are only run after a tool call.
 - ✗d) Input guardrails are ignored if `max_turns` is set.
-

✓Part 2: RunConfig - Global Settings

11. Primary purpose of the `RunConfig` dataclass?

- ✓c) To configure global settings for the entire agent run.
- ✗a) To configure individual agents.

- ☒ b) To define agent instructions.
 - ☒ d) To manage environment variables.
-

12. Effect of setting a model in `RunConfig` on agents?

- ☒ b) It will override the model set on every agent within that run.
 - ☒ a) It is ignored if an agent has its own model specified.
 - ☒ c) It only applies to agents that do not have a model attribute set.
 - ☒ d) It is only used for tracing purposes.
-

13. Role of `model_provider` in `RunConfig`?

- ☒ c) To resolve string model names to actual model implementations.
 - ☒ a) To manage multiple agent instances concurrently.
 - ☒ b) To provide global tracing configurations.
 - ☒ d) To handle handoffs between different agents.
-

14. Interaction between `RunConfig.model_settings` and agent-specific settings?

- ☒ c) Any non-null values in `RunConfig.model_settings` will override agent-specific settings.
 - ☒ a) `RunConfig.model_settings` are ignored if agent-specific settings exist.
 - ☒ b) `RunConfig.model_settings` only apply to agents without specific settings.
 - ☒ d) They are merged, with agent-specific settings taking precedence for conflicts.
-

15. Function of `handoff_input_filter` in `RunConfig`?

- ☒ c) To apply a global filter to inputs sent to all handoffs.
 - ☒ a) To block certain inputs from reaching the initial agent.
 - ☒ b) To filter the output of agents before a handoff.
 - ☒ d) To determine which agent to handoff to.
-

16. Precedence when both `Handoff.input_filter` and `RunConfig.handoff_input_filter` are present?

- ☒ c) `Handoff.input_filter`.
 - ☒ a) `RunConfig.handoff_input_filter`.
 - ☒ b) They are both applied in sequence.
 - ☒ d) An error is raised due to conflict.
-

17. Where do `input_guardrails` from `RunConfig` apply?

- ☒ b) To the initial input of the entire run.
 - ☒ a) To every agent's input in the workflow.
 - ☒ c) To inputs specifically for tools.
 - ☒ d) To inputs after a handoff.
-

18. Where are `output_guardrails` from `RunConfig` applied?

- ☒ c) The final output of the entire run.

- ☒ a) The output of every agent in the workflow.
 - ☒ b) The output of tool calls only.
 - ☒ d) The output of handoffs only.
-

✔Part 3: Tracing Configuration within RunConfig

19. Default status of `tracing_disabled` in RunConfig?

- ✔b) False (tracing is enabled).
 - ☒ a) True (tracing is disabled).
 - ☒ c) It depends on an environment variable.
 - ☒ d) It's ignored if `trace_id` is provided.
-

20. Effect of `trace_include_sensitive_data=False`?

- ✔b) Spans will still be created, but sensitive data (like LLM inputs/outputs, tool data) will not be included.
 - ☒ a) No spans will be created for sensitive events.
 - ☒ c) Tracing will be completely disabled.
 - ☒ d) Only error traces will be recorded.
-

21. Default `workflow_name` if not specified in RunConfig?

- ✔b) 'Agent workflow'
 - ☒ a) The name of the starting agent.
 - ☒ c) None (no workflow name).
 - ☒ d) A random UUID.
-

22. Purpose of providing `trace_id` in RunConfig?

- ✔b) To use a custom, pre-defined trace ID.
 - ☒ a) To automatically generate a new trace ID.
 - ☒ c) To disable automatic trace ID generation.
 - ☒ d) To link runs to a specific user.
-

23. Which attribute links multiple traces in a conversation/process?

- ✔c) `group_id`
 - ☒ a) `trace_id`
 - ☒ b) `workflow_name`
 - ☒ d) `trace_metadata`
-

24. What can be included in `trace_metadata`?

- ✔c) An optional dictionary of additional, arbitrary metadata.
 - ☒ a) Only specific predefined metadata keys.
 - ☒ b) Sensitive data that is otherwise filtered.
 - ☒ d) Only model parameters and versions.
-

25. **Meaning of GuardrailTripwireTriggered exception?**
- ✓c) **A predefined guardrail condition has been violated.**
 - ✗a) The agent has successfully completed its task.
 - ✗b) The maximum number of turns has been reached.
 - ✗d) An external API call has failed.

OpenAI Agents SDK - REPL Module MCQs

✓run_demo_loop Function – MCQs with Answers

1. **What is the primary purpose of the `run_demo_loop` function?**
 - ✓c) **To provide a simple REPL loop for manual testing and debugging of an agent.**
 - ✗a) To deploy agents to a production environment.
 - ✗b) To generate automated test cases for agents.
 - ✗d) To visualize agent workflow diagrams.

2. **What kind of interaction model does `run_demo_loop` facilitate?**
 - ✓c) **A conversational, turn-based loop.**
 - ✗a) One-shot query-response.
 - ✗b) Batch processing of inputs.
 - ✗d) Purely programmatic execution without user input.

3. **How is the conversation state handled across turns within `run_demo_loop`?**
 - ✓b) **It is preserved across turns.**
 - ✗a) It is reset after each turn.
 - ✗c) It is saved to a file after each turn.
 - ✗d) It is handled by the underlying LLM only, not the loop.

4. **Which of the following commands can be used to stop the `run_demo_loop`?**
 - ✓c) **exit or quit**
 - ✗a) stop or end
 - ✗b) close or terminate
 - ✗d) break or halt

5. **What type of argument must be passed as the `agent` parameter to `run_demo_loop`?**
 - ✓d) **An instance of an Agent.**
 - ✗a) A string representing the agent's name.
 - ✗b) An AgentConfig object.
 - ✗c) A function that returns an agent.

6. **By default, how does `run_demo_loop` display the agent's output?**
 - ✓c) **It streams the agent output as it's generated.**

- ☒ a) It waits for the full response and then prints it.
 - ☒ b) It saves the output to a log file.
 - ☒ d) It displays a summary only.
-

7. What is the default value of the `stream` parameter in `run_demo_loop`?

- ☒ b) True
 - ☒ a) False
 - ☒ c) None
 - ☒ d) It's a required parameter with no default.
-

8. If `stream` is set to `False` in `run_demo_loop`, what is the expected behavior for displaying the agent's response?

- ☒ c) The full response will be displayed after it is completely generated.
 - ☒ a) The response will not be displayed.
 - ☒ b) The response will be displayed word by word.
 - ☒ d) Only the final line of the response will be shown.
-

9. From which file path is the `run_demo_loop` function sourced?

- ☒ c) `src/agents/repl.py`
 - ☒ a) `src/agents/main.py`
 - ☒ b) `src/agents/utlis.py`
 - ☒ d) `src/agents/demo.py`
-

10. What is a primary benefit of `run_demo_loop` for developers working with agents?

- ☒ c) It enables quick manual testing and interactive debugging from the command line.
 - ☒ a) It automates unit testing.
 - ☒ b) It provides performance metrics.
 - ☒ d) It's a tool for deploying production agents.
-

OpenAI Agents SDK - Tools MCQs

☒ Part 1: General Tool Concepts and `FunctionTool`

1. What is the `Tool` module-attribute?

- ☒ c) It's a `Union` type defining all possible tool types.
 - ☒ a) It's an abstract base class for all tools.
 - ☒ b) It's a list of all available tool instances.
 - ☒ d) It's a factory for creating new tool instances.
-

2. **What does `FunctionToolResult` encapsulate?**
✓b) **The tool that was run, its output, and the `RunItem` produced.**
✗a) The JSON schema of the tool.
✗c) Only the `FunctionTool` name and description.
✗d) The success or failure status of the tool call.
-

3. **Recommended way to create a `FunctionTool`?**
✓c) **Use the `function_tool` helper.**
✗a) Directly instantiate `FunctionTool()`.
✗b) Use `Tool.create_function()`.
✗d) Load it from a configuration file.
-

4. **Which attribute shows the name to the LLM?**
✓a) **name**
✗b) description
✗c) `params_json_schema`
✗d) `on_invoke_tool`
-

5. **Purpose of the `description` attribute?**
✓c) **To describe the tool's purpose to the LLM.**
✗a) To provide internal documentation.
✗b) To specify the function's arguments.
✗d) To define the tool's return type.
-

6. **How does `on_invoke_tool` receive arguments?**
✓c) **A JSON string.**
✗a) A Python dictionary.
✗b) A list of strings.
✗d) A Pydantic model instance.
-

7. **What must `on_invoke_tool` return?**
✓c) **A string representation (or something `str()` can be called on).**
✗a) An integer.
✗b) A boolean.
✗d) A complex Python object.
-

8. **How can errors be communicated from `on_invoke_tool`?**
✓b) **By raising an `Exception` or returning a string error message.**
✗a) By returning `None`.
✗c) By setting a global error flag.
✗d) By calling `Runner.fail()`.
-

9. **Why is `strict_json_schema=True` recommended?**
- ✓b) **It increases the likelihood of correct JSON input from the LLM.**
 - ✗a) It improves performance.
 - ✗c) It makes the tool easier to debug.
 - ✗d) It enables asynchronous execution.
-

10. **Purpose of `is_enabled` as a `Callable`?**
- ✓c) **To dynamically enable/disable the tool based on the run context and agent.**
 - ✗a) To determine if the tool has been used previously.
 - ✗b) To enable/disable based on user input.
 - ✗d) To indicate if the tool is currently running.
-

✓Part 2: Specialized Hosted Tools

11. **Which tool allows LLM to search a vector store?**
- ✓b) **FileSearchTool**
 - ✗a) WebSearchTool
 - ✗c) CodeInterpreterTool
 - ✗d) ImageGenerationTool
-

12. **FileSearchTool is supported with which API?**
- ✓b) **OpenAI models using the Responses API.**
 - ✗a) All OpenAI models via standard API.
 - ✗c) Locally hosted open-source models.
 - ✗d) Any model with tool-calling capabilities.
-

13. **Which WebSearchTool attribute customizes geography?**
- ✓b) **user_location**
 - ✗a) search_context_size
 - ✗c) max_num_results
 - ✗d) filters
-

14. **What describes the environment in `ComputerTool`?**
- ✓b) **computer**
 - ✗a) shell_executor
 - ✗c) system_interface
 - ✗d) tool_config
-

15. **Role of `on_approval_request` in `HostedMCPTool`?**
- ✓b) **To provide a function for programmatic approval or rejection.**
 - ✗a) To log all tool calls.
 - ✗c) To automatically approve all MCP calls.
 - ✗d) To notify user that tool call has finished.

16. Return type of `MCPToolApprovalFunction`?

- ☒ c) `MCPToolApprovalFunctionResult`
 - ☐ a) `bool`
 - ☐ b) `str`
 - ☐ d) `None`
-

17. Tool that executes code in a sandbox?

- ☒ d) `CodeInterpreterTool`
 - ☐ a) `LocalShellTool`
 - ☐ b) `ComputerTool`
 - ☐ c) `HostedMCPTool`
-

18. Typical function for executor in `LocalShellTool`?

- ☒ b) A function that executes a command on a shell.
 - ☐ a) A function that generates a command string.
 - ☐ c) A function that parses shell output.
 - ☐ d) A function that validates shell commands.
-

☒ Part 3: `function_tool` Helper and Error Handling

19. What does `function_tool` parse automatically?

- ☒ b) The JSON schema for the tool's parameters.
 - ☐ a) The `ToolContext` type.
 - ☐ c) The agent's `output_type`.
 - ☐ d) The `RunConfig` settings.
-

20. How is the tool's description populated by default?

- ☒ c) From the function's docstring.
 - ☐ a) From a YAML file.
 - ☐ b) From a default generic string.
 - ☐ d) From the function's name.
-

21. If `failure_error_function=None` and call fails?

- ☒ c) An `Exception` will be raised, causing the run to fail.
 - ☐ a) A generic error message is sent to the LLM.
 - ☐ b) The tool will retry automatically.
 - ☐ d) The error will be silently ignored.
-

22. Purpose of `use_docstring_info` in `function_tool`?

- ☒ c) To control whether the docstring is used for tool/arg descriptions.
- ☐ a) To enable syntax highlighting.

- ✗b) To enforce docstring style.
 - ✗d) To validate length.
-

23. Requirement for `RunContextWrapper` in wrapped function?

- ✓c) It must match the context type (`TContext`) of the agent.
 - ✗a) It must be `Any`.
 - ✗b) It must be `None`.
 - ✗d) It can be any dataclass.
-

24. Which is NOT a valid Tool type?

- ✓d) `DatabaseQueryTool`
 - ✗a) `ImageGenerationTool`
 - ✗b) `WebSearchTool`
 - ✗c) `ComputerTool`
-

25. When is `reason` used in `MCPToolApprovalFunctionResult`?

- ✓b) When the tool call is rejected.
 - ✗a) When the tool call is approved.
 - ✗c) After the call has completed.
 - ✗d) As a general comment.
-

OpenAI Agents SDK - Results MCQs

✓Part 1: `RunResultBase` (Common Attributes and Methods)

1. What does it mean that `RunResultBase` is an ABC?

- ✓c) It cannot be directly instantiated and defines an interface for its subclasses.
 - ✗a) It can be directly instantiated.
 - ✗b) It provides concrete implementations.
 - ✗d) It is a utility class.
-

2. Which attribute contains all new messages, tool calls, and outputs?

- ✓b) `new_items`
 - ✗a) `input`
 - ✗c) `raw_responses`
 - ✗d) `final_output`
-

3. What does `raw_responses` provide?

- ✓b) The raw, unfiltered LLM responses.

- ☒ a) Processed and filtered LLM outputs.
 - ☒ c) Only the final response.
 - ☒ d) A summary of LLM usage.
-

4. The `final_output` attribute represents:

- ☒ c) The output of the last agent in the workflow.
 - ☒ a) The initial input.
 - ☒ b) Output of every agent.
 - ☒ d) A list of all intermediates.
-

5. What is stored in `input_guardrail_results`?

- ☒ c) Guardrail results for the initial input messages.
 - ☒ a) Guardrails for tool inputs.
 - ☒ b) For handoff inputs.
 - ☒ d) For the final output.
-

6. Meaning of `last_agent` being an `abstractmethod`?

- ☒ c) Subclasses must implement this property.
 - ☒ a) Provides a default agent.
 - ☒ b) It's optional.
 - ☒ d) Refers to the starting agent.
-

7. Purpose of `last_response_id` property?

- ☒ c) Get the ID of the last model response.
 - ☒ a) ID of first model response.
 - ☒ b) Unique ID for the run.
 - ☒ d) User who started the run.
-

8. Type checking in `final_output_as(cls)` by default:

- ☒ b) Static typechecking only, no runtime check by default.
 - ☒ a) Strict runtime checking.
 - ☒ c) Best-effort conversion.
 - ☒ d) Deep object validation.
-

9. How to enforce runtime `TypeError` on mismatch in `final_output_as`?

- ☒ b) `raise_if_incorrect_type`
 - ☒ a) `strict_mode`
 - ☒ c) `enforce_type`
 - ☒ d) `validate_output`
-

10. What does `to_input_list()` do?

- ☒ c) Creates a new input list by merging the original input with all new items.

- ☒ a) Converts final output to input.
 - ☒ b) Filters based on guardrails.
 - ☒ d) Saves input to a file.
-

✔ Part 2: `RunResult` and `RunResultStreaming`

11. Class returned by `Runner.run()` (non-streamed)?

- ✔ a) **`RunResult`**
 - ☒ b) `RunResultBase`
 - ☒ c) `RunResultStreaming`
 - ☒ d) `StreamEvent`
-

12. How does `RunResult` differ from `RunResultBase`?

- ✔ c) **Provides a concrete implementation for `last_agent`.**
 - ☒ a) Adds streaming.
 - ☒ b) More guardrail data.
 - ☒ d) Different I/O attributes.
-

13. Main purpose of `RunResultStreaming`?

- ✔ c) **To stream semantic events during a run.**
 - ☒ a) Store final results.
 - ☒ b) Offline analysis.
 - ☒ d) Manage concurrent runs.
-

14. Initial value of `final_output` in `RunResultStreaming`?

- ✔ b) **None**
 - ☒ a) ""
 - ☒ c) First generated output
 - ☒ d) Placeholder object
-

15. When is `is_complete` True in streaming?

- ✔ c) **When the agent finishes and final output is produced.**
 - ☒ a) After first turn.
 - ☒ b) When `stream_events` is called.
 - ☒ d) When `cancel` is called.
-

16. How to stop `RunResultStreaming` run?

- ✔ c) **Call the `cancel()` method.**
 - ☒ a) `stop_stream()`
 - ☒ b) `SystemExit`
 - ☒ d) Let `max_turns` exceed
-

17. What is yielded by `stream_events()`?
- ✓c) `StreamEvent` objects with type field and data.
 - ✗a) Raw LLM chunks
 - ✗b) Python dicts
 - ✗d) Only error messages
-

18. Exception if `max_turns` is exceeded in stream:
- ✓c) `MaxTurnsExceeded`
 - ✗a) `TurnLimitExceededError`
 - ✗b) `MaxIterationsReached`
 - ✗d) `StreamingLimitError`
-

19. Exception raised if a guardrail is tripped?
- ✓c) `GuardrailTripwireTriggered`
 - ✗a) `GuardrailViolation`
 - ✗b) `SecurityAlert`
 - ✗d) `PolicyViolationError`
-

20. When does `last_agent` hold true final agent (streaming)?
- ✓c) After `is_complete` is `True`.
 - ✗a) Immediately after `run_streamed()`.
 - ✗b) After first `StreamEvent`.
 - ✗d) After `cancel()` is called.
-

OpenAI Agents SDK - Streaming Events MCQs

✓Part: `StreamEvent` & Its Variants

1. What is `StreamEvent` defined as?
- ✓b) A `TypeAlias` that can be one of `RawResponsesStreamEvent`, `RunItemStreamEvent`, or `AgentUpdatedStreamEvent`.
 - ✗a) Concrete class
 - ✗c) Abstract base class
 - ✗d) Generator function
-
2. Which type of `StreamEvent` carries raw LLM deltas?
- ✓b) `RawResponsesStreamEvent`
 - ✗a) `RunItemStreamEvent`
 - ✗c) `AgentUpdatedStreamEvent`
 - ✗d) `SemanticStreamEvent`

3. What is the type of `RawResponsesStreamEvent`?

- ✓c) `'raw_response_event'`
 - ✗a) `raw_event`
 - ✗b) `llm_response`
 - ✗d) `model_data`
-

4. What does a `RunItemStreamEvent` wrap?

- ✓c) Semantic events that wrap a `RunItem`.
 - ✗a) Raw LLM tokens
 - ✗b) Debugging logs
 - ✗d) System updates
-

5. Which attribute classifies a `RunItemStreamEvent` action?

- ✓b) `name`
 - ✗a) `type`
 - ✗c) `data`
 - ✗d) `status`
-

6. Which is NOT a valid `RunItemStreamEvent.name`?

- ✓d) `"run_completed"`
 - ✗a) `message_output_created`
 - ✗b) `tool_called`
 - ✗c) `handoff_occured` ✓ (Valid spelling is with double "r", i.e., `"handoff_occurred"`)
-

7. What does `item` in `RunItemStreamEvent` hold?

- ✓c) The actual `RunItem` object that was created.
 - ✗a) Raw LLM output
 - ✗b) Previous event
 - ✗d) String description
-

8. Purpose of `AgentUpdatedStreamEvent`?

- ✓c) Indicates a new agent is running (after a handoff).
 - ✗a) Agent state changed
 - ✗b) New tool added
 - ✗d) Agent completion
-

9. What's in `new_agent` of `AgentUpdatedStreamEvent`?

- ✓c) The new active `Agent` instance.
- ✗a) Previous agent name
- ✗b) Agent identifier
- ✗d) Config dict

10. Which attribute to check first in a stream event?

- ✓b) type
 - ✗a) name
 - ✗c) item
 - ✗d) data
-

11. If `name = "tool_output"` in `RunItemStreamEvent`, what's in `item`?

- ✓c) Output from a tool call.
 - ✗a) Message object
 - ✗b) Handoff request
 - ✗d) Tool definition
-

12. Primary use of `RawResponsesStreamEvent`?

- ✓b) Access raw LLM tokens/chunks.
 - ✗a) Debug logic
 - ✗c) Track handoffs
 - ✗d) Guardrail monitoring
-

13. What event when an agent uses a tool (e.g. web search)?

- ✓c) `RunItemStreamEvent` (`name="tool_called"`)
 - ✗a) `AgentUpdatedStreamEvent`
 - ✗b) `RawResponsesStreamEvent`
 - ✗d) `ToolExecutionEvent` ✗(Not a real type)
-

14. What does `AgentUpdatedStreamEvent` imply?

- ✓c) Likely a handoff occurred.
 - ✗a) Run completed
 - ✗b) Agent error
 - ✗d) New input to same agent
-

15. Difference: `data` (`RawResponsesStreamEvent`) vs. `item` (`RunItemStreamEvent`)?

- ✓c) `data` = raw LLM stream, `item` = structured semantic action.
 - ✗a) `data` = parsed, `item` = raw
 - ✗b) `data` = string, `item` = dict
 - ✗d) Interchangeable ✗(Not true)
-

✓Part 1: Handoff Concepts and Data Structures

1. **What is the primary purpose of a "handoff" in the Agents SDK?**
 - ✓b) **To delegate a task or conversation to another agent.**
 - ✗a) To save the agent's state to a database.
 - ✗c) To restart the current agent.
 - ✗d) To synchronize multiple agent instances.

 2. **Expected input/output for a HandoffInputFilter function:**
 - ✓c) **Takes HandoffInputData, returns HandoffInputData.**
 - ✗a) Takes str, returns str.
 - ✗b) Takes list[RunItem], returns list[RunItem].
 - ✗d) Takes Agent, returns Agent.

 3. **Attribute containing pre-run conversation history:**
 - ✓b) **input_history**
 - ✗a) new_items
 - ✗c) pre_handoff_items
 - ✗d) current_turn_items

 4. **What does new_items include in HandoffInputData?**
 - ✓c) **Items generated during the current agent turn, including handoff-related items.**
 - ✗a) Only the final output of the previous agent.
 - ✗b) The entire conversation history.
 - ✗d) Only raw LLM responses.

 5. **What RunItem typically triggers a handoff?**
 - ✓c) **The documentation implies a specific RunItem related to the handoff being invoked.**
 - ✗a) message_output_created
 - ✗b) tool_output
 - ✗d) reasoning_item_created
-

✓Part 2: Handoff Dataclass Attributes and Behavior

6. **Purpose of tool_name and tool_description in Handoff:**
 - ✓c) **To allow the LLM to recognize and decide when to "call" this handoff as a tool.**
 - ✗a) For internal logging
 - ✗b) To identify the agent being handed off to
 - ✗d) To generate unique IDs

7. **What must on_invoke_handoff return?**
 - ✓c) **An instance of Agent[TContext]**

- ☒ a) A boolean
- ☒ b) A string
- ☒ d) RunContextWrapper

8. **Precedence between input_filter in RunConfig and Handoff:**

☒ c) **Handoff.input_filter**

- ☒ a) RunConfig
- ☒ b) Both sequentially
- ☒ d) System error

9. **Effect of modifying input history via input_filter during streaming:**

☒ c) **Nothing will be streamed as a result of this function; items already streamed.**

- ☒ a) All events re-streamed
- ☒ b) Next turn only
- ☒ d) InputFilteredEvent sent

10. **Why strict_json_schema=True is recommended for Handoff:**

☒ c) **Increases likelihood of correct JSON from LLM**

- ☒ a) Makes handoffs faster
- ☒ b) Allows more flexibility
- ☒ d) Disables input validation

11. **What does input_json_schema define?**

☒ c) **The JSON schema for inputs to the next agent during handoff.**

- ☒ a) Internal state
- ☒ b) Final output
- ☒ d) RunContextWrapper

☒ **Part 3: handoff Helper Function**

12. **Primary role of handoff helper:**

☒ c) **To simplify the creation of Handoff objects**

- ☒ a) Executes handoff immediately
- ☒ b) Registers with a registry
- ☒ d) Manages approval

13. **What can the agent parameter also be?**

☒ c) **A function that returns an agent**

- ☒ a) Agent name (str)
- ☒ b) Config dictionary
- ☒ d) List of agents

14. **Effect of `on_handoff` on schema and description:**
- ✓c) **Signature and docstring can auto-generate them**
 - ✗a) Must set manually
 - ✗b) Overrides with defaults
 - ✗d) Prevents generation
15. **Purpose of `input_type` parameter:**
- ✓c) **Defines the expected input type for `on_handoff`**
 - ✗a) Return type
 - ✗b) `HandoffInputData`
 - ✗d) Agent type
16. **True statement about `on_handoff` and `input_type`:**
- ✓c) **`input_type` is only relevant if `on_handoff` takes input**
 - ✗a) `on_handoff` always required
 - ✗b) `input_type` always required
 - ✗d) `on_handoff` only works with `input_type=None`
17. **What does `tool_name_override` change?**
- ✓c) **The name of the "tool" used by LLM for handoff**
 - ✗a) Agent name
 - ✗b) Default name from helper
 - ✗d) Function name
18. **Best use-case for handoffs:**
- ✓c) **Multi-agent system solving complex problems**
 - ✗a) Single simple task
 - ✗b) Repetitive calculations
 - ✗d) Single API interaction
19. **What does `handoff_input_data.all_items` contain?**
- ✓c) **`input_history` + `pre_handoff_items` + `new_items` (maybe filtered)**
 - ✗a) `input_history` only
 - ✗b) `new_items` only
 - ✗d) Only raw LLM responses
20. **What does it mean that `handoff` can be used as a decorator?**
- ✓c) **You can use `@handoff(...)` on a function to create the handoff**
 - ✗a) Must call with `()`
 - ✗b) Cannot take parameters
 - ✗d) Returns a boolean

OpenAI Agents SDK - Lifecycle MCQs

✔Part 1: RunHooks (Global Run Lifecycle)

- Primary purpose of RunHooks:**
 - ✔b) To receive callbacks on various lifecycle events for an entire agent run.
 - ✗a) Manage config
 - ✗c) Define state
 - ✗d) Handle tool errors
- When is `on_agent_start` called?**
 - ✔c) Before any agent is invoked, and each time the current active agent changes.
 - ✗a) Only once at start
 - ✗b) After agent task
 - ✗d) When a tool is called
- `on_agent_end` signals what?**
 - ✔b) Agent produced a final output
 - ✗a) Agent started
 - ✗c) About to be garbage collected
 - ✗d) Agent error
- `on_handoff` agent parameters:**
 - ✔c) `from_agent`, `to_agent`
 - ✗a) `current_agent`, `next_agent`
 - ✗b) `initiator_agent`, `receiver_agent`
 - ✗d) `primary_agent`, `secondary_agent`
- When is `on_tool_start` called?**
 - ✔b) Before any tool is invoked by any agent in the run.
 - ✗a) After tool result
 - ✗c) `FunctionTool` only
 - ✗d) Agent requests info
- What does `on_tool_end` provide that `on_tool_start` doesn't?**
 - ✔c) The result (output) of the tool's execution.
 - ✗a) Agent
 - ✗b) Description
 - ✗d) Context

7. **To register global lifecycle callbacks, subclass:**
✓b) **RunHooks**
✗a) AgentHooks
✗c) GlobalHooks
✗d) WorkflowHooks
8. **Implication of async RunHooks methods:**
✓c) **They should be defined using `async def` and can use `await`.**
✗a) Must return None
✗b) Execute synchronously
✗d) Run in separate process
-

✓Part 2: AgentHooks (Specific Agent Lifecycle)

9. **Main difference between AgentHooks vs RunHooks:**
✓c) **AgentHooks = specific agent callbacks, RunHooks = global run callbacks.**
✗a) Debug vs production
✗b) Sync vs async
✗d) Single-turn vs multi-turn
10. **How to register AgentHooks:**
✓c) **Assign to `agent.hooks` attribute.**
✗a) Via `Runner.run()`
✗b) Global config
✗d) Override register
11. **When is AgentHooks.on_start called?**
✓c) **Before the agent becomes active each time.**
✗a) Start of run
✗b) After turn
✗d) On tool call
12. **What does `source` refer to in AgentHooks.on_handoff?**
✓c) **The agent handing off to this agent.**
✗a) Target agent
✗b) Handoff output
✗d) Triggering tool
13. **Correct on_handoff sequence when Agent A → Agent B:**
✓a) **RunHooks.on_handoff(from_agent=A, to_agent=B) AND AgentHooks(B).on_handoff(agent=B, source=A).**
✗b), c), d) are incorrect variations.

14. If a tool is invoked by an agent with AgentHooks:

- ✓c) Both RunHooks.on_tool_start and AgentHooks.on_tool_start are called.
- ✗a), b), d) miss one of them.

15. When is AgentHooks.on_end called?

- ✓b) When this agent produces a final output.
 - ✗a) Reasoning step
 - ✗c) Entire run ends
 - ✗d) Agent deactivation
-

✓Part 3: General Lifecycle Concepts

16. Key benefit of lifecycle hooks:

- ✓c) They enable custom logging, monitoring, and behavior injection.
- ✗a) Auto-optimize performance
- ✗b) Replace code
- ✗d) Built-in UI tools

17. What does TContext represent?

- ✓c) The context type for the agent run.
- ✗a) Agent type
- ✗b) Tool type
- ✗d) StreamEvent type

18. To log every tool call in a multi-agent system:

- ✓b) RunHooks.on_tool_start
- ✗a) Per-agent is harder
- ✗c), d) are unrelated

19. Custom setup when a specific agent becomes active:

- ✓b) AgentHooks.on_start (on that agent)
- ✗a) Global start
- ✗c) Handoff tracking
- ✗d) Tool calls

20. What happens if you use lifecycle methods without async?

- ✓c) Runtime errors or blocking behavior may occur.
 - ✗a) Auto-wrapped
 - ✗b) Run in thread
 - ✗d) Silently ignored
-

OpenAI Agents SDK - Items MCQs

✔✘Part 1: Type Aliases and RunItem Overview

1. What does **TResponse** serve as a type alias for?
a) **RunResult** ✘
b) **AgentResponse** ✘
✔c) **Response from the OpenAI SDK**
d) **StreamEvent** ✘
2. Which **TypeAlias** represents any semantic item generated by an agent?
a) **ToolCallItemTypes** ✘
b) **TResponseOutputItem** ✘
✔c) **RunItem**
d) **ModelResponse** ✘
3. **ToolCallItemTypes** specifically represents:
a) Any type of agent message ✘
b) Any type of handoff event ✘
✔c) **Any type of raw tool call item (function, computer, file search, MCP, image, shell)**
d) Any type of agent output ✘
4. Which of the following is **NOT** a type of **RunItem**?
a) **ReasoningItem** ✘
b) **HandoffOutputItem** ✘
c) **MCPApprovalRequestItem** ✘
✔d) **InputMessageItem**
5. The purpose of **TResponseInputItem** is to:
a) Represent items returned by the model ✘
✔b) **Represent input parameters for the model**
c) Represent items from **ItemHelpers** ✘
d) Represent streaming output ✘

✔✘Part 2: RunItemBase and Subclasses

6. In **RunItemBase[T]**, what does **T** represent?
a) The agent that generated the item ✘
b) The tool type ✘

- ✓c) **The type of the raw_item**
- d) RunContextWrapper ✗

7. Which attribute is common to all subclasses of RunItemBase?

- a) output ✗
- ✓b) **agent**
- c) source_agent ✗
- d) name ✗

8. raw_item attribute in RunItemBase is always:

- a) StreamEvent or ModelResponse ✗
- b) ToolCallItemTypes or HandoffCallItem ✗
- ✓c) **ResponseOutputItem or ResponseInputItemParam**
- d) Agent or Tool ✗

9. RunItem subclass for messages directly from LLM:

- a) ReasoningItem ✗
- b) ToolCallOutputItem ✗
- ✓c) **MessageOutputItem**
- d) HandoffCallItem ✗

10. Unique attributes of HandoffOutputItem:

- a) output and reasoning ✗
- b) tool_name and tool_description ✗
- ✓c) **source_agent and target_agent**
- d) input_history and new_items ✗

11. raw_item type within a HandoffCallItem:

- a) McpCall ✗
- b) TResponseInputItem ✗
- ✓c) **ResponseFunctionToolCall**
- d) ResponseComputerToolCall ✗

12. ToolCallItem raw_item is ToolCallItemTypes, meaning:

- a) Only FunctionTool calls ✗
- ✓b) **Can be function, shell, computer, etc.**
- c) Always tool output ✗
- d) No raw data ✗

13. RunItem subclass with processed tool output separate from raw_item:

- a) HandoffOutputItem ✗
- b) **MessageOutputItem** ✗

- c) ReasoningItem ✗
- ✓d) ToolCallOutputItem

14. What does ReasoningItem represent?

- a) A completed task ✗
- b) An error message ✗
- ✓c) Agent's internal thought step
- d) Request for more input ✗

15. MCPApprovalRequestItem is for:

- a) General human input ✗
- b) Approval of regular tools ✗
- ✓c) MCP-specific tool approval
- d) Reply to previous request ✗

✓✗Part 3: ModelResponse and ItemHelpers

16. ModelResponse.output contains:

- a) Final string output ✗
- b) One token ✗
- ✓c) List of TResponseOutputItems
- d) Reasoning steps ✗

17. Purpose of response_id:

- a) Identifies agent ✗
- b) Tracks turn number ✗
- ✓c) ID for reusing in future model calls
- d) Model name ✗

18. to_input_items() does what?

- a) Converts to string ✗
- b) Filters tool calls ✗
- ✓c) Converts output into valid input items
- d) Saves response ✗

19. Method to extract last content or refusal:

- ✓a) extract_last_content
- b) extract_last_text ✗
- c) text_message_output ✗
- d) get_message_text ✗

20. **Method to get only text (ignore refusal):**
a) `extract_last_content` ✗
✓ b) `extract_last_text`
c) `text_message_outputs` ✗
d) `tool_call_output_item` ✗
21. **Return type of `text_message_output()`:**
a) `list[str]` ✗
b) `MessageOutputItem` ✗
✓ c) `str`
d) `None` ✗
22. **`tool_call_output_item()` creates:**
a) Executes a tool ✗
b) Parses tool call ✗
✓ c) **Creates raw tool call output with string result**
d) Gets tool name ✗
23. **Normalize inputs into `TResponseInputItems` list:**
a) `extract_last_text` ✗
b) `text_message_outputs` ✗
✓ c) **`input_to_new_input_list`**
d) `to_input_items` ✗
24. **Primary purpose of `Items` module:**
a) Database handling ✗
b) UI rendering ✗
✓ c) **Standardize agent data and events**
d) Define agent logic ✗
25. **To inspect messages/tool calls in `RunResult`:**
a) `raw_responses` ✗
b) `final_output` ✗
✓ c) **`new_items`**
d) `input_guardrail_results` ✗

OpenAI Agents SDK - Run Context MCQs

✓✗RunContextWrapper Quiz Answers

1. **What is the primary purpose of `RunContextWrapper`?**
a) To manage LLM model configurations ✗
b) To store agent conversation history for the LLM ✗

- ✓c) To wrap a custom context object and track run usage, accessible by user-implemented code
d) To define the agent's response format ✗

2. Which type parameter does RunContextWrapper use to define the type of the custom context object?
a) TAgent ✗
b) TResult ✗
✓c) TContext
d) TInput ✗

3. Is the context object within RunContextWrapper passed to the LLM as part of the prompt?
a) Yes, always ✗
b) Yes, but only if it's a string ✗
✓c) No, contexts are explicitly not passed to the LLM
d) Only in streaming mode ✗

4. Which attribute of RunContextWrapper is used to track resource consumption like token counts?
a) context ✗
b) response_id ✗
✓c) usage
d) output ✗

5. When using RunResultStreaming, what is true about the usage attribute of RunContextWrapper?
a) It is always perfectly up-to-date ✗
b) It is reset to zero after each stream event ✗
✓c) It will be stale until the last chunk of the stream is processed
d) It only tracks input tokens, not output tokens ✗

6. How is the usage attribute typically initialized if not explicitly provided during the run setup?
a) It starts as None ✗
✓b) Using a default_factory to create a new Usage object
c) It is a static attribute shared across all runs ✗
d) It is lazily loaded only when accessed ✗

7. To what parts of your agent application can the context object within RunContextWrapper provide data or dependencies?
a) Only to the Runner itself ✗
b) Only to the LLM ✗
✓c) Tool functions, callbacks, hooks, and other user-implemented code
d) Only for logging purposes ✗

8. If you want to pass a database connection object to all your agent's tools, where would you typically store this connection?

- a) Directly in the agent's configuration ✗
 - b) As a global variable ✗
 - ✓c) **Within the custom context object passed to Runner.run(), accessible via RunContextWrapper.context**
 - d) In environment variables ✗
9. **What happens if you don't provide a custom context object to Runner.run()?**
- a) The run will fail ✗
 - ✓b) **The context attribute of RunContextWrapper will likely be None (or its default if TContext allows it)**
 - c) A default context object will be automatically generated with random data ✗
 - d) usage tracking will be disabled ✗
10. **The RunContextWrapper helps to adhere to which software design principle by providing a dedicated object for dependencies?**
- a) High Coupling ✗
 - b) Data Duplication ✗
 - ✓c) **Dependency Injection**
 - d) Global State Management ✗
-

OpenAI Agents SDK - Usage MCQs

✓✗Usage Dataclass Quiz Answers

1. **What is the primary purpose of the Usage dataclass?**
- a) To manage agent internal state ✗
 - b) To store conversation history ✗
 - ✓c) **To track resource consumption (like LLM requests and tokens) during an agent run**
 - d) To configure agent behavior ✗
2. **The requests attribute in Usage counts:**
- a) The number of user inputs ✗
 - b) The number of handoffs between agents ✗
 - ✓c) **The total number of API requests made to the LLM**
 - d) The number of tools executed ✗
3. **Which attribute in Usage specifically tracks the total tokens sent to the LLM?**
- a) output_tokens ✗
 - ✓b) **input_tokens**
 - c) total_tokens ✗
 - d) requests ✗

4. **The `input_tokens_details` attribute provides:**
- a) A summary of output tokens ✗
 - ✓b) **More granular information about input tokens, such as cached tokens**
 - c) Details about the LLM model used ✗
 - d) The timestamp of the last input ✗
5. **What kind of information might be found within `output_tokens_details`?**
- a) The latency of the LLM response ✗
 - b) The number of API errors ✗
 - ✓c) **Details about output tokens, potentially including `reasoning_tokens`**
 - d) The number of tools called by the agent ✗
6. **The `total_tokens` attribute is calculated as the sum of:**
- a) requests and `input_tokens` ✗
 - ✓b) **`input_tokens` and `output_tokens`**
 - c) requests and `output_tokens` ✗
 - d) `cached_tokens` and `reasoning_tokens` ✗
7. **By default, what is the initial value for requests, `input_tokens`, `output_tokens`, and `total_tokens` when a `Usage` object is created?**
- a) None ✗
 - b) 1 ✗
 - ✓c) **0**
 - d) `field(default_factory=...)` ✗
8. **Where would you typically find an instance of the `Usage` dataclass during an agent run?**
- a) Directly on the `Agent` object ✗
 - b) As a global variable in the SDK ✗
 - ✓c) **As an attribute of `RunContextWrapper`**
 - d) Within the `StreamEvent` object ✗
9. **Why might the usage metrics be "stale" when dealing with streamed responses?**
- a) Because streaming is only for output, not input ✗
 - b) Because the agent might be paused ✗
 - ✓c) **Because the full usage can only be calculated once all chunks of the stream have been processed**
 - d) Because the LLM does not provide usage details during streaming ✗
10. **Which attribute would you check to understand how many tokens were used for an agent's internal thought processes if the model distinguishes them?**
- a) `input_tokens` ✗
 - b) `total_tokens` ✗
 - c) `requests` ✗
 - ✓d) **`output_tokens_details.reasoning_tokens` (assuming `OutputTokensDetails` contains this field)**

OpenAI Agents SDK - Exceptions MCQs

✔️ Agents SDK - Exception Handling Quiz Answers

1. **What is the base class for all exceptions within the OpenAI Agents SDK?**
 - a) BaseException ❌
 - b) PythonException ❌
 - ✔️ c) **AgentsException**
 - d) SDKError ❌

2. **What is the primary role of RunErrorDetails?**
 - a) To define new types of exceptions ❌
 - b) To handle exceptions automatically ❌
 - ✔️ c) **To collect diagnostic data when an exception occurs during an agent run**
 - d) To log successful agent runs ❌

3. **When is a MaxTurnsExceeded exception raised?**
 - a) When the agent takes too long to respond ❌
 - ✔️ b) **When the agent run exceeds a predefined maximum number of turns (iterations)**
 - c) When the LLM generates a very long output ❌
 - d) When a tool call takes too many turns ❌

4. **Which exception indicates that the underlying LLM has behaved in an unexpected or invalid way, such as calling a non-existent tool?**
 - a) UserError ❌
 - b) AgentsException ❌
 - ✔️ c) **ModelBehaviorError**
 - d) InternalError ❌

5. **If the LLM provides malformed JSON when attempting to call a tool, which exception is most likely to be raised?**
 - a) UserError ❌
 - ✔️ b) **ModelBehaviorError**
 - c) InputGuardrailTripwireTriggered ❌
 - d) MaxTurnsExceeded ❌

6. **A UserError is raised when:**
 - a) The agent cannot understand the user's input ❌
 - b) The LLM generates an incorrect response ❌
 - ✔️ c) **The developer using the SDK has made an error in their code or configuration**
 - d) An external API returns an error ❌

7. Which exception is specifically associated with a safety mechanism that monitors content entering the agent's processing flow?
- a) OutputGuardrailTripwireTriggered ✗
 - ✓b) InputGuardrailTripwireTriggered
 - c) ModelBehaviorError ✗
 - d) UserError ✗
8. Both InputGuardrailTripwireTriggered and OutputGuardrailTripwireTriggered exceptions have a common attribute providing details about the triggered guardrail. What is this attribute called?
- a) error_message ✗
 - b) trigger_details ✗
 - ✓c) guardrail_result
 - d) policy_violation ✗
9. When an agent's generated response violates a predefined safety or policy guideline, which exception would be raised?
- a) InputGuardrailTripwireTriggered ✗
 - b) ModelBehaviorError ✗
 - ✓c) OutputGuardrailTripwireTriggered
 - d) AgentResponseError ✗
10. What is a key benefit of having a specific exception hierarchy like AgentsException and its subclasses?
- a) It makes the SDK faster ✗
 - b) It allows for automatic self-correction of agents ✗
 - ✓c) It enables developers to catch and handle distinct categories of errors more effectively
 - d) It reduces the number of tokens consumed ✗
11. If you wanted to catch any exception that originates specifically from the OpenAI Agents SDK, but not general Python exceptions, which exception type would you catch?
- a) Exception ✗
 - ✓b) AgentsException
 - c) ModelBehaviorError ✗
 - d) UserError ✗
12. An agent is configured with max_turns=5. On its 6th attempt to get a final answer, it is still thinking. What exception will be raised?
- a) ModelBehaviorError ✗
 - b) UserError ✗
 - ✓c) MaxTurnsExceeded
 - d) InputGuardrailTripwireTriggered ✗
13. What type of issue would most likely lead to a UserError?
- a) The LLM providing a refusal message ✗
 - ✓b) A typo in the name of a tool provided in the agent's configuration
 - c) An external API being down ✗
 - d) The agent generating a very long message ✗

14. The `guardrail_result` attribute is important for debugging guardrail-related exceptions because it provides:
- a) The recommended next step for the agent ✗
 - b) The time the guardrail was activated ✗
 - ✓c) **Detailed data about the specific guardrail triggered and the reason for the violation**
 - d) A list of all available guardrails ✗
15. If `InputGuardrailTripwireTriggered` occurs, what does it primarily prevent?
- a) The agent from making tool calls ✗
 - ✓b) **Potentially problematic content from reaching the agent's processing logic**
 - c) The agent from generating a final response ✗
 - d) The LLM from receiving any input ✗
-

OpenAI Agents SDK - Guardrails MCQs

Part 1: Guardrail Concepts and Output

1. **Primary purpose of guardrails?**
- a) To enhance LLM's reasoning capabilities ✗
 - ✓b) **To implement safety, validation, and control checks on agent inputs and outputs**
 - c) To optimize agent response time ✗
 - d) To manage external tool integrations ✗
2. **Significance of `tripwire_triggered: bool`?**
- a) Indicates the guardrail performed a basic check ✗
 - ✓b) **Means a serious violation occurred; agent execution halted**
 - c) Signifies successful operation ✗
 - d) Optional field ✗
3. **`output_info` in `GuardrailFunctionOutput` is:**
- a) Main output of the agent ✗
 - ✓b) **Optional info about guardrail's findings**
 - c) Mandatory success/failure string ✗
 - d) The input that triggered the guardrail ✗
4. **Which dataclass represents the result of `InputGuardrail` run?**
- a) `GuardrailFunctionOutput` ✗
 - b) `OutputGuardrailResult` ✗
 - ✓c) **`InputGuardrailResult`**
 - d) `GuardrailStatus` ✗
-

5. **InputGuardrails check:**
 - a) Agent's final output ✗
 - ✓b) **Input messages or data**
 - c) Agent internal state ✗
 - d) Tool call results ✗

6. **If `tripwire_triggered` is True in InputGuardrail, what is raised?**
 - a) OutputGuardrailTripwireTriggered ✗
 - b) ModelBehaviorError ✗
 - ✓c) **InputGuardrailTripwireTriggered**
 - d) MaxTurnsExceeded ✗

7. **Use case for InputGuardrail?**
 - a) Validate final response ✗
 - b) Ensure tool returns JSON ✗
 - ✓c) **Detect off-topic user input**
 - d) Check LLM token usage ✗

8. **OutputGuardrails run on:**
 - a) Input received ✗
 - b) Intermediate reasoning ✗
 - ✓c) **Final agent output**
 - d) Tool call parameters ✗

9. **Extra attributes in OutputGuardrailResult VS InputGuardrailResult?**
 - a) `guardrail_function` and `name` ✗
 - b) `tripwire_triggered` and `output_info` ✗
 - ✓c) **`agent_output` and `agent`**
 - d) `input_history` and `new_items` ✗

10. **Params for `guardrail_function` in OutputGuardrail:**
 - a) (`context`, `agent`, `input_message`) ✗
 - ✓b) **(`context`, `agent`, `agent_output`)**
 - c) (`output`) ✗
 - d) (`tool`, `result`) ✗

11. **Purpose of `name` attribute in Guardrails:**
 - a) LLM name ✗
 - ✓b) **For tracing and logging**
 - c) Agent owner name ✗
 - d) Guardrail type ✗

12. **Benefit of `@input_guardrail()` / `@output_guardrail()`?**
- a) Auto-fix violations ✗
 - ✓b) **Simplify creation of Guardrail instances from functions**
 - c) Bypass checks ✗
 - d) Only for sync functions ✗
13. **Can `@input_guardrail` function be async?**
- ✓a) **Yes, both sync and async are supported**
 - b) No ✗
 - c) Only if it doesn't return output ✗
 - d) Only with `MaybeAwaitable` ✗
14. **Turn function `check_content_safety` into `OutputGuardrail`?**
- a) `OutputGuardrail(func=...)` ✗
 - b) `output_guardrail_func(...)` ✗
 - ✓c) **`@output_guardrail` above function definition**
 - d) `check_content_safety.as_guardrail()` ✗
15. **Use `@input_guardrail(name="...")` — how must decorator be used?**
- a) Without parentheses ✗
 - ✓b) **With parentheses**
 - c) Not possible ✗
 - d) Must be below function ✗
16. **What happens if `OutputGuardrail` returns `tripwire_triggered=True`?**
- a) Run continues ✗
 - b) Auto-rephrase ✗
 - ✓c) **`OutputGuardrailTripwireTriggered` is raised**
 - d) Human review ✗
17. **`InputGuardrail` receives `str | list[TResponseInputItem]`. Meaning?**
- a) Only simple text ✗
 - ✓b) **Can process string or list of OpenAI-like input items**
 - c) List of `RunItem` ✗
 - d) Only `RunContextWrapper` ✗
18. **`agent_output` in `OutputGuardrailResult` is:**
- a) Input to agent ✗
 - ✓b) **Final output checked by the guardrail**

- c) Raw LLM response ✗
- d) Agent summary ✗

19. Why InputGuardrail runs “in parallel”?

- a) To increase tokens ✗
- b) Instant LLM feedback ✗
- ✓c) Early intervention before resource waste
- d) Real-time stream ✗

20. Primary difference between Input vs Output Guardrails?

- a) Prevent errors vs ensure performance ✗
 - ✓b) Input prevents bad content entering; output prevents bad content leaving
 - c) Input = tool, Output = message ✗
 - d) Input = dev-defined, Output = model-defined ✗
-

OpenAI Agents SDK - Model Settings MCQs

Part 1: Core Parameters & Behavior

1. Primary function of `ModelSettings`?

- a) Define agent name ✗
- b) Manage API keys ✗
- ✓c) Hold optional config for calling an LLM
- d) Store conversation history ✗

2. Controls LLM randomness/creativity?

- a) `top_p` ✗
- ✓b) `temperature`
- c) `frequency_penalty` ✗
- d) `max_tokens` ✗

3. How to make output more deterministic?

- a) Higher temperature (e.g., 1.5) ✗
- ✓b) Lower temperature (e.g., 0.2)

- c) None ✗
- d) No effect ✗

4. **What does `top_p` do?**

- a) Penalize frequent tokens ✗
- b) Limit total output tokens ✗
- ✓ c) **Keep most probable tokens above a cumulative threshold**
- d) Force tool call ✗

5. **Reduce LLM repetition – adjust:**

- a) `presence_penalty` ✗
- ✓ b) **`frequency_penalty`**
- c) `repetition_penalty` ✗ (*not in SDK*)
- d) `token_penalty` ✗ (*non-existent*)

6. **What does `tool_choice="required"` mean?**

- a) Can't call tools ✗
- b) Model decides tool/message ✗
- ✓ c) **Model must call a tool**
- d) Call all tools ✗

7. **Effect of `parallel_tool_calls=True`?**

- a) Synchronous execution ✗
- ✓ b) **Generate multiple tool calls in one turn**
- c) Choose tools by parallelism ✗
- d) Prioritize tool execution ✗

8. **`truncation='auto'` means?**

- a) Always truncate ✗
- b) Raise error if input long ✗
- ✓ c) **Auto-drop old messages if input too long**
- d) Disable truncation ✗

9. **Purpose of `max_tokens`?**

- a) Limit input tokens ✗
- b) Limit agent turns ✗
- ✓ c) **Max output tokens generated by LLM**
- d) Limit request count ✗

10. **Which setting is used for explicit internal thoughts?**

- a) `metadata` ✗
- b) `extra_args` ✗

- c) tool_choice ✗
 - ✓d) reasoning
-

Part 2: Customization & Advanced Behavior

11. Add custom fields to HTTP request body?

- a) extra_query ✗
- ✓b) extra_body
- c) extra_headers ✗
- d) metadata ✗

12. What is extra_args for?

- a) Only numeric values ✗
- b) Only strings ✗
- ✓c) Arbitrary keyword args to provider API
- d) List of tools ✗

13. Default behavior of store attribute?

- a) False ✗
- ✓b) True
- c) Depends on provider ✗
- d) Raises error ✗

14. If include_usage=True, what is included?

- a) Agent name ✗
- b) Conversation ID ✗
- ✓c) Token usage info (e.g., input/output tokens)
- d) Error messages ✗

15. Function of resolve() method?

- a) Reset to default ✗
- b) Random settings ✗
- ✓c) Merge override values into new settings instance
- d) Validate settings ✗

16. Result of resolving override with temperature=0.2, max_tokens=100 on global settings (temperature=0.7)?

- a) temp=0.7, max_tokens=None ✗
- b) temp=0.7, max_tokens=100 ✗
- c) temp=0.2, max_tokens=None ✗
- ✓d) temp=0.2, max_tokens=100

17. **Why check API docs when using ModelSettings?**
a) SDK has bugs ✗
b) Ensure all None ✗
✓c) **Not all models/providers support all parameters**
d) Validate `resolve()` ✗
18. **Encourage novelty – which setting?**
a) `frequency_penalty` ✗
✓b) **`presence_penalty` (set high)**
c) `temperature` (low) ✗
d) `max_tokens` (low) ✗
19. **Send custom HTTP headers – which attribute?**
a) `metadata` ✗
b) `extra_body` ✗
✓c) **`extra_headers`**
d) `extra_args` ✗
20. **What config pattern does `resolve()` enable?**
a) Flat ✗
b) Static ✗
✓c) **Hierarchical/layered**
d) Random ✗

OpenAI Agents SDK - Agent Output MCQs

✓✗AgentOutputSchemaBase & AgentOutputSchema – Answer Key with Wrong Answers

1. **What is the primary goal of the AgentOutputSchemaBase and AgentOutputSchema classes?**
✗a) To store the agent's internal reasoning.
✗b) To configure the LLM's temperature.
✓c) **To define, validate, and parse the structured output produced by the LLM.**
✗d) To manage external tool definitions.
2. **AgentOutputSchemaBase is an ABC. What does this imply?**
✗a) It can be directly instantiated.
✓b) **It must be subclassed, and its abstract methods must be implemented.**
✗c) It is a concrete implementation.
✗d) It is only used for plain text outputs.

3. The `is_plain_text()` method in `AgentOutputSchemaBase` determines:
- ☒ a) If the input to the agent is plain text.
 - ☒ b) If the agent can only generate plain text.
 - ☒ c) If the expected output type is simple plain text (vs. a JSON object).
 - ☒ d) If the LLM supports plain text mode.
4. Which method in `AgentOutputSchemaBase` is responsible for returning the JSON schema of the expected output?
- ☒ a) `get_schema()`
 - ☒ b) `json_schema()`
 - ☒ c) `get_output_format()`
 - ☒ d) `schema_definition()`
5. If `is_plain_text()` returns `False`, which other method of `AgentOutputSchemaBase` becomes relevant?
- ☒ a) `name()`
 - ☒ b) `validate_json()`
 - ☒ c) `json_schema()`
 - ☒ d) `is_strict_json_schema()`
6. When `validate_json(json_str: str)` is called, what should it do if the `json_str` is invalid or doesn't conform to the schema?
- ☒ a) Return `None`.
 - ☒ b) Return an empty dictionary.
 - ☒ c) Raise a `ModelBehaviorError`.
 - ☒ d) Log a warning and continue.
7. What is the main parameter passed to the `AgentOutputSchema` constructor?
- ☒ a) `json_schema_dict`
 - ☒ b) `output_type: type[Any]`
 - ☒ c) `agent_name`
 - ☒ d) `validation_function`
8. Why does the documentation "strongly recommend" setting `strict_json_schema=True` in `AgentOutputSchema`?
- ☒ a) Because it makes the agent run faster.
 - ☒ b) Because it reduces LLM token usage.
 - ☒ c) Because it increases the likelihood of the LLM producing correct and valid JSON output.
 - ☒ d) Because it allows for more flexible schema definitions.
9. If `strict_json_schema=True`, what impact does it generally have on the JSON schema communicated to the LLM?
- ☒ a) It makes the schema more complex.
 - ☒ b) It allows for any JSON structure.

- ✓c) It constrains the JSON schema features used, simplifying it for the LLM.
- ✗d) It forces the LLM to output plain text.

10. What does the `AgentOutputSchema` use internally to generate the `json_schema()` based on the `output_type`?

- ✗a) It fetches a predefined schema from a server.
- ✗b) It requires manual schema definition for each type.
- ✓c) It dynamically generates the schema from the provided Python `output_type` (e.g., a Pydantic model).
- ✗d) It reads the schema from a `.json` file.

11. If an `AgentOutputSchema` is initialized with `output_type=str`, what would its `is_plain_text()` method return?

- ✗a) False
- ✓b) True
- ✗c) None
- ✗d) It would raise an error.

12. The `name()` method in `AgentOutputSchema` returns:

- ✗a) A fixed string "AgentOutput".
- ✗b) The name of the agent.
- ✓c) The name of the `output_type` (e.g., "str", "dict", "MyPydanticModel").
- ✗d) A randomly generated UUID.

13. You define an `AgentOutputSchema` with a Pydantic model as its `output_type`. When the LLM returns a JSON string, which method will convert that string into an instance of your Pydantic model and validate it?

- ✗a) `json_schema()`
- ✗b) `is_strict_json_schema()`
- ✗c) `name()`
- ✓d) `validate_json()`

14. What is the consequence if `validate_json()` raises a `ModelBehaviorError`?

- ✗a) The agent will automatically re-attempt the LLM call.
- ✗b) The LLM will receive a warning message.
- ✓c) The agent run will likely halt or enter an error state, as the LLM failed to produce valid output.
- ✗d) The output will be silently dropped.

15. What is the benefit of defining a strict output schema for an agent?

- ✗a) It makes the agent's responses always shorter.
- ✗b) It allows the LLM more freedom in its output format.
- ✓c) It makes the agent's outputs more predictable and easier for downstream code to consume reliably.
- ✗d) It is primarily for performance optimization.

16. Which scenario would most strongly indicate the need to define an `AgentOutputSchema` that is not `is_plain_text`?
- ☒ a) The agent needs to respond with a simple "Hello!".
 - ☒ b) The agent is writing a long essay.
 - ☒ c) **The agent needs to provide structured data, like a user's address with separate fields for street, city, and zip code.**
 - ☒ d) The agent is generating a code snippet.
17. If you create a custom class that needs to handle agent output schema and validation, what must it inherit from?
- ☒ a) `AgentOutputSchema`
 - ☒ b) **`AgentOutputSchemaBase`**
 - ☒ c) `ResponseOutputItem`
 - ☒ d) `ModelSettings`
18. The `AgentOutputSchema` typically uses the provided `output_type` to:
- ☒ a) Determine the agent's personality.
 - ☒ b) Generate the necessary API keys.
 - ☒ c) **Infer the expected JSON structure and validation rules.**
 - ☒ d) Choose the best LLM provider.
19. Why is it important for `AgentOutputSchema` to potentially raise a `ModelBehaviorError` during validation?
- ☒ a) To signal a problem with the user's input.
 - ☒ b) To indicate a network issue.
 - ☒ c) **To clearly identify when the LLM itself has failed to adhere to the expected output format.**
 - ☒ d) To suggest a different LLM model.
20. What is the relationship between `AgentOutputSchemaBase` and `AgentOutputSchema`?
- ☒ a) `AgentOutputSchemaBase` extends `AgentOutputSchema`.
 - ☒ b) They are unrelated but serve similar purposes.
 - ☒ c) **`AgentOutputSchema` is a concrete implementation of the abstract `AgentOutputSchemaBase`.**
 - ☒ d) `AgentOutputSchema` is a utility class for `AgentOutputSchemaBase`.

OpenAI Agents SDK - Function Schema MCQs

✓✗FuncSchema Quiz – Answer Key with Wrong Answers

1. What is the main purpose of the `FuncSchema` dataclass?
- ☒ a) To define how agents interact with users.
 - ☒ b) To manage external API keys for tools.

- ✓c) To capture the schema of a Python function for an LLM to use as a tool.
 - ✗d) To store the execution history of tool calls.
2. Which attribute of `FuncSchema` is crucial for defining the types and validation rules for a function's parameters?
- ✗a) `name`
 - ✗b) `description`
 - ✓c) `params_pydantic_model`
 - ✗d) `params_json_schema`
3. The `params_json_schema` attribute in `FuncSchema` is derived directly from what?
- ✗a) The function's docstring.
 - ✗b) A manual JSON file.
 - ✓c) The `params_pydantic_model`.
 - ✗d) The LLM's capabilities.
4. Setting `strict_json_schema: bool = True` in `FuncSchema` is strongly recommended because it:
- ✗a) Reduces the number of tool calls the LLM makes.
 - ✗b) Makes the tool execution faster.
 - ✓c) Increases the likelihood of the LLM providing correct JSON input for the tool.
 - ✗d) Allows the LLM to call any function.
5. What does the `to_call_args()` method of `FuncSchema` do?
- ✗a) Converts a function call into a JSON string.
 - ✓**b) Converts validated data from the Pydantic model into `(*args, kwargs)` for function execution.
 - ✗c) Generates the Pydantic model from arguments.
 - ✗d) Validates the raw JSON output from the LLM.
6. `FuncDocumentation` is a dataclass used to hold metadata primarily extracted from where?
- ✗a) Function signatures.
 - ✓b) Function docstrings.
 - ✗c) External configuration files.
 - ✗d) LLM descriptions.
7. Which function is used to extract detailed metadata (name, description, parameter descriptions) from a Python function's docstring?
- ✗a) `function_schema`
 - ✓b) `generate_func_documentation`
 - ✗c) `get_function_metadata`
 - ✗d) `extract_docstring`

8. What is the main utility function for creating a `FuncSchema` object from a Python function?
- ☐ a) `FuncSchema.__init__`
 - ☐ b) `generate_func_documentation`
 - ☒ c) `function_schema`
 - ☐ d) `create_tool_schema`
9. If you want the LLM to see a different name for your tool than your Python function's `__name__`, which parameter of `function_schema` would you use?
- ☐ a) `use_docstring_info`
 - ☐ b) `strict_json_schema`
 - ☒ c) `name_override`
 - ☐ d) `docstring_style`
10. The `takes_context: bool` attribute in `FuncSchema` indicates whether the function expects:
- ☐ a) A history of previous tool calls.
 - ☒ b) A `RunContextWrapper` argument as its first parameter.
 - ☐ c) A list of all available agents.
 - ☐ d) A database connection.
11. If `use_docstring_info=False` in `function_schema`, what is the consequence?
- ☐ a) The function will not be able to be called.
 - ☐ b) The `params_pydantic_model` will not be generated.
 - ☒ c) The description and `param_descriptions` for the `FuncSchema` will not be derived from the docstring.
 - ☐ d) The `strict_json_schema` will automatically become False.
12. What does `FuncSchema.signature` hold?
- ☐ a) A hash of the function's code.
 - ☐ b) A unique ID for the function.
 - ☒ c) The `inspect.Signature` object of the original Python function.
 - ☐ d) The LLM's "signature" for generating calls.
13. The `params_json_schema` is the actual JSON representation sent to the LLM. What is its primary role from the LLM's perspective?
- ☐ a) To determine the function's execution priority.
 - ☒ b) To inform the LLM about the required arguments and their types for the tool call.
 - ☐ c) To tell the LLM if the function is asynchronous.
 - ☐ d) To help the LLM decide if it should generate a message or a tool call.
14. When would you typically set `description_override` in `function_schema`?
- ☐ a) When the function has no docstring.
 - ☐ b) When you want the LLM to ignore the function.
 - ☒ c) When the docstring's description isn't ideal for how the LLM should interpret the tool.
 - ☐ d) When the function's name is too long.

15. What type of model must `params_pydantic_model` inherit from?
- ☐ a) `dict`
 - ☐ b) `Any`
 - ☒ c) `pydantic.BaseModel`
 - ☐ d) `FunctionArguments`
16. Which of the following is NOT a direct attribute of `FuncSchema`?
- ☒ d) `docstring_style`
 - ☐ a) `name`
 - ☐ b) `signature`
 - ☐ c) `takes_context`
17. If `generate_func_documentation()` cannot auto-detect the docstring style, what can you do?
- ☐ a) It will raise an error.
 - ☒ b) You can explicitly provide the style parameter.
 - ☐ c) The description will remain `None`.
 - ☐ d) You must rewrite the docstring.
18. What is the benefit of `FuncSchema` being able to handle functions that `takes_context`?
- ☐ a) It allows the LLM to directly modify the context.
 - ☐ b) It simplifies the tool execution by removing the need for context.
 - ☒ c) It allows tool functions to access shared run-specific data and dependencies.
 - ☐ d) It is purely for logging purposes.
19. In the overall agent workflow, where does the `FuncSchema` primarily get used?
- ☐ a) It's returned as the final output of the agent.
 - ☐ b) It's part of the LLM's internal state.
 - ☒ c) It's sent to the LLM to define available tools for function calling.
 - ☐ d) It's used by `InputGuardrails` for validation.
20. What is the role of `FuncDocumentation` in the creation of a `FuncSchema`?
- ☐ a) It replaces the `FuncSchema` entirely.
 - ☒ b) It acts as an intermediate step to extract natural language descriptions from docstrings for the `FuncSchema`.
 - ☐ c) It validates the `FuncSchema`.
 - ☐ d) It converts `FuncSchema` back into a Python function.
-

OpenAI Agents SDK - Model Interface MCQs

✓✗Agents SDK – Model Interface Quiz Answer Key

1. **What is the main benefit of the "Model interface" in the Agents SDK?**
 - ✗a) It hardcodes specific LLM providers.
 - ✗b) It always uses the latest OpenAI model.
 - ✓c) **It decouples the agent's logic from specific LLM implementations, allowing for easy swapping of models/providers.**
 - ✗d) It automatically optimizes LLM calls for cost.
2. **Which Enum is used to configure the level of logging and data capture for LLM interactions?**
 - ✗a) LogLevel
 - ✓b) **ModelTracing**
 - ✗c) TraceMode
 - ✗d) ModelConfig
3. **If you want to enable tracing for LLM calls but specifically exclude the actual input/output data, which ModelTracing option would you choose?**
 - ✗a) DISABLED
 - ✗b) ENABLED
 - ✓c) **ENABLED_WITHOUT_DATA**
 - ✗d) VERBOSE
4. **The `Model` class is an ABC. What does this indicate about it?**
 - ✗a) It is a concrete implementation of an LLM.
 - ✓b) **It defines the abstract contract for any LLM wrapper and cannot be directly instantiated.**
 - ✗c) It represents a specific LLM provider.
 - ✗d) It is a utility class for managing model settings.
5. **Which `Model` method is designed for making a single, complete LLM request and waiting for the full response?**
 - ✗a) `stream_response`
 - ✓b) **`get_response`**
 - ✗c) `send_query`
 - ✗d) `process_input`
6. **The `get_response` method of `Model` returns what type of object?**
 - ✗a) `AsyncIterator[TResponseStreamEvent]`
 - ✗b) `str`
 - ✓c) **`ModelResponse`**
 - ✗d) `dict`

7. The `stream_response` method of `Model` returns what type of object?
- ☐ a) `ModelResponse`
 - ☐ b) `str`
 - ☐ c) `list[TResponseStreamEvent]`
 - ☒ d) `AsyncIterator[TResponseStreamEvent]`
8. Both `get_response` and `stream_response` methods of `Model` accept which common parameter to configure LLM behavior (e.g., temperature, max tokens)?
- ☐ a) `system_instructions`
 - ☐ b) `tools`
 - ☒ c) `model_settings`
 - ☐ d) `tracing`
9. What is the purpose of the `output_schema` parameter in the `Model` methods?
- ☐ a) To define the input format for the LLM.
 - ☐ b) To specify the data type of `system_instructions`.
 - ☒ c) To guide the LLM on the expected structured output format using `AgentOutputSchemaBase`.
 - ☐ d) To limit the total number of turns.
10. The `ModelProvider` base interface is responsible for:
- ☐ a) Executing LLM calls directly.
 - ☒ b) Looking up and providing `Model` instances by name.
 - ☐ c) Handling all agent-level business logic.
 - ☐ d) Managing conversation history.
11. If you wanted to switch from using OpenAI's `gpt-4o` to Google's `gemini-pro`, which part of the model interface would primarily be swapped or configured?
- ☐ a) The `ModelSettings`.
 - ☐ b) The `system_instructions`.
 - ☒ c) The `ModelProvider` implementation.
 - ☐ d) The `tracing` setting.
12. The `tools` parameter passed to `Model` methods (`get_response`, `stream_response`) contains:
- ☐ a) References to external APIs.
 - ☐ b) Configuration for internal agent mechanisms.
 - ☒ c) A list of `Tool` objects that the LLM is aware of and can call.
 - ☐ d) User input messages.
13. The `Model` methods are async. What does this imply about how they should typically be called in Python?
- ☐ a) They must be called in a separate thread.
 - ☒ b) They should be awaited.
 - ☐ c) They return immediately.
 - ☐ d) They are synchronous by default.

14. What type of input does the `input` parameter of `Model` methods expect?
- ☐ a) Only `str`
 - ☒ b) `str | list[TResponseInputItem]`
 - ☐ c) `ModelResponse`
 - ☐ d) `AsyncIterator`
15. `ModelTracing.ENABLED` includes:
- ☐ a) Only request headers.
 - ☐ b) Only response bodies.
 - ☒ c) All data, including inputs and outputs.
 - ☐ d) No data, only timestamps.
16. Which of the following is NOT a parameter commonly passed to `Model`'s `get_response` or `stream_response` methods?
- ☐ a) `system_instructions`
 - ☐ b) `model_settings`
 - ☒ c) `run_context`
 - ☐ d) `tools`
17. If a custom LLM integration needs to be built for the Agents SDK, what base class must the LLM wrapper inherit from?
- ☐ a) `ModelProvider`
 - ☒ b) `Model`
 - ☐ c) `Agent`
 - ☐ d) `BaseModel`
18. What is the role of `TResponseStreamEvent` in the `stream_response` method's return type?
- ☐ a) It represents the final, complete response.
 - ☐ b) It's an error type.
 - ☒ c) It represents individual chunks or events received during a streaming response from the LLM.
 - ☐ d) It's the type of the tool call.
19. The `previous_response_id` parameter is described as "Generally not used by the model, except for the OpenAI Responses API." What does this suggest about its primary purpose?
- ☐ a) It's for debugging purposes only.
 - ☐ b) It's a universal identifier for all LLM calls.
 - ☒ c) It's an optional, provider-specific identifier used to maintain context or state in certain APIs.
 - ☐ d) It's used for caching.
20. What is the core responsibility of the `get_model(model_name: str | None)` method within `ModelProvider`?

- ☒ a) To train a new LLM.
 - ☒ b) To initialize the ModelSettings for a given model.
 - ☒ c) **To retrieve a specific LLM Model instance by its name.**
 - ☒ d) To validate the model's output.
-

OpenAI Agents SDK - OpenAIChatCompletionsModel MCQs

1. **The OpenAIChatCompletionsModel class is a concrete implementation for interacting with which specific API?**
 - ☒ a) OpenAI Assistants API
 - ☒ b) Google Gemini API
 - ☒ c) **OpenAI Chat Completions API**
 - ☒ d) Local LLM API

2. **What is the base class that OpenAIChatCompletionsModel inherits from?**
 - ☒ a) ABC
 - ☒ b) ModelProvider
 - ☒ c) **Model**
 - ☒ d) ChatModelBase

3. **The stream_response method in OpenAIChatCompletionsModel is designed to:**
 - ☒ a) Return a single, complete response from the LLM.
 - ☒ b) **Yield partial messages and usage information as they are generated.**
 - ☒ c) Only process input messages without generating output.
 - ☒ d) Store the entire conversation history in a stream.

4. **When stream_response yields usage information, how accurate is it for ongoing streams?**
 - ☒ a) It is always perfectly accurate and final.
 - ☒ b) It is unavailable until the stream completes.
 - ☒ c) **It is updated incrementally but might not be final until the last chunk is processed.**
 - ☒ d) It only tracks input tokens during streaming.

5. **Which of the following parameters is *not* listed as an input to the stream_response method of OpenAIChatCompletionsModel?**
 - ☒ a) model_settings
 - ☒ b) tools
 - ☒ c) **api_key**
 - ☒ d) output_schema

☐ *API keys are handled at the client/provider level, not passed directly.*

6. The `stream_response` method returns an `AsyncIterator` yielding `TResponseStreamEvent`. What do these events represent?
- ☒ a) Complete LLM responses.
 - ☒ b) Errors encountered during the stream.
 - ☒ c) **Individual chunks or partial data received during the streaming process.**
 - ☒ d) Summaries of the conversation.
7. By inheriting from `Model`, `OpenAIChatCompletionsModel` guarantees that it provides a consistent interface for:
- ☒ a) Only streaming interactions.
 - ☒ b) Only non-streaming interactions.
 - ☒ c) **Both streaming and non-streaming (implied `get_response` implementation) interactions with LLMs.**
 - ☒ d) Only local model inference.
8. The `tools` parameter passed to `stream_response` would be translated by `OpenAIChatCompletionsModel` into what format for the OpenAI API?
- ☒ a) A plain string of tool names.
 - ☒ b) A list of Python function objects.
 - ☒ c) **OpenAI's specific function-calling (or tool-calling) JSON format.**
 - ☒ d) A dictionary of tool descriptions.
9. What is the role of `model_settings` when calling `OpenAIChatCompletionsModel.stream_response`?
- ☒ a) To define the `system_instructions`.
 - ☒ b) **To configure LLM-specific parameters like temperature, max_tokens, and tool choice for the OpenAI call.**
 - ☒ c) To specify the tracing level.
 - ☒ d) To manage handoff points.
10. If an agent is built using the generic `Model` interface, could it seamlessly use `OpenAIChatCompletionsModel`?
- ☒ a) **Yes, because `OpenAIChatCompletionsModel` implements the `Model` interface.**
 - ☒ b) No, it would require significant code changes.
 - ☒ c) Only if streaming is disabled.
 - ☒ d) Only if no tools are used.

OpenAI Agents SDK - OpenAIResponsesModel & Converter MCQs

☒ OpenAIResponsesModel Quiz – Answer Key

1. `OpenAIResponsesModel` is a concrete implementation of which base class?
- ☒ a) `OpenAIChatCompletionsModel`
 - ☒ b) **`Model`**
 - ☒ c) `ModelProvider`
 - ☒ d) `BaseModel`

2. What is the key distinction between `OpenAIChatCompletionsModel` and `OpenAIResponsesModel`?
☒ a) One supports streaming, the other does not.
☒ b) One supports tools, the other does not.
☒ c) They interact with different OpenAI APIs (Chat Completions vs. Responses API).
☒ d) One is for input, the other for output.
3. The `stream_response` method of `OpenAIResponsesModel` returns an `AsyncIterator` yielding what type of events?
☒ a) `TResponseStreamEvent`
☒ b) `ResponseStreamEvent`
☒ c) `OpenAIEvent`
☒ d) `StreamChunk`
4. When `OpenAIResponsesModel.stream_response` yields data, what information does it include besides partial messages?
☒ a) Agent internal state
☒ b) Tool call history
☒ c) Usage information
☒ d) Current time
5. What is the primary role of the `Converter` class within the `openai_responses.py` module?
☒ a) To manage API authentication.
☒ b) To cache LLM responses.
☒ c) To translate data formats between the Agents SDK and the OpenAI Responses API.
☒ d) To perform sentiment analysis on LLM outputs.
6. Which of the following is most likely a task performed by the `Converter`?
☒ a) Executing Python functions called by the LLM.
☒ b) Determining the temperature for an LLM call.
☒ c) Transforming a `Tool` object into the OpenAI Responses API's tool definition JSON.
☒ d) Handling network retries for API calls.
7. If the Agents SDK uses a generic `ModelSettings` object, but the OpenAI Responses API expects a specific JSON structure for model parameters, which component is responsible for this translation?
☒ a) `OpenAIResponsesModel` directly
☒ b) The `Model` base class
☒ c) The `Converter` class
☒ d) `AgentOutputSchema`
8. What is implied about the `get_response` method for `OpenAIResponsesModel` given its inheritance from `Model`?
☒ a) It is not implemented.
☒ b) It only supports streaming.
☒ c) It must also be implemented to handle non-streaming requests.
☒ d) It is handled automatically by the base class.

9. **The `ResponseStreamEvent` type, as used by `OpenAIResponsesModel.stream_response`, suggests:**
- ☒ a) It's a generic event type for any streaming API.
 - ☒ b) It's an error-only event type.
 - ☒ c) **It's a specialized, potentially richer event type specific to the OpenAI Responses API.**
 - ☒ d) It only contains token counts.
10. **The consistency of parameters (`system_instructions`, `model_settings`, `tools`) across Model implementations (e.g., `OpenAIChatCompletionsModel` and `OpenAIResponsesModel`) enables:**
- ☒ a) Automatic model selection by the SDK.
 - ☒ b) Lower API costs.
 - ☒ c) **Easy interchangeability of LLM backends without major code changes.**
 - ☒ d) Faster LLM response times.
11. **Why would a developer choose `OpenAIResponsesModel` over `OpenAIChatCompletionsModel`?**
- ☒ a) `OpenAIChatCompletionsModel` does not support tools.
 - ☒ b) `OpenAIResponsesModel` is generally faster.
 - ☒ c) **To leverage specific features, richer streaming events, or advanced prompt paradigms offered by the OpenAI Responses API.**
 - ☒ d) `OpenAIResponsesModel` does not require an OpenAI API key.
12. **The `Converter` helps to bridge the gap between:**
- ☒ a) User input and agent output.
 - ☒ b) Synchronous and asynchronous calls.
 - ☒ c) **The SDK's generic interfaces and a specific LLM provider's API formats.**
 - ☒ d) Different agent personalities.
13. **If `OpenAIResponsesModel` receives an `AgentOutputSchemaBase` object, which component is most likely to translate this into the precise format the OpenAI Responses API expects for structured output?**
- ☒ a) The `ModelSettings`
 - ☒ b) The `tools` list
 - ☒ c) **The `Converter`**
 - ☒ d) The `previous_response_id`
14. **What is the role of `previous_response_id` in the `OpenAIResponsesModel.stream_response` method?**
- ☒ a) To specify the maximum length of the response.
 - ☒ b) To identify the user making the request.
 - ☒ c) **To provide contextual continuity for the OpenAI Responses API.**
 - ☒ d) To track the total usage.
15. **The `OpenAIResponsesModel` is part of a broader design pattern where `Model` is an abstraction. What does this allow the SDK to be?**
- ☒ a) Tightly coupled to OpenAI.
 - ☒ b) Limited to a single LLM.
 - ☒ c) **Flexible and extensible to support various LLM providers.**
 - ☒ d) Strictly for chat-based applications.

OpenAI Agents SDK - MCP Servers MCQs

✔✕MCP Server Quiz – Answer Key

1. What is the primary purpose of the `MCPServer` base class?
 - ✕a) To manage LLM context windows.
 - ✔b) To define a common interface for interacting with external services hosting tools.
 - ✕c) To configure OpenAI API keys.
 - ✕d) To store agent conversation history.
2. Which of the following is an abstract method that all concrete `MCPServer` implementations must provide?
 - ✕a) `get_version()`
 - ✕b) `send_heartbeat()`
 - ✕c) `process_message()`
 - ✔d) `call_tool()`
3. The `connect()` and `cleanup()` methods in `MCPServer` are both:
 - ✕a) Synchronous
 - ✔b) Asynchronous
 - ✕c) Optional
 - ✕d) Used only for local servers
4. What does the `list_tools()` method of `MCPServer` return?
 - ✕a) A dictionary of tool names and descriptions.
 - ✕b) A list of `FuncSchema` objects.
 - ✔c) A list of `Tool` objects.
 - ✕d) The current status of the server.
5. `MCPServerStdioParams` is used to configure which type of MCP server?
 - ✕a) Remote HTTP server
 - ✕b) Server-Sent Events server
 - ✔c) Standard input/output (subprocess) server
 - ✕d) WebSocket server
6. In `MCPServerStdioParams`, what does the `command` attribute specify?
 - ✕a) The name of the tool to call.
 - ✕b) The URL of the server.
 - ✔c) The executable to run to start the server (e.g., `python`).
 - ✕d) A command to send to the LLM.

7. The `cache_tools_list` parameter in `MCPServerStdio` (and others) is primarily for:
- ☒ a) Ensuring tool definitions are always up-to-date.
 - ☒ b) Reducing memory usage on the client side.
 - ☒ c) **Improving latency by avoiding repeated round-trips to list tools if they are static.**
 - ☒ d) Forcing the server to reload its tools.
8. Which method would you call to force a refresh of the cached tool list on an `MCPServer`?
- ☒ a) `refresh_tools()`
 - ☒ b) **`invalidate_tools_cache()`**
 - ☒ c) `reset_cache()`
 - ☒ d) `clear_tools()`
9. `MCPServerSseParams` requires which essential attribute to define the server connection?
- ☒ a) `command`
 - ☒ b) `encoding`
 - ☒ c) **`url`**
 - ☒ d) `cwd`
10. Which MCP server implementation is suitable for connecting to a remote service that streams updates over HTTP?
- ☒ a) `MCPServerStdio`
 - ☒ b) **`MCPServerSse`**
 - ☒ (Also accepted: c) `MCPServerStreamableHttp`
 - ☒ d) `MCPServerLocal`
11. What is the role of the `sse_read_timeout` attribute in `MCPServerSseParams`?
- ☒ a) Timeout for the initial HTTP request.
 - ☒ b) Timeout for connecting to the server process.
 - ☒ c) **The timeout for the Server-Sent Events (SSE) connection itself.**
 - ☒ d) Timeout for a single tool call.
12. `MCPServerStreamableHttp` likely uses a transport protocol that is a specialization or evolution of:
- ☒ a) Stdio
 - ☒ b) **HTTP with Server-Sent Events (SSE)**
 - ☒ c) WebSockets
 - ☒ d) UDP
13. If `cache_tools_list` is `True` and the tools on the server actually change, what must be done for the client to see the new tools?
- ☒ a) Restart the agent entirely.
 - ☒ b) Call `connect()` again.
 - ☒ c) **Call `invalidate_tools_cache()`.**
 - ☒ d) It updates automatically.

14. The `call_tool()` method returns a `CallToolResult`. What does this object typically contain?
- ☒ a) The name of the tool called.
 - ☒ b) The arguments passed to the tool.
 - ☒ c) The result or outcome of the tool invocation.
 - ☒ d) The tool's description.
15. What type of parameter is `terminate_on_close` specifically associated with?
- ☒ a) `MCPServerStdioParams`
 - ☒ b) `MCPServerSseParams`
 - ☒ c) `MCPServerStreamableHttpParams`
 - ☒ d) All `MCPServer` parameters
16. Which attribute in `MCPServerStdioParams` is used to specify the initial working directory for the subprocess?
- ☒ a) `env`
 - ☒ b) `command`
 - ☒ c) `cwd`
 - ☒ d) `args`
17. The `name` property of `MCPServer` is:
- ☒ a) An optional field that defaults to `None`.
 - ☒ b) An abstract property that must return a readable string.
 - ☒ c) Used only for debugging purposes.
 - ☒ d) The unique identifier for the server.
18. What would be a good use case for `MCPServerStdio`?
- ☒ a) Connecting to a public web API.
 - ☒ b) Integrating a local Python script running as a separate process that exposes tools.
 - ☒ c) Receiving real-time stock updates from a remote server.
 - ☒ d) Communicating with an LLM provider.
19. The `client_session_timeout_seconds` parameter refers to:
- ☒ a) The timeout for the `connect()` method.
 - ☒ b) The maximum duration a tool call can take.
 - ☒ c) The read timeout for the underlying MCP ClientSession.
 - ☒ d) The time before the server automatically cleans up.
20. The Model Context Protocol (MCP) aims to facilitate communication between agents and:
- ☒ a) Different LLM providers.
 - ☒ b) Other agents in a multi-agent system.
 - ☒ c) External, self-contained services or "tools."
 - ☒ d) Human users for feedback.

OpenAI Agents SDK - MCPUtil MCQs

1. **What is the primary role of the `MCPUtil` class?**
 - ☐ a) To manage MCP server authentication.
 - ☒ b) **To provide utilities for interoperability between MCP and Agents SDK tools.**
 - ☐ c) To implement new MCP transport protocols.
 - ☐ d) To store MCP server configuration.
2. **The `get_all_function_tools` method is used to retrieve tools from:**
 - ☐ a) A single MCP server.
 - ☒ b) **A list of multiple MCP servers.**
 - ☐ c) Local Python functions only.
 - ☐ d) Directly from the LLM.
3. **What is the return type of `get_all_function_tools`?**
 - ☐ a) `list[FuncSchema]`
 - ☐ b) `list[MCPServer]`
 - ☒ c) **`list[Tool]`**
 - ☐ d) `list[str]`
4. **The `convert_schemas_to_strict` parameter in `MCPUtil` methods is important for:**
 - ☐ a) Reducing network latency.
 - ☐ b) Caching tool definitions.
 - ☒ c) **Ensuring LLMs generate correct arguments by using strict JSON schemas.**
 - ☐ d) Changing the tool's name.
5. **Which method would you use if you only wanted to retrieve tools from one specific MCP server?**
 - ☐ a) `get_all_function_tools`
 - ☒ b) **`get_function_tools`**
 - ☐ c) `to_function_tool`
 - ☐ d) `invoke_mcp_tool`
6. **The `to_function_tool` method converts a generic `Tool` object into what specific type?**
 - ☐ a) `FuncSchema`
 - ☐ b) `CallToolResult`
 - ☒ c) **`FunctionTool`**
 - ☐ d) `MCPServer`

7. **Why is the `server: MCPServer` parameter needed in `to_function_tool`?**
☒ a) To get the server's name for logging.
☒ b) To check if the server is currently connected.
☒ c) **To establish the link back to the specific MCP server that will execute the tool.**
☒ d) To invalidate the server's tool cache.
8. **What is the purpose of the `invoke_mcp_tool` method?**
☒ a) To discover available tools on an MCP server.
☒ b) To connect to an MCP server.
☒ c) **To execute a specific tool on an MCP server with provided arguments.**
☒ d) To convert an MCP tool's result to a different format.
9. **The `input_json` parameter in `invoke_mcp_tool` is:**
☒ a) The raw user input message.
☒ b) The tool's description.
☒ c) **The arguments for the tool call, provided as a JSON string.**
☒ d) A list of previous tool call results.
10. **What is the return type of `invoke_mcp_tool`?**
☒ a) `dict[str, Any]`
☒ b) **`CallToolResult`**
☒ c) `Any`
☒ d) `str`
11. **All methods within `MCPUtil` are defined as:**
☒ a) Instance methods.
☒ b) **Class methods.**
☒ c) Static methods.
☒ d) Abstract methods.
12. **Sequence of `MCPUtil` operations to use an MCP tool:**
☒ a) `invoke_mcp_tool` → `get_function_tools`
☒ b) `to_function_tool` → `invoke_mcp_tool` → `get_function_tools`
☒ c) **`get_function_tools` (or `get_all_function_tools`) → `to_function_tool` → `invoke_mcp_tool`**
☒ d) `get_all_function_tools` → `invoke_mcp_tool`
13. **The context: `RunContextWrapper[Any]` parameter in `invoke_mcp_tool` is for:**
☒ a) Storing the tool's output.
☒ b) Managing network connections.
☒ c) **Providing the tool with access to the agent's current runtime context.**
☒ d) Specifying the LLM model to use.

14. If `convert_schemas_to_strict` is `False`, potential drawback:
- ☒ a) Slower tool execution.
 - ☒ b) Inability to discover tools.
 - ☒ c) **LLMs might struggle to generate valid JSON arguments due to a less constrained schema.**
 - ☒ d) Tools will not be able to return a result.
15. MCPUtil bridges:
- ☒ a) Python and JavaScript code.
 - ☒ b) Synchronous and asynchronous operations.
 - ☒ c) **The Agents SDK's internal tool representation and external MCP-compliant services.**
 - ☒ d) Text-based and image-based inputs.
16. Which is NOT a direct parameter to `invoke_mcp_tool`?
- ☒ a) server
 - ☒ b) tool
 - ☒ c) **output_schema**
 - ☒ d) input_json
17. Which method benefits most from tool list caching?
- ☒ a) `invoke_mcp_tool`
 - ☒ b) `to_function_tool`
 - ☒ c) **`get_function_tools` / `get_all_function_tools`**
 - ☒ d) None of them directly...
18. MCPUtil methods are async where appropriate. This means:
- ☒ a) Block execution until result is available.
 - ☒ b) **Perform non-blocking I/O operations, typically when interacting with MCP servers.**
 - ☒ c) Run in a separate thread.
 - ☒ d) Execute multiple tools in parallel.
19. Why does MCPUtil exist as a separate class?
- ☒ a) To avoid circular dependencies.
 - ☒ b) To keep `MCPServer` focused on its core...
 - ☒ c) To provide a central point...
 - ☒ d) **All of the above.**
20. If `invoke_mcp_tool` returns a `str`, what is the most likely format?
- ☒ a) Plain text summary
 - ☒ b) **JSON string**
 - ☒ c) XML string
 - ☒ d) Python literal
-

Tracing

OpenAI Agents SDK - Tracing Module MCQs

1. What is the main purpose of the Tracing module in the Agents SDK?

- ☒ c) To provide observability and debugging for agent workflows.
 - ☐ a) To manage API authentication.
 - ☐ b) To handle asynchronous operations.
 - ☐ d) To store conversation history persistently.
-

2. A Trace in the tracing module represents:

- ☒ b) A complete logical workflow or request.
 - ☐ a) A single operation within a workflow.
 - ☐ c) An error log entry.
 - ☐ d) A configuration setting for the agent.
-

3. A Span in the tracing module represents:

- ☒ c) A single operation or unit of work within a trace.
 - ☐ a) The entire conversation flow.
 - ☐ b) A collection of related traces.
 - ☐ d) An external API call only.
-

4. Which abstract base class defines the interface for creating and managing traces and spans?

- ☒ c) **TraceProvider**
 - ☐ a) TracingProcessor
 - ☐ b) SpanData
 - ☐ d) Trace
-

5. Which method on TraceProvider would you use to add a component that processes trace and span events?

- ☒ c) **register_processor()**
 - ☐ a) create_span()
 - ☐ b) get_current_trace()
 - ☐ d) set_disabled()
-

6. What is the primary responsibility of a TracingProcessor?

- ☒ c) **To consume and process trace and span events (e.g., for logging or export).**
- ☐ a) To generate unique trace IDs.

- ☒ b) To start and finish spans.
 - ☒ d) To disable tracing dynamically.
-

7. Which TracingProcessor method should not block or raise exceptions to avoid interfering with application logic?

- ☒ d) on_span_end()
 - ☒ a) shutdown()
 - ☒ b) force_flush()
 - ☒ c) on_trace_start()
-

8. GenerationSpanData is a specific type of SpanData that would typically capture details about:

- ☒ c) LLM model input/output messages, model configuration, and usage.
 - ☒ a) Tool invocation inputs and outputs.
 - ☒ b) Agent handoffs.
 - ☒ d) Custom, arbitrary data.
-

9. If you want to create a span that records arbitrary, structured data for a custom operation, which SpanData type would you use?

- ☒ c) CustomSpanData
 - ☒ a) AgentSpanData
 - ☒ b) FunctionSpanData
 - ☒ d) ResponseSpanData
-

10. What are the two primary methods used to define the boundaries (start and end) of a Span or Trace?

- ☒ c) start() and finish()
 - ☒ a) create() and destroy()
 - ☒ b) init() and close()
 - ☒ d) open() and shut()
-

11. What is the benefit of using a with statement (context manager) with span creation functions like tracing.generation_span() ?

- ☒ c) It automatically calls start() when entering and finish() when exiting the block.
 - ☒ a) It makes the span immutable.
 - ☒ b) It automatically exports the span to a database.
 - ☒ d) It disables the span if an error occurs.
-

12. If you create a new Trace or Span but do not call its start() method or use it as a context manager, what happens?

- ☒ b) It will not be recorded by the tracing system.
 - ☒ a) It will automatically start when the first child span is created.
 - ☒ c) It will only record errors.
 - ☒ d) It will throw an error immediately.
-

13. The `parent` parameter in span creation functions (e.g., `agent_span()`) is used to:
- ✓c) Explicitly link the new span to an existing parent trace or span, forming a hierarchy.
 - ✗a) Specify the LLM model to be used.
 - ✗b) Define the ownership of the span in a multi-user system.
 - ✗d) Determine if the span should be disabled.
-

14. To temporarily stop all tracing without changing individual span creation calls, which global function would you use?
- ✓d) `set_tracing_disabled(True)`
 - ✗a) `force_flush()`
 - ✗b) `set_trace_provider(None)`
 - ✗c) `remove_all_processors()`
-

15. What format does `time_iso()` method on `TraceProvider` return the current time in?
- ✓c) ISO 8601 format
 - ✗a) Unix timestamp
 - ✗b) Milliseconds since epoch
 - ✗d) A localized date and time string
-

16. Which `SpanData` type would be most appropriate for recording a control transfer between two different agents in a multi-agent system?
- ✓b) `HandoffSpanData`
 - ✗a) `AgentSpanData`
 - ✗c) `CustomSpanData`
 - ✗d) `ResponseSpanData`
-

17. The `trace_id` property on a `Trace` object ensures:
- ✓c) The trace can be uniquely identified and correlated across different operations.
 - ✗a) The trace is always processed synchronously.
 - ✗b) The trace can be easily exported to an external API key.
 - ✗d) The trace is immutable once created.
-

18. If a `TracingProcessor` queues spans before sending them to an external system, which method would you call to ensure all pending spans are immediately dispatched?
- ✓c) `force_flush()`
 - ✗a) `shutdown()`
 - ✗b) `on_span_end()`
 - ✗d) `register_processor()`
-

19. What is the primary difference between `generation_span` and `response_span` creation utilities?
- ✓c) `generation_span` captures full model generation details (input, output, model, usage), while `response_span` might focus on a simpler model response identifier.
 - ✗a) `generation_span` is for tool calls, `response_span` is for LLM calls.

- ☒ b) `generation_span` captures only errors, `response_span` captures success.
 - ☒ d) There is no functional difference; they are aliases.
-

20. What is the recommended way to generate a new unique span identifier if you need to provide one explicitly?
- ☒ c) Use `util.gen_span_id()`.
 - ☒ a) Hardcode a random string.
 - ☒ b) Use `uuid.uuid4().hex`.
 - ☒ d) Rely on the `TraceProvider` to generate it implicitly.
-

OpenAI Agents SDK - Trace and Span Creation MCQs

1. When creating a new trace using `tracing.trace()`, which parameter is required?
- ☒ c) `workflow_name`
 - ☒ a) `trace_id`
 - ☒ b) `group_id`
 - ☒ d) `metadata`
-
2. A trace created with `tracing.trace(...)` will automatically start recording upon creation.
- ☒ b) False
 - ☒ a) True
-
3. What is the recommended way to ensure a trace or span automatically starts and finishes, even if exceptions occur?
- ☒ c) Use it as a context manager (with `trace(...):`)
 - ☒ a) Always call `start()` and `finish()` in a `try...finally` block.
 - ☒ b) Set `auto_start=True` during creation.
 - ☒ d) Call `tracing.get_current_trace().start()` manually.
-
4. Which function would you use to create a span specifically designed to track an LLM's reasoning process, including its input and output messages?
- ☒ c) `generation_span()`
 - ☒ a) `function_span()`
 - ☒ b) `agent_span()`
 - ☒ d) `response_span()`
-
5. The `disabled=True` parameter in span/trace creation functions will:
- ☒ b) Return a Trace or Span object, but it will not be recorded.
 - ☒ a) Cause an error if tracing is globally enabled.
 - ☒ c) Only record errors within that trace/span.
 - ☒ d) Prevent the `start()` method from being called.
-

6. Which function is used to create a span for a tool or function invocation?
- ✓b) `function_span()`
 - ✗a) `agent_span()`
 - ✗c) `custom_span()`
 - ✗d) `mcp_tools_span()`
-
7. To explicitly link a newly created span to an existing trace or span as its parent, which parameter should be used?
- ✓c) `parent`
 - ✗a) `workflow_name`
 - ✗b) `group_id`
 - ✗d) `trace_id`
-
8. What is the primary difference in data captured by `generation_span` compared to `response_span`?
- ✓c) `generation_span` captures detailed input messages, model config, and usage; `response_span` focuses more on the final output/response object.
 - ✗a) `response_span` includes tool calls; `generation_span` does not.
 - ✗b) `response_span` is for agent communication; `generation_span` is for user output.
 - ✗d) `generation_span` records only errors; `response_span` records successes.
-
9. Which function allows you to create a versatile span for any custom operation and attach arbitrary data?
- ✓c) `custom_span()`
 - ✗a) `agent_span()`
 - ✗b) `function_span()`
 - ✗d) `guardrail_span()`
-
10. If you need to retrieve the currently active span in the execution context, which function would you call?
- ✓b) `get_current_span()`
 - ✗a) `get_current_trace()`
 - ✗c) `get_active_context()`
 - ✗d) `get_global_span()`
-
11. The `metadata` parameter in `tracing.trace()` allows you to attach:
- ✓b) Any arbitrary dictionary of user-defined information.
 - ✗a) Only pre-defined key-value pairs.
 - ✗c) Only security credentials.
 - ✗d) A list of child span IDs.
-
12. A `handoff_span` is specifically designed to track:
- ✓c) The transfer of control or responsibility between different agents.
 - ✗a) Data transfer between a model and a database.
 - ✗b) The parsing of user input.
 - ✗d) The loading of agent configurations.

-
13. Which of the following span creation functions would be most appropriate for recording whether a "PII Detector" rule was activated?
- ✓d) `guardrail_span()`
 - ✗a) `function_span()`
 - ✗b) `custom_span()`
 - ✗c) `agent_span()`
-
14. The `result` parameter in `mcp_tools_span()` typically contains:
- ✓c) A list of tool names or descriptions discovered from the MCP server.
 - ✗a) The input query to the MCP server.
 - ✗b) The duration of the MCP call.
 - ✗d) The error message from the MCP server.
-
15. In `transcription_span()` and `speech_span()`, what does the `input` parameter typically represent?
- ✓c) Base64 encoded audio bytes (for transcription) or text (for speech synthesis).
 - ✗a) The LLM prompt.
 - ✗b) The name of the audio file.
 - ✗d) A URL to an audio stream.
-
16. If you omit the `span_id` parameter when creating a span, what generally happens?
- ✓c) The system will automatically generate a unique ID for the span.
 - ✗a) An error is raised, requiring a `span_id`.
 - ✗b) The span is automatically disabled.
 - ✗d) The span will inherit the `trace_id` as its `span_id`.
-
17. The `group_id` parameter in `tracing.trace()` is useful for:
- ✓b) Linking multiple, distinct traces that belong to the same overarching conversation or process.
 - ✗a) Categorizing traces by their `workflow_name`.
 - ✗c) Setting access control for traces.
 - ✗d) Defining the maximum number of spans a trace can contain.
-
18. What is the common parameter found in almost all the span creation functions that allows for immediate non-recording?
- ✓c) `disabled`
 - ✗a) `sync`
 - ✗b) `enabled`
 - ✗d) `silent`
-
19. A `SpeechGroupSpanData` would likely encapsulate:
- ✓b) A logical grouping of related speech operations, possibly representing a spoken turn.
 - ✗a) A single word generated by TTS.

- ☒ c) The configuration settings for a speech model.
 - ☒ d) An error during speech processing.
-

20. When calling `function_span`, if the input and output parameters are `None`, but the name is provided, will the span still be recorded (assuming tracing is enabled and it's started)?
- ☒ c) Yes, the span will be recorded with its name and timestamps, but without input/output data.
 - ☒ a) No, input and output are required fields for `function_span`.
 - ☒ b) Only if a parent span is explicitly provided.
 - ☒ d) It depends on the `TracingProcessor` implementation.
-

OpenAI Agents SDK - Trace MCQs

1. What is the primary role of a `Trace` object in the tracing module?
- ☒ c) To represent a complete logical workflow or end-to-end operation.
 - ☒ a) To manage individual operation details.
 - ☒ b) To process and export trace data.
 - ☒ d) To generate unique identifiers for spans.
-
2. Which of the following is an abstract property of the `Trace` class?
- ☒ c) `name`
 - ☒ a) `duration`
 - ☒ b) `start_time`
 - ☒ d) `children`
-
3. The `trace_id` property of a `Trace` is used for:
- ☒ c) Uniquely identifying and correlating all related spans within a workflow.
 - ☒ a) Sorting traces by their creation time.
 - ☒ b) Defining the type of workflow.
 - ☒ d) Enabling or disabling the trace.
-
4. When `trace.start(mark_as_current=True)` is called:
- ☒ c) The trace is set as the active trace for subsequent implicit parent-child relationships.
 - ☒ a) The trace is immediately exported.
 - ☒ b) All previously started traces are finished.
 - ☒ d) Tracing is globally disabled.
-
5. What happens when `trace.finish(reset_current=True)` is called on a trace that was previously marked as current?
- ☒ c) The trace is removed from the current tracing context.
 - ☒ a) The trace is restarted.

- ☒ b) It forces a flush of all trace processors.
 - ☒ d) An error is raised because a trace cannot be reset.
-

6. The `export()` method on the `Trace` abstract base class is intended for:

- ☒ c) **Converting the trace's data into a dictionary for serialization and external consumption.**
 - ☒ a) Printing the trace details to the console.
 - ☒ b) Storing the trace in a temporary cache.
 - ☒ d) Loading trace data from a file.
-

7. Which class is a "no-operation" (dummy) implementation of `Trace`?

- ☒ d) **`NoOpTrace`**
 - ☒ a) `TraceProvider`
 - ☒ b) `TracingProcessor`
 - ☒ c) `TraceImpl`
-

8. When would a `NoOpTrace` typically be returned by the `TraceProvider`?

- ☒ c) **When tracing is globally disabled or explicitly disabled for that specific trace.**
 - ☒ a) When an error occurs during trace creation.
 - ☒ b) When a trace successfully completes its execution.
 - ☒ d) Only in development environments.
-

9. What is the primary benefit of using `NoOpTrace` when tracing is disabled?

- ☒ c) **It provides a consistent interface, preventing the need for `if tracing_enabled:` checks in application code.**
 - ☒ a) It consumes less memory than `TraceImpl`.
 - ☒ b) It allows for different export formats.
 - ☒ d) It speeds up the `start()` and `finish()` methods.
-

10. Which class is the concrete implementation of `Trace` that actively records data and interacts with `TracingProcessors`?

- ☒ c) **`TraceImpl`**
 - ☒ a) `Trace`
 - ☒ b) `NoOpTrace`
 - ☒ d) `BaseTrace`
-

11. What action does `TraceImpl.start()` method invoke?

- ☒ b) **It records the start time and notifies `TracingProcessors` via `on_trace_start()`.**
 - ☒ a) It immediately calls `export()`.
 - ☒ c) It creates new child spans automatically.
 - ☒ d) It throws an exception if a `trace_id` is not provided.
-

12. **TraceImpl** notifies registered **TracingProcessors** at which points in its lifecycle?

- ☒ c) At the beginning (`on_trace_start`) and end (`on_trace_end`) of the trace.
 - ☐ a) Only when an error occurs.
 - ☐ b) Only upon `export()`.
 - ☐ d) Only when a new span is added to it.
-

13. If `trace.export()` on a **NoOpTrace** is called, what will it return?

- ☒ b) None
 - ☐ a) An empty dictionary.
 - ☐ c) An error.
 - ☐ d) A dictionary containing only the `trace_id`.
-

14. The `name` property of a **Trace** helps in:

- ☒ b) Providing a human-readable description for the type of workflow being traced.
 - ☐ a) Distinguishing between **TraceImpl** and **NoOpTrace**.
 - ☐ c) Determining the parent-child relationship of spans.
 - ☐ d) Setting the priority level for trace processing.
-

15. What type of object does the `export()` method on **TraceImpl** return?

- ☒ c) A `dict[str, Any]` (dictionary).
 - ☐ a) A list of strings.
 - ☐ b) A serialized JSON string.
 - ☐ d) A custom **TraceData** object.
-

16. A **Trace** is described as the "root level object" because:

- ☒ c) It is the top-most container for a logical workflow, encompassing multiple spans.
 - ☐ a) It has no parent.
 - ☐ b) It always runs first.
 - ☐ d) It is stored in the root directory.
-

17. If you create a **Trace** instance but do not call its `start()` method, what is the consequence for tracing?

- ☒ b) It will not be recorded by the tracing system, as its lifecycle events won't be triggered.
 - ☐ a) It will automatically start when the first span is added.
 - ☐ c) It will be recorded as a pending trace.
 - ☐ d) Only its `trace_id` will be recorded.
-

18. The `reset_current` parameter in `finish()` is primarily concerned with:

- ☒ c) Managing the global context of the currently active trace.
 - ☐ a) Resetting the trace's start time.
 - ☐ b) Clearing all child spans from the trace.
 - ☐ d) Re-exporting the trace data.
-

19. Why are Trace properties like `trace_id` and `name` abstract methods?
- ✓c) To enforce that concrete Trace implementations must provide these essential pieces of information.
 - ✗a) To allow for different data types.
 - ✗b) To indicate they can be changed dynamically.
 - ✗d) To make them optional for specific trace types.
-

20. If an application uses `TraceImpl` objects, which other component is essential for processing and potentially storing the trace data?
- ✓c) `TracingProcessor`
 - ✗a) `SpanData`
 - ✗b) `gen_trace_id` utility
 - ✗d) `ResponseSpanData`
-

OpenAI Agents SDK - Span MCQs

1. What does a Span fundamentally represent in the tracing module?
- ✓b) A single unit of work or an individual operation within a trace.
 - ✗a) A complete workflow.
 - ✗c) A global configuration setting.
 - ✗d) An error log.
-
2. The `Span` class is defined as `Generic[TSpanData]`. What is the purpose of `TSpanData`?
- ✓c) It allows different span types to carry specific, structured data relevant to their operation.
 - ✗a) It represents the unique ID of the span.
 - ✗b) It indicates whether the span is currently active.
 - ✗d) It defines the parent of the span.
-
3. Which of the following methods is abstract in the `Span` class?
- ✓c) `export()`
 - ✗a) `get_id()`
 - ✗b) `start()`
 - ✗d) `get_duration()`
-
4. When `span.start(mark_as_current=True)` is called:
- ✓c) The span is set as the active span in the tracing context, making it the default parent for new child spans.
 - ✗a) The span is immediately sent to all processors.
 - ✗b) The span's parent is automatically reset.
 - ✗d) Tracing is temporarily disabled for this span.
-
5. What is the primary effect of calling `span.finish()`?
- ✓b) It marks the end of the span's execution and potentially notifies processors.
 - ✗a) It restarts the span's timer.

- ☒ c) It creates a new child span.
 - ☒ d) It exports the entire trace.
-

6. Which class is a "no-operation" (dummy) implementation of Span?

- ☒ b) NoOpSpan
 - ☒ a) SpanProcessor
 - ☒ c) SpanData
 - ☒ d) SpanInterface
-

7. NoOpSpan objects are primarily used when:

- ☒ b) Tracing is disabled, to provide a consistent API without actual recording.
 - ☒ a) A span has no parent.
 - ☒ c) A span represents an error condition.
 - ☒ d) Only asynchronous operations are being traced.
-

8. What is the main advantage of NoOpSpan for developers?

- ☒ c) It simplifies application code by removing the need for `if tracing_enabled:` checks around span operations.
 - ☒ a) It provides detailed debugging information even when disabled.
 - ☒ b) It allows for dynamic enabling/disabling of individual spans at runtime.
 - ☒ d) It automatically flushes trace data to a persistent store.
-

9. Which class is the concrete implementation of Span that actively records data and interacts with TracingProcessors?

- ☒ c) SpanImpl
 - ☒ a) Span
 - ☒ b) NoOpSpan
 - ☒ d) BaseSpan
-

10. What does SpanImpl do when its `start()` method is called?

- ☒ c) It records the start time and notifies registered TracingProcessors.
 - ☒ a) It calculates the span's total duration.
 - ☒ b) It exports the span data immediately.
 - ☒ d) It recursively starts all its child spans.
-

11. If `span.finish(reset_current=True)` is called on a SpanImpl that was the current span, what is the likely outcome regarding the tracing context?

- ☒ c) The current span context will revert to its parent (or None if it had no parent).
 - ☒ a) The span will remain the current span.
 - ☒ b) The entire trace will be finished.
 - ☒ d) A new span will be automatically created.
-

12. **SpanImpl maintains a reference to its parent for what purpose?**
☒ c) **To establish the hierarchical relationship within a trace.**
☐ a) To determine if it should be recorded.
☐ b) To set its name property.
☐ d) To control its disabled status.
-
13. **What type of information does a SpanImpl object typically hold through its TSpanData instance?**
☒ c) **Specific details relevant to the operation it represents (e.g., model input/output for a generation span).**
☐ a) Global tracing configuration.
☐ b) A list of all traces in the system.
☐ d) The API key for the tracing backend.
-
14. **TracingProcessors receive notifications from SpanImpl at which points?**
☒ c) **When the span starts (`on_span_start`) and when it finishes (`on_span_end`).**
☐ a) Only when an error occurs within the span.
☐ b) Only when the `export()` method is called.
☐ d) Periodically while the span is active.
-
15. **What is the return type of the `finish()` method on the Span abstract class?**
☒ c) **None**
☐ a) `bool`
☐ b) `dict`
☐ d) `Span`
-
16. **If you create a Span instance using a helper function (like `tracing.agent_span`) but do not use a `with` statement or call `start()/finish()` manually, what will happen?**
☒ b) **The span will not be recorded by the tracing system because its lifecycle events are not triggered.**
☐ a) It will automatically start and finish when the Python interpreter exits.
☐ c) It will only record errors.
☐ d) It will implicitly become a child of the root trace.
-
17. **The `mark_as_current` parameter in `start()` is a `bool`. If set to `True` for a Span, it means:**
☒ c) **This span becomes the default parent for any subsequent spans created without an explicit parent.**
☐ a) This span is the root of a new trace.
☐ b) This span will run in a separate thread.
☐ d) This span should be exported immediately.
-
18. **Span objects are designed to be:**
☒ c) **Hierarchical and nested within a Trace.**
☐ a) Independent and unrelated.
☐ b) Only for error logging.
☐ d) Immutable after creation.
-

19. Which of the following is NOT a method or property directly defined on the Span abstract base class?
- ✓c) **export()** (*This belongs to `SpanData`, not directly to `Span`*)
 - ✗a) `start()`
 - ✗b) `finish()`
 - ✗d) `mark_as_current` (*not a method, it's a parameter in `start()`*)
-
20. When dealing with tracing, the distinction between `SpanImpl` and `NoOpSpan` is handled by the:
- ✓b) **TraceProvider**
 - ✗a) `TracingProcessor`
 - ✗c) `SpanData` types
 - ✗d) The application code itself with `if` statements.
-

OpenAI Agents SDK - Processor and Exporter MCQs

1. What is the primary role of a `TracingProcessor`?
- ✓c) **To receive and process lifecycle events (start/end) from traces and spans.**
 - ✗a) To create new traces and spans.
 - ✗b) To store trace data persistently.
 - ✗d) To generate unique IDs for tracing.
-
2. Which method on `TracingProcessor` is called when a trace begins?
- ✓b) **`on_trace_start()`**
 - ✗a) `on_span_start()`
 - ✗c) `start_trace()`
 - ✗d) `process_trace_begin()`
-
3. The `on_span_end()` method in `TracingProcessor` has a crucial recommendation: it "Should not block or raise exceptions." Why is this important?
- ✓c) **To prevent interference with the application's critical path and avoid performance degradation or crashes.**
 - ✗a) To ensure all data is immediately exported.
 - ✗b) To reduce the memory footprint of the processor.
 - ✗d) To allow for easier debugging.
-
4. When should the `shutdown()` method of a `TracingProcessor` typically be called?
- ✓c) **When the application or tracing system is stopping, for cleanup.**
 - ✗a) Every time a span finishes.
 - ✗b) When a new trace is created.
 - ✗d) After every `force_flush()` operation.
-
5. What is the purpose of the `force_flush()` method on a `TracingProcessor`?
- ✓c) **To compel the processor to immediately send any buffered or queued traces/spans.**
 - ✗a) To disable all ongoing tracing.

- ☒ b) To clear all previously exported data.
 - ☒ d) To restart the trace collection process.
-

6. What is the main responsibility of a `TracingExporter`?

- ☒ c) To send processed trace and span data to an external destination (e.g., console, backend).
 - ☒ a) To define the structure of trace and span data.
 - ☒ b) To manage the lifecycle of traces and spans.
 - ☒ d) To filter out sensitive information from traces.
-

7. The `export()` method of `TracingExporter` takes which type of argument?

- ☒ d) A `list[Trace | Span[Any]]` (a list of traces and/or spans).
 - ☒ a) A single `Trace` object.
 - ☒ b) A single `Span` object.
 - ☒ c) A dict representation of a trace or span.
-

8. Which interface is more general, dealing with events, while the other is more specialized, dealing with sending data out?

- ☒ b) `TracingProcessor` is general; `TracingExporter` is specialized.
 - ☒ a) `TracingExporter` is general; `TracingProcessor` is specialized.
 - ☒ c) Both are equally general.
 - ☒ d) Neither is considered general; they serve distinct, unrelated purposes.
-

9. A `TracingProcessor` might use a `TracingExporter` internally.

- ☒ a) True
 - ☒ b) False
-

10. If a `TracingProcessor` buffers spans for batch sending, which method would clear that buffer and send the data immediately?

- ☒ c) `force_flush()`
 - ☒ a) `shutdown()`
 - ☒ b) `on_span_end()`
 - ☒ d) `on_trace_end()`
-

11. What is passed as an argument to `on_span_start()`?

- ☒ c) The `Span[Any]` object that just started.
 - ☒ a) The `span_id` string.
 - ☒ b) The `SpanData` object.
 - ☒ d) The `Trace` object that the span belongs to.
-

12. The `on_trace_end()` method is typically where a `TracingProcessor` would get the complete data for a trace. Why?

- ☒ b) Because at this point, all child spans should have finished, and the trace's full context is available.

- ☒ a) Because traces are stateless.
 - ☒ c) Because the trace ID is only generated at the end.
 - ☒ d) To prevent resource leaks.
-

13. Which scenario would most benefit from implementing a custom `TracingProcessor`?

- ☒ b) You want to send trace data to a specific third-party observability platform.
 - ☒ a) You need to create a new type of span.
 - ☒ c) You want to define the hierarchy of spans.
 - ☒ d) You want to start and stop traces automatically.
-

14. The `TracingProcessor` methods (`on_trace_start`, `on_trace_end`, etc.) are abstract because:

- ☒ b) They define a contract that concrete processor implementations must fulfill.
 - ☒ a) They are optional methods.
 - ☒ c) They are placeholders for future functionality.
 - ☒ d) They can be overridden by users at runtime.
-

15. If a `TracingExporter` fails to send data to its backend (e.g., network error), what is generally the desired behavior?

- ☒ c) To log the error and possibly retry later without blocking the main application thread.
 - ☒ a) To crash the application.
 - ☒ b) To retry indefinitely, blocking the application.
 - ☒ d) To immediately disable all tracing.
-

16. What does `Span[Any]` signify in the method signatures of `TracingProcessor` (e.g., `on_span_start(span: Span[Any])`)?

- ☒ c) The span's generic `TSpanData` type parameter is not constrained, meaning it can be any concrete `SpanData` type.
 - ☒ a) The span can be of any type, but its data is unknown.
 - ☒ b) The span's `SpanData` is optional.
 - ☒ d) The span is only for Any type of errors.
-

17. If a `TracingProcessor` implements a queuing mechanism for spans, where would it typically add a span to the queue?

- ☒ b) In `on_span_end()`.
 - ☒ a) In `on_trace_start()`.
 - ☒ c) In `shutdown()`.
 - ☒ d) In `force_flush()`.
-

18. `TracingProcessor` and `TracingExporter` are found in the `processor_interface.py` file, indicating they are:

- ☒ c) Abstract interfaces defining contracts for custom implementations.
- ☒ a) Concrete implementations ready for direct use.
- ☒ b) Utility functions for internal use only.
- ☒ d) Configuration classes.

-
19. **Calling `trace.finish()` or `span.finish()` triggers which methods on registered `TracingProcessors`?**
✓c) **`on_trace_end()` or `on_span_end()` respectively.**
✗a) `force_flush()`
✗b) `shutdown()`
✗d) `export()`
-
20. **Is it possible for a single `Trace` to be processed by multiple `TracingProcessors`?**
✓a) **Yes, multiple processors can be registered with the `TraceProvider`.**
✗b) No, only one processor can handle a trace at a time.
✗c) Only if they implement different `TracingExporters`.
✗d) Only if the `Trace` is marked as `disabled=False`.
-

OpenAI Agents SDK - Concrete Processors & Exporters MCQs

-
1. **Which concrete `TracingExporter` is best suited for quickly viewing trace and span data during local development or debugging?**
✓b) **`ConsoleSpanExporter`**
✗a) `BackendSpanExporter`
✗c) `BatchTraceProcessor`
✗d) `default_exporter()`
-
2. **The `BackendSpanExporter` is designed to send tracing data via which protocol?**
✓c) **HTTP**
✗a) FTP
✗b) WebSocket
✗d) SMTP
-
3. **Which environment variable does `BackendSpanExporter` primarily check for authentication by default?**
✓c) **`OPENAI_API_KEY`**
✗a) `TRACING_API_KEY`
✗b) `OPENAI_AUTH_TOKEN`
✗d) `BACKEND_TRACE_KEY`
-
4. **What is the main benefit of `BatchTraceProcessor`'s use of a background thread for exporting spans?**
✓c) **It minimizes performance impact on the main application thread by offloading network I/O.**
✗a) It allows for real-time, synchronous data transfer.
✗b) It simplifies the configuration of exporters.
✗d) It guarantees immediate delivery of all spans.
-

5. If the queue size in a BatchTraceProcessor reaches max_queue_size, what might happen to new spans?
✓d) They might be dropped to prevent memory exhaustion.
✗a) They are immediately exported individually.
✗b) They cause the processor to raise an exception.
✗c) They trigger a shutdown() call.
-
6. The max_retries parameter in BackendSpanExporter is used to:
✓c) Specify the maximum number of attempts to resend a failed HTTP request.
✗a) Limit the number of spans in a batch.
✗b) Define how many times the export() method can be called.
✗d) Set the maximum number of times a trace can be processed.
-
7. What is the default endpoint for BackendSpanExporter?
✓c) <https://api.openai.com/v1/traces/ingest>
✗a) <http://localhost:8080/traces>
✗b) <https://api.openai.com/v1/traces/export>
✗d) <http://tracing.example.com/api>
-
8. The schedule_delay parameter in BatchTraceProcessor controls:
✓b) The interval (in seconds) at which the background thread checks the queue for spans to export.
✗a) How long the main application thread pauses before generating new spans.
✗c) The delay before retrying a failed export.
✗d) The time before the processor shuts down automatically.
-
9. Which method on BackendSpanExporter should be called when the application is exiting to ensure all resources are properly released?
✓a) shutdown()
✗b) force_flush()
✗c) close()
✗d) disconnect()
-
10. What does the export_trigger_ratio in BatchTraceProcessor help to achieve?
✓b) It defines a threshold (based on queue fullness) to trigger an immediate export, even if the scheduled delay hasn't passed.
✗a) It sets the maximum time a span can stay in the queue.
✗c) It determines the proportion of spans that should be dropped.
✗d) It controls the growth of the exponential backoff delay.
-
11. What is the primary advantage of BatchTraceProcessor over simply calling exporter.export() for every on_span_end event?
✓c) Reduced network overhead and improved application performance by batching requests.
✗a) Simpler configuration.
✗b) Guaranteed real-time delivery.
✗d) Automatic error recovery without retries.

12. The `default_exporter()` function returns an instance of which class?

- ☒ b) **BackendSpanExporter**
 - ☐ a) ConsoleSpanExporter
 - ☐ c) BatchTraceProcessor
 - ☐ d) TracingExporter (abstract class)
-

13. The `default_processor()` function returns an instance of which class?

- ☒ c) **BatchTraceProcessor**
 - ☐ a) ConsoleSpanExporter
 - ☐ b) BackendSpanExporter
 - ☐ d) TracingProcessor (abstract class)
-

14. If you want to configure the `BackendSpanExporter` with a specific OpenAI project ID, which parameter would you use in its constructor?

- ☒ c) **project**
 - ☐ a) organization
 - ☐ b) api_key
 - ☐ d) endpoint
-

15. `BackendSpanExporter` implements retry logic for network requests. This feature primarily contributes to:

- ☒ c) **Improved reliability and fault tolerance.**
 - ☐ a) Faster data transfer.
 - ☐ b) Lower CPU utilization.
 - ☐ d) Reduced memory consumption.
-

16. What kind of objects does `ConsoleSpanExporter` typically receive in its `export()` method?

- ☒ c) **Trace and Span objects (which it then converts to dictionaries for printing).**
 - ☐ a) Already formatted JSON strings.
 - ☐ b) Raw binary data.
 - ☐ d) File paths.
-

17. If a `BatchTraceProcessor` receives many spans rapidly, and `export_trigger_ratio` is set to 0.9, what happens when the queue reaches 90% of `max_queue_size`?

- ☒ c) **An export attempt is triggered immediately.**
 - ☐ a) Spans are immediately dropped.
 - ☐ b) The processor halts.
 - ☐ d) The `schedule_delay` is increased.
-

18. What is the type of object that `BatchTraceProcessor` takes as its primary argument in its constructor?

- ☒ b) **TracingExporter**
- ☐ a) TracingProcessor

- ☒ c) Span
 - ☒ d) Trace
-

19. Calling `force_flush()` on a `BatchTraceProcessor` causes:

- ☒ c) Any queued spans/traces to be immediately sent through its exporter.
 - ☒ a) The background thread to stop.
 - ☒ b) All previously exported data to be deleted.
 - ☒ d) The `schedule_delay` to be reset.
-

20. When setting up tracing for a production application that sends data to an OpenAI backend, which combination of components is generally recommended?

- ☒ c) `default_processor()` (which uses `BatchTraceProcessor` and `default_exporter()`).
 - ☒ a) `ConsoleSpanExporter` directly in application code.
 - ☒ b) `BackendSpanExporter` with manual `export()` calls.
 - ☒ d) Custom `TracingProcessor` and `TracingExporter` from scratch.
-

OpenAI Agents SDK - Scope MCQs

1. What is the primary function of the `Scope` module in the tracing system?

- ☒ b) To manage the currently active span and trace in the execution context.
 - ☒ a) To manage API keys for backend exporters.
 - ☒ c) To define abstract interfaces for processors.
 - ☒ d) To export trace data to external systems.
-

2. Why is the concept of a "current" span/trace, managed by `Scope`, important for tracing?

- ☒ c) It enables the automatic establishment of parent-child relationships for newly created spans.
 - ☒ a) It allows for faster processing of trace data.
 - ☒ b) It determines the output format of exported traces.
 - ☒ d) It defines the maximum depth of a trace.
-

3. If you create a new span without explicitly providing a parent argument, how does the tracing system typically determine its parent?

- ☒ c) It uses the currently active span or trace managed by `Scope` as the parent.
 - ☒ a) It assigns a random parent from existing traces.
 - ☒ b) It creates a new root trace for the span.
 - ☒ d) The span will not be associated with any parent.
-

4. Which Python mechanism is commonly used by `Scope` (or similar context management systems) to store the active trace/span for the current execution flow?

- ☒ c) Thread-local storage or `contextvars`
- ☒ a) Global variables

- ☒ b) Class attributes
 - ☒ d) Database sessions
-

5. The `get_current_trace()` and `get_current_span()` utility functions directly rely on which module to retrieve the active tracing context?

- ☒ c) Scope (implicitly, via the `TraceProvider` which uses it)
 - ☒ a) `TracingProcessor`
 - ☒ b) `TracingExporter`
 - ☒ d) `SpanData`
-

6. Without the functionality provided by `Scope`, how would you typically have to link spans to their parents?

- ☒ c) You would need to explicitly pass the parent argument to every span creation function.
 - ☒ a) All spans would automatically be top-level traces.
 - ☒ b) Spans would be linked by their name property.
 - ☒ d) Spans would only link to the `workflow_name`.
-

7. The management of the "current" trace and span helps ensure that traces are:

- ☒ c) Well-formed and their spans are correctly nested hierarchically.
 - ☒ a) Always disabled for production.
 - ☒ b) Limited to a fixed number of spans.
 - ☒ d) Sorted alphabetically by their name.
-

8. When a `TraceImpl` or `SpanImpl` calls `start(mark_as_current=True)`, it interacts with `Scope` to:

- ☒ c) Set itself as the active trace/span in the current context.
 - ☒ a) Begin sending data to an exporter.
 - ☒ b) Calculate its duration.
 - ☒ d) Register a new processor.
-

9. When a `TraceImpl` or `SpanImpl` calls `finish(reset_current=True)`, it interacts with `Scope` to:

- ☒ b) Remove itself as the active trace/span and potentially revert to its parent's context.
 - ☒ a) Export its data.
 - ☒ c) Mark itself as an error.
 - ☒ d) Increase the `max_queue_size`.
-

10. Is `Scope` typically something an end-user developer directly interacts with by calling its methods?

- ☒ b) No, it's usually managed internally by the tracing library (e.g., via `TraceImpl`, `SpanImpl`, and context managers).
 - ☒ a) Yes, it's frequently used for manual context switching.
 - ☒ c) Only during debugging.
 - ☒ d) Only when defining custom span types.
-

11. Which of the following best describes Scope's role in making tracing "ergonomic"?
- ✓c) It simplifies instrumentation by automatically propagating the parent context, reducing boilerplate.
 - ✗a) It provides default values for all span parameters.
 - ✗b) It optimizes the performance of network calls.
 - ✗d) It filters out unnecessary trace data.
-
12. If Scope fails to manage the current context correctly, what might be a visible symptom in your traced data?
- ✓b) Spans appearing as root spans when they should be children, or incorrect nesting.
 - ✗a) Traces being too short.
 - ✗c) Exported data being corrupted.
 - ✗d) Processors failing to shut down.
-
13. The concept of Scope ensures that the tracing context is:
- ✓c) Specific to the current execution flow or thread.
 - ✗a) Globally immutable.
 - ✗b) Shared across all processes.
 - ✗d) Only active when a ConsoleSpanExporter is used.
-
14. Scope contributes to ensuring that a Trace is a well-formed:
- ✓c) Tree (or directed acyclic graph - DAG)
 - ✗a) List
 - ✗b) Flat sequence
 - ✗d) Queue
-
15. If a NoOpTrace or NoOpSpan is created, does Scope still manage its "current" status if mark_as_current=True is used?
- ✓b) No, NoOpTrace and NoOpSpan do not interact with the actual context management provided by Scope for recording.
 - ✗a) Yes, but it will still perform no-op operations.
 - ✗c) Only if explicitly configured.
 - ✗d) It depends on the Python version.
-
16. Which parameter directly influences Scope's state when calling trace.start() or span.start()?
- ✓c) mark_as_current
 - ✗a) workflow_name
 - ✗b) metadata
 - ✗d) disabled
-
17. The primary benefit of Scope is related to:
- ✓c) Context propagation.
 - ✗a) Data serialization.
 - ✗b) Asynchronous operations.
 - ✗d) Error handling.

-
18. If a Trace is started with `mark_as_current=True`, and then a Span is started within that trace without an explicit parent, what object does Scope ensure the span is associated with?
- ✓c) The currently active Trace (if no Span is current).
 - ✗a) A new, independent trace.
 - ✗b) The global root of all traces.
 - ✗d) A randomly selected parent.
-
19. The Scope mechanism helps in correctly attributing spans to their parent operations, which is vital for:
- ✓c) Visualizing the flow of execution and understanding dependencies.
 - ✗a) Reducing API call latency.
 - ✗b) Compressing trace data.
 - ✗d) Encrypting sensitive information.
-
20. The abstract base class Trace and Span define methods (`start`, `finish`) that, when implemented by `TraceImpl` and `SpanImpl`, interact with Scope to:
- ✓c) Update the current active tracing context.
 - ✗a) Validate input parameters.
 - ✗b) Choose the correct exporter.
 - ✗d) Manage concurrent access to resources.
-

OpenAI Agents SDK - Trace Provider Setup MCQs

1. What is the main purpose of `set_trace_provider()`?
- ✓c) To globally configure the central `TraceProvider` instance for the tracing utilities.
 - ✗a) To create a new trace.
 - ✗b) To get the currently active span.
 - ✗d) To export trace data immediately.
-
2. The `TraceProvider` is described as the "central, overarching component" responsible for:
- ✓c) Creating Trace and Span objects, and managing the tracing context.
 - ✗a) Only printing traces to the console.
 - ✗b) Only sending data to a backend.
 - ✗d) Implementing `on_trace_start` and `on_span_end`.
-
3. If `set_trace_provider()` is called, which subsequent tracing operation will NOT use the newly set provider?
- ✓d) All of the above will use the newly set provider.
 - ✗a) `tracing.trace()`
 - ✗b) `tracing.agent_span()`
 - ✗c) `tracing.get_current_trace()`
-
4. Why is it beneficial to be able to `set_trace_provider()`?
- ✓c) It allows for flexible configuration and dependency injection of tracing behavior.
-

- ☒ a) It makes traces run faster.
 - ☒ b) It reduces the memory footprint of traces.
 - ☒ d) It automatically handles all network retries.
-

5. What type of object does `set_trace_provider()` expect as its provider argument?

- ☒ c) **TraceProvider**
 - ☒ a) TracingProcessor
 - ☒ b) TracingExporter
 - ☒ d) SpanData
-

6. What does `get_trace_provider()` return?

- ☒ c) **The currently configured global TraceProvider instance.**
 - ☒ a) A new TraceProvider instance every time.
 - ☒ b) The default NoOpTraceProvider.
 - ☒ d) A list of all active traces.
-

7. A common use case for `get_trace_provider()` is:

- ☒ c) **To inspect or potentially dynamically modify the behavior of the active provider (e.g., add/remove processors).**
 - ☒ a) To directly create a new Trace without using `tracing.trace()`.
 - ☒ b) To clear all existing trace data.
 - ☒ d) To force a flush of all pending exports.
-

8. If `set_trace_provider()` has not been explicitly called, what kind of TraceProvider might `get_trace_provider()` return by default in a typical tracing setup?

- ☒ b) **A default TraceProvider (e.g., TraceProviderImpl) or a NoOpTraceProvider if tracing is off.**
 - ☒ a) A ConsoleSpanExporter.
 - ☒ c) An error because no provider is set.
 - ☒ d) None.
-

9. The functions in `setup.py` emphasize the importance of having a single, global TraceProvider.

- ☒ a) **True**
 - ☒ b) False
-

10. What is the primary benefit of `set_trace_provider()` in terms of testing your application?

- ☒ c) **It allows you to inject mock or disabled TraceProvider instances for isolated testing.**
 - ☒ a) It automatically generates test cases.
 - ☒ b) It highlights performance bottlenecks.
 - ☒ d) It provides a built-in testing framework.
-

11. Where in an application's lifecycle would `set_trace_provider()` typically be called?

- ☒ b) **During application initialization or startup.**

- ☒ a) Inside every function that creates a span.
 - ☒ c) Immediately before the application shuts down.
 - ☒ d) Only when an error occurs.
-

12. If `set_trace_provider(NoOpTraceProvider())` is called, what will be the effect on `tracing.trace()` and `tracing.agent_span()` calls?

- ☒ c) They will return `NoOpTrace` and `NoOpSpan` objects, effectively disabling tracing.
 - ☒ a) They will continue to record data but won't export it.
 - ☒ b) They will raise an exception.
 - ☒ d) They will buffer all traces and spans indefinitely.
-

13. `get_trace_provider()` is used to access the provider, not to create it.

- ☒ a) True
 - ☒ b) False
-

14. Why is managing the `TraceProvider` globally via `set_trace_provider` and `get_trace_provider` preferable to passing it around as an argument to every function?

- ☒ b) It simplifies application code by providing a globally accessible tracing context without explicit argument passing.
 - ☒ a) It makes the code more complex.
 - ☒ c) It only works for single-threaded applications.
 - ☒ d) It reduces the number of `TracingProcessors`.
-

15. If you have a custom `TraceProvider` implementation that includes unique debugging methods, how would you access those methods after setting your custom provider?

- ☒ c) You can retrieve it using `get_trace_provider()` and then call your custom methods.
 - ☒ a) You cannot, only standard methods are accessible.
 - ☒ b) You must cast the result of `get_trace_provider()` to your custom type.
 - ☒ d) Custom methods are automatically called.
-

16. What kind of parameters does `set_trace_provider()` take?

- ☒ c) An instance of a class that implements `TraceProvider`.
 - ☒ a) A string representing the provider name.
 - ☒ b) A dictionary of configuration settings.
 - ☒ d) A boolean to enable/disable tracing.
-

17. If a `TraceProvider` is not set, attempts to create traces or spans using the tracing utilities might:

- ☒ b) Result in `NoOpTrace` or `NoOpSpan` objects, or potentially an error depending on the default provider's behavior.
 - ☒ a) Automatically set up a default, fully functional provider.
 - ☒ c) Crash the application immediately.
 - ☒ d) Export data to a temporary file.
-

18. The `setup.py` module essentially provides the API for:
- ✓c) Bootstrapping the tracing system within an application.
 - ✗a) Defining new `SpanData` structures.
 - ✗b) Implementing `TracingProcessor` logic.
 - ✗d) Converting trace data to JSON.
-

19. Calling `set_trace_provider()` multiple times with different providers during an application's runtime would:
- ✓b) Change the active `TraceProvider` to the most recently set one, potentially altering tracing behavior mid-execution.
 - ✗a) Be an error, as it can only be set once.
 - ✗c) Cause all previous providers to flush their data.
 - ✗d) Have no effect after the first call.
-

20. Which of the following is NOT a direct responsibility of the `set_trace_provider()` or `get_trace_provider()` functions themselves?
- ✓c) Performing the actual collection or export of trace data.
 - ✗a) Allowing global configuration of tracing.
 - ✗b) Providing access to the current `TraceProvider`.
 - ✗d) Enabling dependency injection for the tracing core.
-

OpenAI Agents SDK - SpanData MCQs

1. What is the primary role of `SpanData` in the tracing module?

- ✗a) To manage the parent-child relationships of spans.
 - ✗b) To define the start and end times of a span.
 - ✓c) To represent the specific payload or contextual attributes of a particular type of span.
 - ✗d) To export traces to a backend.
-

2. Which of the following is an abstract property of the `SpanData` class?

- ✗a) `duration`
- ✗b) `parent_id`
- ✓c) `type`

☐ d) timestamp

3. The `export()` method on `SpanData` is responsible for:

- ☐ a) Sending the data to a `TracingExporter`.
 - ☒ b) **Serializing the span's specific data into a dictionary format.**
 - ☐ c) Printing the span data to the console.
 - ☐ d) Calculating the cost of the span.
-

4. Which `SpanData` implementation would typically include input messages, output messages, model, and usage?

- ☐ a) `FunctionSpanData`
 - ☐ b) `ResponseSpanData`
 - ☒ c) **`GenerationSpanData`**
 - ☐ d) `AgentSpanData`
-

5. If you want to track the arguments passed to and results returned from a custom tool your agent calls, which `SpanData` type would be most appropriate?

- ☐ a) `AgentSpanData`
 - ☒ b) **`FunctionSpanData`**
 - ☐ c) `CustomSpanData`
 - ☐ d) `MCPListToolsSpanData`
-

6. Which `SpanData` type is designed for flexible, user-defined operations where you can attach arbitrary key-value pairs?

- ☐ a) `AgentSpanData`
 - ☐ b) `FunctionSpanData`
 - ☒ c) **`CustomSpanData`**
 - ☐ d) `SpeechGroupSpanData`
-

7. A HandoffSpanData specifically captures information about:

- ☐ a) Data transfer between a model and a database.
 - ☒ b) **The transfer of control or responsibility between different agents.**
 - ☐ c) The initial user prompt.
 - ☐ d) The final response generated by the agent.
-

8. What key piece of information does GuardrailSpanData typically record?

- ☐ a) The LLM's confidence score.
 - ☐ b) The duration of the guardrail check.
 - ☒ c) **The name of the guardrail and its triggered status (boolean).**
 - ☐ d) The policy document used by the guardrail.
-

9. TranscriptionSpanData and SpeechSpanData are specifically for operations related to:

- ☐ a) Image processing.
 - ☐ b) Data parsing.
 - ☒ c) **Speech-to-text and text-to-speech.**
 - ☐ d) Database queries.
-

10. Which SpanData type would be used to represent the overall activity of an agent, potentially encompassing multiple internal steps?

- ☒ a) **AgentSpanData**
 - ☐ b) CustomSpanData
 - ☐ c) ResponseSpanData
 - ☐ d) SpeechGroupSpanData
-

11. The ResponseSpanData object primarily focuses on:

- ☐ a) The agent's internal reasoning steps.
- ☐ b) Tool calls made by the agent.
- ☒ c) **The final output or message generated by the agent for the user/system.**

☐ d) The configuration of the agent.

12. The type property of a SpanData object allows consumers of trace data to:

- ☐ a) Determine the original source file of the span.
 - ☒ b) **Understand the specific kind or category of operation represented by the span.**
 - ☐ c) Authenticate the span's origin.
 - ☐ d) Calculate the exact duration of the span.
-

13. MCPListToolsSpanData would be used when an agent interacts with a system to:

- ☐ a) Execute a specific tool.
 - ☐ b) Monitor MCP server health.
 - ☒ c) **Retrieve a list of available tools from the MCP server.**
 - ☐ d) Configure the MCP server.
-

14. What distinguishes SpanData from Span itself?

- ☐ a) SpanData manages lifecycle; Span carries payload.
 - ☒ b) **Span manages lifecycle and hierarchy; SpanData carries the specific, contextual payload.**
 - ☐ c) SpanData is abstract; Span is concrete.
 - ☐ d) SpanData is only for errors; Span is for successes.
-

15. If a Span is created using tracing.generation_span(), what type of SpanData object will it typically hold internally?

- ☐ a) AgentSpanData
 - ☐ b) FunctionSpanData
 - ☒ c) **GenerationSpanData**
 - ☐ d) CustomSpanData
-

16. What is the benefit of having a generic Span[TSpanData] for the Span class, instead of just Span[Any]?

- ☐ a) It makes the code run faster.

- ☒ b) It allows for different export formats.
 - ☒ c) **It provides type-safety, ensuring that a Span instance is associated with the correct, structured data for its type.**
 - ☒ d) It disables tracing if the data type is incorrect.
-

17. The `SpeechGroupSpanData` suggests a purpose of:

- ☒ a) Recording individual spoken words.
 - ☒ b) Storing audio file paths.
 - ☒ c) **Grouping related speech operations, possibly representing a complete spoken turn or interaction segment.**
 - ☒ d) Translating speech into different languages.
-

18. All concrete `SpanData` implementations inherit from `SpanData` and must implement its abstract methods/properties.

- ☒ a) **True**
 - ☒ b) False
-

19. Which method is consistently defined across all `SpanData` implementations for serialization purposes?

- ☒ a) `to_json()`
 - ☒ b) **`export()`**
 - ☒ c) `serialize()`
 - ☒ d) `get_data()`
-

20. When would you typically instantiate a `SpanData` object directly in your application code?

- ☒ a) Every time you call `span.start()`.
 - ☒ b) **You generally don't; you use tracing helper functions (e.g., `tracing.agent_span`) which create the appropriate `SpanData` internally.**
 - ☒ c) Only when using `ConsoleSpanExporter`.
 - ☒ d) When retrieving data from a backend.
-

OpenAI Agents SDK - Utility Functions MCQs

1. What is the primary purpose of the `time_iso()` function?

- ☐ a) To calculate the duration of a span.
 - ☐ b) To set the system's time zone.
 - ☒ c) **To return the current time in ISO 8601 format.**
 - ☐ d) To measure CPU usage.
-

2. Why is using ISO 8601 format important for timestamps in tracing?

- ☐ a) It makes the timestamps shorter.
 - ☐ b) It encrypts the time information.
 - ☒ c) **It ensures consistent ordering, accurate duration calculation, and interoperability with other systems.**
 - ☐ d) It only records local time.
-

3. The `gen_trace_id()` function generates an ID that is primarily used for:

- ☐ a) Identifying individual operations within a function.
 - ☐ b) Linking a span to its direct parent.
 - ☒ c) **Correlating all spans that belong to a single logical workflow.**
 - ☐ d) Categorizing different types of spans.
-

4. What characteristic is crucial for a `trace_id` to ensure proper correlation across distributed services?

- ☐ a) It must be a short integer.
 - ☐ b) It must be human-readable.
 - ☒ c) **It must be globally unique (or have a very high probability of uniqueness).**
 - ☐ d) It must be sorted alphabetically.
-

5. The `gen_span_id()` function generates an ID that is primarily used for:

- ☐ a) Linking multiple traces together.
- ☒ b) **Uniquely identifying an individual operation (unit of work) within a trace.**

☐ c) Specifying the type of data within a span.

☐ d) Setting the name of a span.

6. Which function is used to link multiple distinct traces that belong to a broader, ongoing conversation or process?

☐ a) `gen_trace_id()`

☐ b) `gen_span_id()`

☒ c) **`gen_group_id()`**

☐ d) `time_iso()`

7. A `group_id` differs from a `trace_id` in that a `group_id` links:

☐ a) Spans to other spans.

☐ b) Spans to their trace.

☒ c) **Traces to other traces within a larger logical conversation.**

☐ d) Workflows to external systems.

8. If `time_iso()` returns `2025-06-24T06:39:01.123456Z`, what does the `z` at the end signify?

☐ a) Zone-specific time.

☐ b) Zero milliseconds.

☒ c) **Zulu time (UTC).**

☐ d) Daylight Saving Time.

9. Which of these functions typically produces a UUID or similar high-entropy random string?

☐ a) `time_iso()`

☐ b) `ConsoleSpanExporter`

☒ c) **`gen_trace_id()` and `gen_span_id()`**

☐ d) `set_trace_provider()`

10. What is a primary use case for `gen_group_id()` in an AI assistant application?

- ☐ a) To uniquely identify each LLM call.
 - ☐ b) To identify individual tool invocations.
 - ☒ c) **To link all traces from a single, multi-turn conversational session.**
 - ☐ d) To generate names for new agents.
-

11. The utility functions in `util.py` are essential because they provide the necessary:

- ☐ a) Configuration settings
 - ☐ b) Network communication
 - ☒ c) **Unique identifiers and temporal context**
 - ☐ d) Data compression algorithms
-

12. The output of `time_iso()` is a:

- ☐ a) datetime object
 - ☐ b) float representing milliseconds since epoch
 - ☒ c) **str**
 - ☐ d) int representing Unix timestamp
-

13. If `gen_span_id()` is called twice within the same trace, what is guaranteed about the two IDs generated?

- ☐ a) They will be identical.
 - ☐ b) They will be sequential numbers.
 - ☒ c) **They will be unique.**
 - ☐ d) They will share the same prefix.
-

14. `gen_trace_id()`, `gen_span_id()`, and `gen_group_id()` are designed to create IDs that are:

- ☐ a) Short and easy to remember.
- ☐ b) Always sequential.
- ☒ c) **Highly unlikely to collide (be duplicate).**

☐ d) Directly convertible to integers for arithmetic.

15. When a Span finishes, its duration is calculated using its start and end timestamps. These timestamps are typically formatted using which utility function's output?

☐ a) `gen_trace_id()`

☒ b) `time_iso()`

☐ c) `gen_span_id()`

☐ d) `gen_group_id()`

16. What would be a potential problem if `gen_trace_id()` did not produce globally unique IDs?

☐ a) Spans would not be able to link to their parent.

☐ b) TracingProcessors would fail to export data.

☒ c) It would be difficult to correctly correlate and view all operations belonging to a single workflow, especially in distributed systems.

☐ d) The BackendSpanExporter would use the wrong endpoint.

17. The `util.py` module contains:

☐ a) Concrete TraceProvider implementations

☐ b) Abstract interfaces for tracing

☒ c) Helper functions for ID generation and time formatting

☐ d) Classes for managing tracing context

18. If you were building a custom trace visualization tool, what would be the most important piece of information to properly reconstruct the trace hierarchy?

☐ a) Only `trace_id`

☐ b) Only `span_id`

☒ c) `trace_id`, `span_id`, and a `parent_span_id` (or direct parent reference)

☐ d) `group_id` and `time_iso()`

19. Which function would you *not* expect to directly use `time_iso()` in its internal logic?

- ☐ a) `TraceImpl.start()`
 - ☐ b) `SpanImpl.finish()`
 - ☒ c) `gen_span_id()` (ID generation is separate from timestamping)
 - ☐ d) `BatchTraceProcessor` (for scheduling or flushing)
-

20. The functions in `util.py` are generally meant to be:

- ☐ a) Modified by end-users
 - ☐ b) Used only for internal testing
 - ☒ c) Called by the tracing library's core components (e.g., `TraceImpl`, `SpanImpl`) to ensure consistency
 - ☐ d) Passed as arguments to `TracingExporters`
-

Voice

OpenAI Agents SDK - VoicePipeline MCQs

1. What are the three main steps orchestrated by the VoicePipeline?

- ☐ a) Text-to-Text, Voice Recognition, Audio Saving
 - ☐ b) Audio Playback, Image Generation, Text Analysis
 - ☒ c) **Transcribe audio to text, Run workflow (text-to-text), Convert text to streaming audio**
 - ☐ d) Record Audio, Process Video, Send Email
-

2. The VoicePipeline is described as "opinionated." What does this imply?

- ☐ a) It allows for complete customization of every internal component
 - ☐ b) It provides recommendations but no concrete implementation
 - ☒ c) **It offers a predefined, common structure suitable for typical voice agent use cases**
 - ☐ d) It only works with specific voice assistants
-

3. Which component is explicitly required when initializing a VoicePipeline?

- ☐ a) stt_model
 - ☐ b) tts_model
 - ☒ c) **workflow**
 - ☐ d) config
-

4. The `workflow` parameter in `VoicePipeline`'s constructor must be an instance of a class inheriting from:

- ☐ a) AudioInput
 - ☐ b) STTModel
 - ☒ c) **VoiceWorkflowBase**
 - ☐ d) StreamedAudioResult
-

5. What happens if `stt_model` is not provided when creating a `VoicePipeline`?

- ☐ a) The pipeline will only work with text input
 - ☒ b) **An error will be raised**
 - ☐ c) A default OpenAI Speech-to-Text model will be used
 - ☐ d) The pipeline will automatically try to detect the model
-

6. The `run()` method of `VoicePipeline` is an async method. Why is this significant?

- ☐ a) It means the method runs on a separate CPU core
 - ☒ b) **It indicates support for non-blocking operations, essential for real-time voice interactions**
 - ☐ c) It signifies that it only processes static audio files
 - ☐ d) It requires a specific file system for operation
-

7. Which type of audio input allows for continuous feeding of audio data to the `VoicePipeline`?

- ☐ a) AudioInput
 - ☒ b) **StreamedAudioInput**
 - ☐ c) AudioBuffer
 - ☐ d) StaticAudio
-

8. What is the return type of the `VoicePipeline.run()` method?

- ☐ a) str (the transcribed text)
 - ☐ b) AudioInput (the processed audio)
 - ☒ c) **StreamedAudioResult (for playing audio as it's generated)**
 - ☐ d) dict (containing pipeline statistics)
-

9. The ability to stream audio output via `StreamedAudioResult` primarily helps to:

- ☐ a) Reduce network bandwidth usage
 - ☒ b) **Improve the perceived responsiveness of the agent by playing audio as it's synthesized**
 - ☐ c) Ensure perfect audio fidelity
 - ☐ d) Enable offline processing
-

10. What is the role of `VoiceWorkflowBase` within the `VoicePipeline`?

- ☐ a) To perform speech recognition
 - ☐ b) To convert text to audio
 - ☒ c) **To act as the core logic or "brain" of the agent, processing text and generating text responses**
 - ☐ d) To manage audio input/output devices
-

11. The `VoicePipelineConfig` parameter allows for configuration of:

- ☐ a) The specific workflow to use
 - ☐ b) The `api_key` for models
 - ☒ c) **Technical parameters like audio sample rates or streaming chunk sizes**
 - ☐ d) The name of the voice agent
-

12. If a `VoicePipeline` is created without specifying `tts_model`, what will be used?

- ☐ a) No text-to-speech conversion will occur
 - ☐ b) An error will be thrown during `run()`
 - ☒ c) **A default OpenAI Text-to-Speech model**
 - ☐ d) A console-based text output
-

13. `AudioInput` is used for:

- ☐ a) Live microphone input
 - ☒ b) **A single, static buffer of audio data**
 - ☐ c) A continuous stream of audio
 - ☐ d) Textual input converted to audio
-

14. The `VoicePipeline` directly encapsulates the logic for:

- ☐ a) Database interactions
 - ☐ b) Sending emails
 - ☒ c) **The complete audio-to-text-to-audio conversational loop**
 - ☐ d) User authentication
-

15. What is the benefit of an "opinionated" pipeline design for developers?

- ☐ a) It forces developers to implement all components from scratch
 - ☒ b) **It provides a quick and robust starting point for common use cases, reducing initial setup complexity**
 - ☐ c) It makes it harder to integrate with external services
 - ☐ d) It removes the possibility of using different STT/TTS models
-

16. Can `stt_model` and `tts_model` be provided as simple string names (e.g., "whisper") instead of explicit model objects?

- ☒ a) **Yes**
 - ☐ b) No, they must be full model objects
-

17. The `VoicePipeline` is found in `src/agents/voice/pipeline.py`, indicating its focus on:

- ☐ a) General purpose AI agents
 - ☐ b) Multi-modal agents combining text and images
 - ☒ c) **Agents primarily interacting via spoken language**
 - ☐ d) Backend data processing
-

18. If a `VoiceWorkflowBase` produces multiple text responses sequentially, how does `VoicePipeline` handle them for output?

- ☐ a) It concatenates them into a single response
 - ☐ b) It only uses the first response
 - ☐ c) It discards all but the last response
 - ☒ d) **It converts each text response into streaming audio, sending them in sequence**
-

19. Which of these is NOT a direct function of the `VoicePipeline` itself, but rather part of the workflow it runs?

- ☐ a) Transcribing audio
 - ☐ b) Converting text to speech
 - ☒ c) **Deciding what to say based on the transcribed text (e.g., calling tools, reasoning)**
 - ☐ d) Managing audio input streams
-

20. The `VoicePipeline`'s design aims to facilitate:

- ☐ a) Batch processing of large text documents
 - ☐ b) Offline training of AI models
 - ☒ c) **Building real-time, interactive conversational voice agents**
 - ☐ d) Scientific data analysis
-

• OpenAI Agents SDK - Voice Workflow MCQs

1. What is the main purpose of `VoiceWorkflowBase`?

- ☐ a) To manage audio input/output devices
 - ☐ b) To convert audio to text
 - ☒ c) **To define the interface for the core logic ("brain") of a voice agent**
 - ☐ d) To handle network communication for models
-

2. The `run` method of `VoiceWorkflowBase` receives transcription: `str` as input. What does this transcription represent?

- ☐ a) The agent's generated text response
 - ☐ b) A configuration setting for the workflow
 - ☒ c) **The text converted from the user's spoken audio**
 - ☐ d) A unique ID for the workflow
-

3. The `run` method of `VoiceWorkflowBase` returns an `AsyncIterator[str]`. What is the significance of `AsyncIterator` and `str`?

- ☐ a) It means the workflow returns a single, complete string asynchronously
 - ☐ b) It means the workflow processes a list of strings in a batch
 - ☒ c) **It allows the workflow to stream multiple textual responses over time, enabling incremental text-to-speech**
 - ☐ d) It indicates that the workflow only handles error messages
-

4. What is the recommended approach for implementing the `run` method of a custom `VoiceWorkflowBase`?

- ☐ a) Directly call a Text-to-Speech model
 - ☐ b) Perform complex database operations
 - ☒ c) **Call `Runner.run_streamed()` with an Agent and yield text events from its result**
 - ☐ d) Save the transcription to a file
-

5. What is the primary function of `VoiceWorkflowHelper.stream_text_from()`?

- ☐ a) To transcribe audio from a `RunResultStreaming`
 - ☒ b) **To extract and yield only the text events from a `RunResultStreaming` object**
 - ☐ c) To convert text events into audio
 - ☐ d) To manage the history of an agent
-

6. `SingleAgentVoiceWorkflow` is a concrete implementation of which abstract base class?

- ☐ a) `VoicePipeline`
 - ☒ b) **`VoiceWorkflowBase`**
 - ☐ c) `TracingProcessor`
 - ☐ d) `Agent`
-

7. For which scenario is `SingleAgentVoiceWorkflow` most suitable?

- ☐ a) Workflows requiring multiple `Runner` calls
 - ☒ b) **Simple voice agents that directly wrap a single `Agent` instance**
 - ☐ c) Workflows with complex custom message history logic
 - ☐ d) Workflows that only perform speech-to-text
-

8. What key conversational feature does `SingleAgentVoiceWorkflow` provide out-of-the-box?

- ☐ a) Advanced error detection
 - ☐ b) Real-time language translation
 - ☒ c) **Automatic addition of transcription and agent result to the agent's input history**
 - ☐ d) Dynamic model selection
-

9. The `SingleAgentWorkflowCallbacks.on_run()` method is called:

- ☐ a) After the agent has finished its response
 - ☐ b) When the `VoicePipeline` starts
 - ☒ c) **When the `SingleAgentVoiceWorkflow` begins processing a new transcription**
 - ☐ d) Periodically during the workflow's execution
-

10. If you need to implement a voice agent that involves dynamic branching logic based on the user's input, which class should you directly subclass?

- ☐ a) `SingleAgentVoiceWorkflow`
 - ☐ b) `VoiceWorkflowHelper`
 - ☒ c) **`VoiceWorkflowBase`**
 - ☐ d) `Agent`
-

11. What is the type of the main component passed to the constructor of `SingleAgentVoiceWorkflow`?

- ☐ a) `VoicePipelineConfig`
 - ☒ b) **`Agent[Any]`**
 - ☐ c) `AsyncIterator[str]`
 - ☐ d) `SingleAgentWorkflowCallbacks`
-

12. The `AsyncIterator` return type of `VoiceWorkflowBase.run()` directly supports which feature of `VoicePipeline`?

- ☐ a) Background processing of audio
 - ☐ b) Synchronous execution
 - ☒ c) **Streaming audio output (Text-to-Speech)**
 - ☐ d) Batch processing of transcriptions
-

13. `VoiceWorkflowHelper.stream_text_from()` is a classmethod because:

- ☐ a) It needs access to instance-specific data
 - ☒ b) **It operates on a `RunResultStreaming` object and doesn't require an instance of `VoiceWorkflowHelper` itself**
 - ☐ c) It needs to be overridden by subclasses
 - ☐ d) It handles global configuration
-

14. What does the term "workflow" signify in the context of `VoiceWorkflowBase`?

- ☐ a) A predefined sequence of audio files
 - ☐ b) The process of training a new AI model
 - ☒ c) **Any code logic that receives a transcription and yields text to be spoken**
 - ☐ d) A data storage mechanism
-

15. `SingleAgentVoiceWorkflow` is described as a "simple" workflow. This implies it might lack advanced features like:

- ☐ a) The ability to run an agent
 - ☐ b) Text-to-speech conversion
 - ☒ c) **Custom message history management or multiple calls to `Runner` within one turn**
 - ☐ d) Basic conversational turns
-

16. The `VoiceWorkflowBase` provides the interface for the text-to-text core of the voice agent.

- ☒ a) **True**
 - ☐ b) False
-

17. If a `VoiceWorkflowBase` implementation yields multiple strings, how does the `VoicePipeline` handle them?

- ☐ a) It concatenates them into one long string
 - ☐ b) It ignores all but the first string
 - ☒ c) **It converts each yielded string into a segment of streaming audio in sequence**
 - ☐ d) It buffers them until the entire conversation is complete
-

18. `SingleAgentWorkflowCallbacks` provides `on_run()` method to allow developers to insert custom logic at what point?

- ☐ a) After the agent produces its final response
 - ☐ b) When the entire pipeline shuts down
 - ☒ c) **At the very beginning of the `SingleAgentVoiceWorkflow`'s execution for a new transcription**
 - ☐ d) When the text-to-speech model starts
-

19. Why might a developer choose to subclass `VoiceWorkflowBase` instead of using `SingleAgentVoiceWorkflow`?

- ☐ a) To avoid using any agents
 - ☐ b) To only process static audio input
 - ☒ c) **To implement custom multi-agent orchestration, complex state management, or conditional logic**
 - ☐ d) To print all trace data to the console
-

20. The `VoiceWorkflowBase`'s design, with its `AsyncIterator[str]` return type, aligns with the concept of:

- ☐ a) Batch processing
 - ☒ b) **Stream processing and responsiveness**
 - ☐ c) Synchronous blocking operations
 - ☐ d) One-shot function calls
-

:

OpenAI Agents SDK - Audio Input MCQs

1. Which class is designed for providing static, complete audio data to the `VoicePipeline`?

- ☐ a) `StreamedAudioInput`
 - ☒ b) **`AudioInput`**
 - ☐ c) `AudioBuffer`
 - ☐ d) `VoicePipelineConfig`
-

2. What type of data is primarily stored in the `buffer` attribute of `AudioInput`?

- ☐ a) A Python list of integers
 - ☐ b) A base64 encoded string
 - ☒ c) **A NumPy array (NDArray) of int16 or float32**
 - ☐ d) A raw bytes object without format
-

3. The `frame_rate` attribute of `AudioInput` defaults to:

- ☐ a) 8000 Hz
 - ☐ b) 16000 Hz
 - ☒ c) **24000 Hz (implied by `DEFAULT_SAMPLE_RATE`)**
 - ☐ d) 44100 Hz
-

4. Which method of `AudioInput` would you use if you needed to send the audio data as part of a JSON payload that expects embedded binary data?

- ☐ a) `to_audio_file()`
 - ☒ b) `to_base64()`
 - ☐ c) `get_buffer()`
 - ☐ d) `save_to_disk()`
-

5. What is the primary use case for `StreamedAudioInput`?

- ☐ a) Processing a pre-recorded audio file from a database
 - ☒ b) **Handling real-time, continuous audio input from sources like a live microphone**
 - ☐ c) Converting text to speech
 - ☐ d) Analyzing audio for specific keywords in a batch
-

6. Which method is used to add new audio data to an ongoing `StreamedAudioInput`?

- ☐ a) `set_buffer()`
 - ☐ b) `add_audio()`
 - ☒ c) `append_audio_data()`
 - ☐ d) `push_audio()`
-

7. The `add_audio()` method in `StreamedAudioInput` is async. This indicates its suitability for:

- ☐ a) Blocking the main thread until audio is fully processed
 - ☒ b) **Non-blocking operations, crucial for maintaining responsiveness in real-time applications**
 - ☐ c) Synchronous file I/O
 - ☐ d) Executing only once per pipeline run
-

8. If you have a `.wav` file stored on disk that you want to process through `VoicePipeline`, which input class would you typically use?

- ☐ a) `StreamedAudioInput`
 - ☒ b) **`AudioInput`**
 - ☐ c) Both equally
 - ☐ d) Neither, you'd need a custom input handler
-

9. The `to_audio_file()` method returns a tuple containing filename, bytes (as `BytesIO`), and `content_type`. This is useful for:

- ☐ a) Performing in-memory audio analysis
 - ☐ b) Directly playing the audio through speakers
 - ☒ c) **Preparing the audio for API calls that expect file-like uploads (e.g., as multipart/form-data)**
 - ☐ d) Compressing the audio data
-

10. What does `sample_width: int = 2` typically signify for audio data in `AudioInput`?

- ☐ a) The number of audio channels
 - ☒ b) **That each audio sample occupies 2 bytes (e.g., 16-bit audio)**
 - ☐ c) The audio's sampling rate
 - ☐ d) The duration of the audio in seconds
-

11. The `channels` attribute of `AudioInput` represents:

- ☐ a) The frequency range of the audio
 - ☒ b) **The number of audio channels (e.g., mono or stereo)**
 - ☐ c) The depth of the audio buffer
 - ☐ d) The volume level
-

12. A key difference between `AudioInput` and `StreamedAudioInput` is that `AudioInput` deals with audio that is:

- ☐ a) Always very short
 - ☐ b) Processed only in the background
 - ☒ c) **Entirely available upfront**
 - ☐ d) Encrypted
-

13. What would be a scenario where `StreamedAudioInput` is indispensable?

- ☐ a) Analyzing a pre-recorded podcast
 - ☐ b) Converting a large audio file to text offline
 - ☒ c) **Building a voice assistant that responds while the user is still speaking**
 - ☐ d) Generating an audio summary from a textual transcript
-

14. `AudioInput` is defined as a `dataclass`. This implies its primary role is to:

- ☐ a) Provide complex business logic
 - ☒ b) **Serve as a simple container for data with auto-generated methods like `__init__`, `__repr__`, etc.**
 - ☐ c) Manage external file system operations
 - ☐ d) Implement abstract methods
-

15. Both `AudioInput` and `StreamedAudioInput` expect audio data in their buffer or `add_audio` methods to be:

- ☐ a) Raw Python lists
 - ☐ b) Base64 encoded strings
 - ☒ c) **NumPy arrays (NDArray)**
 - ☐ d) Bytes objects
-

16. Which of the following is NOT a characteristic of StreamedAudioInput?

- ☒ c) It is typically used for pre-recorded, static audio files
 - ☐ a) It supports pushing audio chunks incrementally
 - ☐ b) Its add_audio method is asynchronous
 - ☐ d) It is designed for real-time audio processing
-

17. If to_base64 () were called on an AudioInput object containing 16-bit mono audio at 24kHz for 1 second, the resulting string would be:

- ☐ a) A very short string, as it's compressed
 - ☐ b) Always "audio_data_base64"
 - ☒ c) A longer string, as base64 encoding expands binary data by about 33%
 - ☐ d) An integer value representing the audio length
-

18. The DEFAULT_SAMPLE_RATE is a common audio characteristic that defines:

- ☐ a) The volume level
 - ☒ b) How many samples per second are taken from an analog audio signal
 - ☐ c) The number of channels
 - ☐ d) The bit depth of each sample
-

19. Which class would you use to provide input to the VoicePipeline if your audio comes from a continuous byte stream without a known end?

- ☐ a) AudioInput (after buffering the entire stream)
 - ☒ b) StreamedAudioInput
 - ☐ c) VoicePipelineConfig
 - ☐ d) STTModel
-

20. The functions to_audio_file () and to_base64 () of AudioInput serve as:

- ☐ a) Primary methods for audio playback
 - ☐ b) Methods for altering the audio content
 - ☒ c) Utility methods for different ways of representing or transmitting the static audio data
 - ☐ d) Internal-only methods for pipeline processing
-

OpenAI Agents SDK - StreamedAudioResult MCQs

. What is StreamedAudioResult designed to be?

- ☐ a) The input to the VoicePipeline.
- ☐ b) A configuration object for TTS models.

- ✓c) **The output of a VoicePipeline, streaming events and audio.**
 - ✗d) A class for recording audio from a microphone.
-

2. The primary benefit of `StreamedAudioResult`'s streaming nature is:

- ✗a) Reduced overall processing time.
 - ✗b) Simplified debugging.
 - ✓c) **Improved perceived latency and responsiveness in conversational AI.**
 - ✗d) Guaranteed perfect audio fidelity.
-

3. Which type of model is explicitly required in the `StreamedAudioResult` constructor?

- ✗a) STTModel
 - ✓b) **TTSModel**
 - ✗c) Agent
 - ✗d) VoiceWorkflowBase
-

4. Why does `StreamedAudioResult` need a `TTSModel` in its constructor?

- ✗a) To manage the conversation history.
 - ✗b) To transcribe the input audio.
 - ✓c) **Because `StreamedAudioResult` is responsible for converting the workflow's text output into audio.**
 - ✗d) To generate unique IDs for audio segments.
-

5. What is the return type of the `StreamedAudioResult.stream()` method?

- ✗a) `str` (a single audio file path)
 - ✗b) `bytes` (a complete audio buffer)
 - ✓c) **`AsyncIterator[VoiceStreamEvent]`**
 - ✗d) `dict` (containing audio metadata)
-

6. `VoiceStreamEvent` objects yielded by `stream()` can represent:

- ✗a) Only raw audio bytes.
 - ✗b) Only error messages.
 - ✓c) **Chunks of audio data, textual parts of the response, and end-of-speech signals.**
 - ✗d) Configuration changes during streaming.
-

7. The `stream()` method being `async` implies it's designed for:

- ✗a) Synchronous, blocking execution.
- ✗b) Batch processing of large audio files.
- ✓c) **Non-blocking, real-time delivery of events.**
- ✗d) Offloading computation to a different process.

8. Which parameter in the `StreamedAudioResult` constructor dictates specific voice characteristics like speaking rate or pitch?

- ☐ a) `tts_model`
 - ☒ b) `tts_settings`
 - ☐ c) `voice_pipeline_config`
 - ☐ d) `workflow`
-

9. What would typically cause a `VoiceStreamEvent` of type `AudioEvent` to be yielded?

- ☐ a) A new user transcription is received.
 - ☐ b) The `VoicePipeline` is initialized.
 - ☒ c) A chunk of text from the workflow has been synthesized into audio by the TTS model.
 - ☐ d) The agent calls a tool.
-

10. If the `VoiceWorkflowBase` yields text incrementally, how does `StreamedAudioResult` leverage this?

- ☐ a) It waits for all text to be available before starting TTS.
 - ☐ b) It discards all but the last piece of text.
 - ☒ c) It converts each yielded text chunk into audio in real-time, sending it as `AudioEvents`.
 - ☐ d) It sends the text directly to the user without conversion.
-

11. The `voice_pipeline_config` parameter in `StreamedAudioResult`'s constructor is relevant for:

- ☐ a) Defining the agent's behavior logic.
 - ☐ b) Setting the API key for the TTS model.
 - ☒ c) Providing overall pipeline settings like audio format or chunking for the output stream.
 - ☐ d) Managing the input audio buffer.
-

12. A `StreamedAudioResult` is designed to be consumed:

- ☐ a) Only once, then it's exhausted.
 - ☒ b) As an asynchronous stream, allowing an external client to process events as they arrive.
 - ☐ c) By saving all its content to a file first.
 - ☐ d) Through synchronous blocking calls.
-

13. What kind of event might signal to the consuming application that the agent has finished its spoken response for the current turn?

- ☐ a) An `AudioEvent` with an empty buffer.
- ☐ b) A `TextEvent` with an empty string.
- ☒ c) An `EndOfSpeechEvent` (or similar signal within `VoiceStreamEvent`).
- ☐ d) A `TimeoutEvent`.

14. The `StreamedAudioResult` effectively acts as the orchestrator for which step of the `VoicePipeline`?

- ☐ a) Transcription (STT).
 - ☐ b) Workflow execution.
 - ☒ c) **Text-to-Speech (TTS) and output streaming.**
 - ☐ d) Input audio buffering.
-

15. What is `StreamedAudioResult`'s relationship to `VoicePipeline`?

- ☐ a) `StreamedAudioResult` is a component inside the `VoicePipeline`.
 - ☐ b) `StreamedAudioResult` configures the `VoicePipeline`.
 - ☒ c) **`StreamedAudioResult` is the return value of the `VoicePipeline.run()` method.**
 - ☐ d) They are unrelated classes.
-

16. What would happen if `StreamedAudioResult` did not have access to the `tts_model` and `tts_settings`?

- ☐ a) It would still stream events, but without any audio data.
 - ☐ b) It would fallback to a default TTS model.
 - ☒ c) **It would not be able to perform text-to-speech conversion and thus couldn't stream audio.**
 - ☐ d) It would only stream `TextEvents`.
-

17. The design of `StreamedAudioResult` directly supports which common requirement for modern voice interfaces?

- ☐ a) Offline processing of large audio datasets.
 - ☐ b) Batch processing of hundreds of requests simultaneously.
 - ☒ c) **Low-latency, interactive conversational experiences.**
 - ☐ d) Manual audio editing capabilities.
-

18. `StreamedAudioResult` is located in `src/agents/voice/result.py`, reinforcing its role as:

- ☐ a) A utility for audio file manipulation.
 - ☐ b) A backend database connector.
 - ☒ c) **A component focused on the outcome/output of voice processing.**
 - ☐ d) An input formatter.
-

19. The use of an `AsyncIterator` for `stream()` allows the consuming application to:

- ☐ a) Block until all audio is generated.
 - ☐ b) Request specific events by index.
 - ☒ c) **Process events as they are produced, without waiting for the entire stream to finish.**
 - ☐ d) Store the entire stream in memory before processing.
-

20. Which of these is not a direct responsibility of StreamedAudioResult?

- ☒ a) Converting text to audio.
 - ☒ b) Streaming generated audio data.
 - ☒ c) **Deciding the agent's textual response based on business logic.**
 - ☒ d) Emitting events related to the voice stream.
-

OpenAI Agents SDK - VoicePipelineConfig MCQs

1. What is the primary purpose of the VoicePipelineConfig dataclass?

- ☒ a) To store the actual audio data for the pipeline.
 - ☒ b) To define the abstract interface for a voice workflow.
 - ☒ c) **To provide configurable options for customizing a VoicePipeline.**
 - ☒ d) To manage the streaming output of the pipeline.
-

2. By default, VoicePipelineConfig uses which model provider?

- ☒ a) CustomVoiceModelProvider
 - ☒ b) GoogleCloudVoiceModelProvider
 - ☒ c) **OpenAIVoiceModelProvider**
 - ☒ d) AzureVoiceModelProvider
-

3. If tracing_disabled is set to True in VoicePipelineConfig, what is the effect?

- ☒ a) Only audio data tracing is disabled.
 - ☒ b) Only sensitive data tracing is disabled.
 - ☒ c) **Tracing for the VoicePipeline's operations is completely disabled.**
 - ☒ d) Tracing for the internal VoiceWorkflow is also disabled, regardless of its own settings.
-

4. The trace_include_sensitive_data attribute in VoicePipelineConfig specifically refers to sensitive data from:

- ☒ a) The user's input audio.
 - ☒ b) The VoiceWorkflow's internal operations.
 - ☒ c) **The VoicePipeline's own processing (e.g., transcriptions generated by the pipeline).**
 - ☒ d) All data within the entire application.
-

5. What is the default value for trace_include_sensitive_audio_data?

- ☒ a) **False**
- ☒ b) True
- ☒ c) None
- ☒ d) Depends on the model_provider.

6. Which attribute allows you to link multiple traces from the same conversation or process together?

- ☐ a) workflow_name
 - ☐ b) trace_metadata
 - ☐ c) stt_settings
 - ☒ d) group_id
-

7. If group_id is not explicitly provided in VoicePipelineConfig, what happens?

- ☐ a) It defaults to an empty string.
 - ☒ b) A random group_id is generated using gen_group_id().
 - ☐ c) An error is raised.
 - ☐ d) It reuses the workflow_name as the group ID.
-

8. VoicePipelineConfig includes separate settings objects for:

- ☐ a) Input and output channels.
 - ☐ b) Audio format and sample rate.
 - ☒ c) Speech-to-Text (STT) and Text-to-Speech (TTS) models.
 - ☐ d) Tracing and non-tracing operations.
-

9. The workflow_name attribute in VoicePipelineConfig defaults to:

- ☐ a) Default Voice Workflow
 - ☐ b) Main Agent Workflow
 - ☒ c) Voice Agent
 - ☐ d) The name of the VoiceWorkflowBase instance.
-

10. What kind of data can be added to a trace using the trace_metadata attribute?

- ☐ a) Only predefined integer values.
 - ☐ b) Only string identifiers.
 - ☒ c) An optional dictionary of additional custom key-value pairs.
 - ☐ d) Raw audio buffers.
-

11. The stt_settings attribute within VoicePipelineConfig allows customization of:

- ☐ a) The voice of the TTS model.
 - ☐ b) The number of channels for audio output.
 - ☒ c) Parameters specific to the Speech-to-Text model (e.g., language).
 - ☐ d) The maximum queue size for audio input.
-

12. The `VoicePipelineConfig` is a dataclass which means:

- ☐ a) It must contain abstract methods.
 - ☒ b) It's **primarily a container for data with auto-generated methods like `__init__`, `__repr__`, etc.**
 - ☐ c) It is designed for complex procedural logic.
 - ☐ d) It can only be inherited from, not instantiated directly.
-

13. Which of these attributes provides a global kill switch for tracing all pipeline operations?

- ☐ a) `trace_include_sensitive_data`
 - ☐ b) `trace_metadata`
 - ☒ c) **`tracing_disabled`**
 - ☐ d) `group_id`
-

14. The `VoicePipelineConfig` is passed to the `VoicePipeline`'s constructor to:

- ☐ a) Trigger an immediate run of the pipeline.
 - ☐ b) Return the final audio result.
 - ☒ c) **Configure the pipeline's behavior and its internal components.**
 - ☐ d) Only set up the `tts_model`.
-

15. If `trace_include_sensitive_audio_data` is `False`, but `trace_include_sensitive_data` is `True`, what will be traced from the pipeline?

- ☐ a) No sensitive data at all.
 - ☒ b) **Non-audio sensitive data (e.g., text) will be traced, but audio data will not.**
 - ☐ c) Only audio data will be traced.
 - ☐ d) Both audio and non-audio sensitive data will be traced.
-

16. What is the type of the default value provided for `model_provider`?

- ☐ a) `TTSModel`
 - ☐ b) `STTModel`
 - ☒ c) **`VoiceModelProvider`**
 - ☐ d) `VoicePipelineConfig`
-

17. The `tts_settings` attribute focuses on configuring:

- ☐ a) The input audio sample rate.
 - ☐ b) The transcription language.
 - ☒ c) **The characteristics of the synthesized speech (e.g., voice, pitch, rate).**
 - ☐ d) The buffer size for streaming input.
-

18. `VoicePipelineConfig` allows for control over the `VoicePipeline`'s interaction with which observability feature?

- ☐ a) Logging
 - ☐ b) Metrics
 - ☒ c) **Tracing**
 - ☐ d) Alerts
-

19. Setting `group_id` helps in:

- ☐ a) Decreasing the latency of audio processing.
 - ☐ b) Increasing the accuracy of transcription.
 - ☒ c) **Providing contextual linkage across multiple conversation turns/traces in a monitoring system.**
 - ☐ d) Reducing the size of exported trace data.
-

20. When creating a `VoicePipeline`, if you don't provide a `VoicePipelineConfig` object, what happens?

- ☐ a) An error is raised.
 - ☐ b) The pipeline runs with no configuration, leading to unpredictable behavior.
 - ☒ c) **A default `VoicePipelineConfig` instance with all its default values is used.**
 - ☐ d) The `VoicePipeline` prompts the user for configuration details.
-

OpenAI Agents SDK - VoiceStreamEvent MCQs

1. What is `VoiceStreamEvent` defined as?

- ☐ a) A concrete class for all voice stream events.
 - ☒ b) **A `TypeAlias (Union)` of different specific voice stream event types.**
 - ☐ c) An abstract base class for streaming.
 - ☐ d) A configuration object for audio.
-

2. From which method are `VoiceStreamEvents` primarily streamed?

- ☐ a) `VoicePipeline.run()`
 - ☒ b) **`StreamedAudioResult.stream()`**
 - ☐ c) `VoiceWorkflowBase.run()`
 - ☐ d) `AudioInput.add_audio()`
-

3. Which of the following is NOT a defined type of `VoiceStreamEvent` in the provided snippet?

- ☒ a) **`VoiceStreamEventAudio`**
 - ☒ b) **`VoiceStreamEventLifecycle`**
 - ☒ c) **`VoiceStreamEventError`**
 - ☐ d) **`VoiceStreamEventText` ← Correct answer: d, not defined in the Union in the source.**
-

4. What is the main content of a `VoiceStreamEventAudio` object?

- ☒ a) A string representing transcribed text.
 - ☒ b) An integer indicating audio volume.
 - ☒ c) **A NumPy array (NDArray) containing raw audio data.**
 - ☒ d) A boolean indicating if audio is playing.
-

5. The `type` attribute of `VoiceStreamEventAudio` will always have which literal string value?

- ☒ a) **"audio"**
 - ☒ b) "stream_audio"
 - ☒ c) "voice_stream_event_audio"
 - ☒ d) "audio_data"
-

6. Which `VoiceStreamEvent` type is used to signal the beginning or end of a conversational segment?

- ☒ a) `VoiceStreamEventAudio`
 - ☒ b) **`VoiceStreamEventLifecycle`**
 - ☒ c) `VoiceStreamEventError`
 - ☒ d) `VoiceStreamEventMetadata`
-

7. The `event` attribute of `VoiceStreamEventLifecycle` can take which values?

- ☒ a) "start", "stop", "pause"
 - ☒ b) "audio_started", "audio_stopped", "audio_error"
 - ☒ c) **"turn_started", "turn_ended", "session_ended"**
 - ☒ d) "processing", "responding", "idle"
-

8. When would a `VoiceStreamEventError` typically be yielded?

- ☒ a) Whenever the `VoicePipeline` starts.
 - ☒ b) **When an unhandled exception occurs within the pipeline's streaming process.**
 - ☒ c) When the audio input is empty.
 - ☒ d) After the `VoicePipeline` successfully completes.
-

9. What information is contained within a `VoiceStreamEventError` object?

- ☒ a) A simple error code.
 - ☒ b) A string message only.
 - ☒ c) **The actual Python Exception object that occurred.**
 - ☒ d) The timestamp of the error.
-

10. The `type` attribute across all `VoiceStreamEvent` dataclasses serves what purpose?

- ☐ a) It determines the size of the event.
 - ☐ b) It specifies the encoding of the data.
 - ☒ c) **It provides a clear, programmatic way to identify the specific kind of event.**
 - ☐ d) It indicates the priority of the event.
-

11. Why is it beneficial to stream `VoiceStreamEvents` instead of providing a single, complete output?

- ☐ a) It consumes less memory on the server.
 - ☐ b) It guarantees lower network latency.
 - ☒ c) **It enables real-time responsiveness and allows client applications to react incrementally.**
 - ☐ d) It simplifies the internal logic of the `VoicePipeline`.
-

12. A client receiving `VoiceStreamEventLifecycle` with `event="turn_started"` might:

- ☐ a) Immediately stop listening for user input.
 - ☐ b) Start playing a pre-recorded message.
 - ☒ c) **Update its UI to indicate the agent is processing/responding.**
 - ☐ d) Close the connection.
-

13. The `data` attribute in `VoiceStreamEventAudio` can be `None`. When might this occur?

- ☐ a) Only if an error happened during audio synthesis.
 - ☒ b) **Potentially to signal silence, a pause, or the end of an audio segment without actual samples.**
 - ☐ c) Only if the `tts_model` is disabled.
 - ☐ d) If the `frame_rate` is zero.
-

14. `VoiceStreamEventLifecycle` with `event="session_ended"` would typically be used to signal:

- ☐ a) That a single turn of conversation has finished.
 - ☒ b) **That the entire voice interaction session has concluded.**
 - ☐ c) That an error occurred during the session.
 - ☐ d) That audio data is about to be sent.
-

15. What type of Python construct defines `VoiceStreamEvent`?

- ☐ a) class
 - ☐ b) enum
 - ☒ c) **TypeAlias (a union of dataclasses)**
 - ☐ d) protocol
-

16. Which `VoiceStreamEvent` is crucial for playing back the agent's response to the user?

- ☐ a) `VoiceStreamEventLifecycle`
 - ☒ b) **`VoiceStreamEventAudio`**
 - ☐ c) `VoiceStreamEventError`
 - ☐ d) All of them are equally crucial.
-

17. If a `VoiceStreamEventError` is received, what should a robust client application ideally do?

- ☐ a) Ignore it and continue processing.
 - ☐ b) Immediately crash.
 - ☒ c) **Log the error, potentially display a user-friendly message, and decide whether to terminate or attempt recovery.**
 - ☐ d) Send the error back to the `VoicePipeline`.
-

18. The `Literal` type hints used for `type` and `event` attributes mean:

- ☐ a) The values can be any string.
 - ☒ b) **The values are restricted to a specific set of predefined string constants.**
 - ☐ c) The values are numerical.
 - ☐ d) The values are optional.
-

19. The events defined in `src/agents/voice/events.py` directly support the responsiveness provided by:

- ☐ a) `AudioInput`
 - ☐ b) `VoiceWorkflowBase`
 - ☒ c) **`StreamedAudioResult`**
 - ☐ d) `VoicePipelineConfig`
-

20. A `VoiceStreamEvent` does not directly contain:

- ☐ a) Audio data.
 - ☐ b) Lifecycle signals.
 - ☐ c) Error objects.
 - ☒ d) **The transcribed user input.**
-

OpenAI Agents SDK - STTWebsocketConnectionError MCQs

1. From which base class does `STTWebsocketConnectionError` directly inherit?

- ☐ a) `Exception`

- ☒ b) **AgentsException**
 - ☐ c) WebSocketError
 - ☐ d) VoiceException
-

2. What specific type of connection failure does STTWebSocketConnectionError indicate?

- ☐ a) An HTTP API connection failure for text processing.
 - ☐ b) A database connection failure.
 - ☒ c) **A WebSocket connection failure for Speech-to-Text services.**
 - ☐ d) A connection timeout during file download.
-

3. In the context of a VoicePipeline, when would this exception typically be raised?

- ☐ a) When the Text-to-Speech model fails to synthesize audio.
 - ☐ b) When the VoiceWorkflowBase encounters an internal logic error.
 - ☒ c) **When the real-time connection to the Speech-to-Text service cannot be established or is lost.**
 - ☐ d) When the VoicePipelineConfig is invalid.
-

4. Which of the following is a common reason for an STTWebSocketConnectionError?

- ☐ a) Incorrect grammar in the user's speech.
 - ☐ b) The STT model providing a low confidence score for transcription.
 - ☒ c) **A firewall blocking the WebSocket connection.**
 - ☐ d) Running out of disk space on the client machine.
-

5. This exception primarily focuses on an issue with the:

- ☐ a) Quality of transcription.
 - ☐ b) Authentication of the user.
 - ☒ c) **Underlying communication channel.**
 - ☐ d) Performance of the agent's response.
-

6. If you catch an STTWebSocketConnectionError, what is the most immediate area you should investigate?

- ☐ a) The Text-to-Speech model's configuration.
 - ☐ b) The logic within your VoiceWorkflowBase.
 - ☒ c) **Network connectivity and the STT service's status/URL.**
 - ☐ d) The amount of RAM available on your system.
-

7. Which type of STT interaction is most likely to rely on WebSockets and thus potentially raise this error?

- ☐ a) Batch processing of large, pre-recorded audio files.
- ☒ b) **Real-time, continuous speech transcription.**
- ☐ c) Offline language model training.
- ☐ d) Single, short audio clip analysis.

8. By inheriting from `AgentsException`, `STTWebSocketConnectionError` allows for:

- ☐ a) Automatic retry mechanisms.
 - ☒ b) **More granular error handling specific to the SDK's functionalities.**
 - ☐ c) Direct access to the VoicePipeline's internal state.
 - ☐ d) Automatic logging to a central server.
-

9. Which of these scenarios would less likely directly cause an `STTWebSocketConnectionError` (though it might be a subsequent issue)?

- ☐ a) The STT server going offline.
 - ☐ b) An invalid WebSocket endpoint URL.
 - ☒ c) **The user speaking very quietly, leading to poor audio input.**
 - ☐ d) A proxy server interfering with the WebSocket handshake.
-

10. The presence of "WebSocket" in the exception name is significant because:

- ☐ a) It means the error only occurs on web browsers.
 - ☐ b) It indicates the error is related to web security.
 - ☒ c) **It points to a persistent, bidirectional communication protocol used for streaming.**
 - ☐ d) It refers to the use of webhooks for notifications.
-

11. What would be an appropriate troubleshooting step if you encounter this exception?

- ☐ a) Change the `tts_model` setting.
 - ☐ b) Adjust the `workflow_name` in `VoicePipelineConfig`.
 - ☒ c) **Ping the STT service's domain or check its status page.**
 - ☐ d) Modify the `VoiceWorkflowHelper` logic.
-

12. This exception is part of the `src/agents/voice/exceptions.py` module, indicating its specific relevance to:

- ☐ a) General Python errors.
 - ☐ b) Database errors.
 - ☒ c) **Voice-related functionalities within the Agents SDK.**
 - ☐ d) File system operations.
-

13. Could an incorrect API key or authentication token lead to an `STTWebSocketConnectionError`?

- ☒ a) **Yes, as the service might refuse the WebSocket handshake due to authentication failure.**
 - ☐ b) No, authentication issues would raise a different error.
 - ☐ c) Only if the API key is too long.
 - ☐ d) Only if the API key is too short.
-

14. If you write a try...except block to handle errors from VoicePipeline, which exception type would catch all SDK-specific errors including this one?

- ☐ a) Exception
 - ☐ b) RuntimeError
 - ☒ c) **AgentsException**
 - ☐ d) STTException
-

15. STTWebSocketConnectionError suggests that the problem is preventing the agent from:

- ☐ a) Responding to the user in audio.
 - ☐ b) Calling external tools.
 - ☒ c) **Receiving and processing spoken input from the user.**
 - ☐ d) Storing conversation history.
-

16. What kind of communication is typical over a WebSocket connection used for STT?

- ☐ a) Only client sending audio to server.
 - ☐ b) Only server sending transcription to client.
 - ☒ c) **Bidirectional streaming of audio from client and text from server.**
 - ☐ d) One-time request-response cycles.
-

17. If a developer uses a Speech-to-Text service that relies on HTTP POST requests rather than WebSockets, would this specific exception likely be raised for connection failures?

- ☐ a) Yes, because it's a general STT error.
 - ☒ b) **No, because the error is specific to WebSocket connections.**
 - ☐ c) Only if the HTTP POST request times out.
 - ☐ d) It depends on the VoicePipelineConfig.
-

18. The term "connection" in the exception name refers to the link between:

- ☐ a) The VoicePipeline and the VoiceWorkflowBase.
 - ☐ b) The STTModel and the TTSTModel.
 - ☒ c) **The SDK client (running your agent) and the remote STT service.**
 - ☐ d) The agent and the user's local audio device.
-

19. Which Python module would you look into if you wanted to implement specific error handling for this exception?

- ☐ a) src/agents/voice/pipeline.py
 - ☐ b) src/agents/voice/workflow.py
 - ☒ c) **src/agents/voice/exceptions.py**
 - ☐ d) src/agents/voice/input.py
-

20. What is a key characteristic that distinguishes `STTWebSocketConnectionError` from other general Exception types?

- ☐ a) It is always a fatal error.
 - ☐ b) It provides a detailed stack trace of the entire application.
 - ☒ c) **It provides specific context about a failure in the STT WebSocket communication, aiding targeted debugging.**
 - ☐ d) It automatically retries the connection.
-

OpenAI Agents SDK - Voice Model MCQs

What is the primary purpose of the `TTSVoice` type alias?

- ☐ a) To define the temperature setting for TTS models.
 - ☐ b) To specify the audio buffer size for streaming.
 - ☒ c) **To list predefined, literal string values for available TTS voices.**
 - ☐ d) To indicate the language of the TTS output.
-

2. Which attribute in `TTSModelSettings` allows for post-processing of audio data before it's streamed?

- ☐ a) voice
 - ☐ b) buffer_size
 - ☐ c) instructions
 - ☒ d) **transform_data**
-

3. The `text_splitter` attribute in `TTSModelSettings` is crucial for:

- ☐ a) Detecting conversational turns in audio.
 - ☒ b) **Chunking text into smaller pieces for streaming Text-to-Speech.**
 - ☐ c) Splitting audio data into separate channels.
 - ☐ d) Transforming the data type of the audio.
-

4. The `TTSModel.run()` method returns an `AsyncIterator[bytes]`. What does bytes represent in this context?

- ☐ a) The raw text input.
 - ☐ b) A base64 encoded string.
 - ☒ c) **Chunks of audio data in PCM format.**
 - ☐ d) An error message.
-

5. What is the main role of the `TTSModel` abstract base class?

- ☐ a) To configure the input audio settings.
- ☐ b) To provide a factory for STT models.

- ☒ c) To define the interface for any Text-to-Speech model.
 - ☐ d) To manage the conversation history.
-

6. **StreamedTranscriptionSession** is designed for:

- ☐ a) Transcribing static audio files.
 - ☒ b) **Streaming text transcriptions from continuous audio input.**
 - ☐ c) Converting text to speech in real-time.
 - ☐ d) Managing database connections for transcription data.
-

7. Which attribute in **STTModelSettings** can be used to guide the STT model's transcription, especially for domain-specific vocabulary?

- ☐ a) language
 - ☐ b) temperature
 - ☒ c) **prompt**
 - ☐ d) turn_detection
-

8. The **STTModel.transcribe()** method is used for:

- ☐ a) Real-time, continuous audio input.
 - ☒ b) **Static AudioInput (complete audio buffers).**
 - ☐ c) Generating audio from text.
 - ☐ d) Detecting silence in an audio stream.
-

9. What is the purpose of **VoiceModelProvider**?

- ☐ a) To directly perform STT and TTS conversions.
 - ☐ b) To define the settings for STT and TTS models.
 - ☒ c) **To act as a factory for creating STT and TTS model instances by name.**
 - ☐ d) To store the configuration of the entire VoicePipeline.
-

10. If **TTSModelSettings.speed** is set to 0.5, how would the TTS model read the text?

- ☐ a) Twice as fast as normal.
 - ☐ b) At a random speed.
 - ☒ c) **At half the normal speed.**
 - ☐ d) It would ignore the setting.
-

11. **STTModel.create_session()** returns an instance of:

- ☐ a) STTModelSettings
- ☐ b) AudioInput
- ☒ c) **StreamedTranscriptionSession**

- ☒ d) VoiceModelProvider
-

12. The instructions attribute in TTSMModelSettings can help control the TTS model's:

- ☒ a) Input audio language.
 - ☒ b) Transcription accuracy.
 - ☒ c) **Tone or how it handles partial sentences.**
 - ☒ d) Network connectivity.
-

13. Which of these tracing parameters are passed to both STTModel.transcribe() and STTModel.create_session()?

- ☒ a) group_id
 - ☒ b) workflow_name
 - ☒ c) **trace_include_sensitive_data and trace_include_sensitive_audio_data**
 - ☒ d) buffer_size
-

14. What does the AsyncIterator[str] return type of StreamedTranscriptionSession.transcribe_turns() signify?

- ☒ a) It returns a single, complete transcription string.
 - ☒ b) **It yields textual transcriptions incrementally, turn by turn.**
 - ☒ c) It returns a list of all detected turns at once.
 - ☒ d) It produces a stream of audio bytes.
-

15. TTSMModelSettings.dtype = int16 indicates:

- ☒ a) The voice model uses 16 different voices.
 - ☒ b) The speed of the audio is 1.6x.
 - ☒ c) **The audio data will be returned in 16-bit integer format.**
 - ☒ d) The buffer size is limited to 16 bytes.
-

16. What is the role of STTModelSettings.turn_detection?

- ☒ a) To identify the speaker's accent.
 - ☒ b) To determine the language of the audio.
 - ☒ c) **To configure how conversational turns are recognized in streamed audio.**
 - ☒ d) To set the overall volume of the transcribed text.
-

17. If you have a custom STT model you want to integrate, which abstract class would you need to implement?

- ☒ a) VoiceModelProvider
 - ☒ b) TTSMModel
 - ☒ c) **STTModel**
 - ☒ d) StreamedTranscriptionSession
-

18. The VoiceModelProvider interface is designed for:

- ☐ a) Directly converting audio to text.
 - ☐ b) Applying settings to models.
 - ☒ c) **Abstracting the access and instantiation of specific STT/TTS model implementations.**
 - ☐ d) Handling errors during model execution.
-

19. Why is StreamedTranscriptionSession.close() an async method?

- ☐ a) It needs to perform immediate, blocking operations.
 - ☒ b) **To allow for asynchronous resource cleanup (e.g., closing network connections).**
 - ☐ c) It only operates on local files.
 - ☐ d) It is meant to be called at the very start of a session.
-

20. The model_name property is present in both TTSModel and STTModel to:

- ☐ a) Specify the file path of the model.
 - ☐ b) Indicate the version of the model.
 - ☒ c) **Provide a unique identifier for the specific model being used.**
 - ☐ d) Store a description of the model's capabilities.
-

OpenAI Agents SDK - Text Utility MCQs

1. What is the primary function of get_sentence_based_splitter?

- ☐ a) To concatenate multiple text strings into one.
 - ☐ b) To perform sentiment analysis on text.
 - ☒ c) **To return a function that splits text into chunks based on sentence boundaries.**
 - ☐ d) To convert text into audio.
-

2. The get_sentence_based_splitter function itself returns a:

- ☐ a) str
 - ☐ b) list[str]
 - ☒ c) **Callable (a function)**
 - ☐ d) tuple[str, str]
-

3. Why is sentence-based splitting particularly useful for Text-to-Speech (TTS)?

- ☐ a) It reduces the overall size of the text.
- ☐ b) It simplifies grammatical analysis.
- ☒ c) **It helps ensure natural intonation and pacing by sending complete thoughts/sentences to the TTS model.**

- ☒ d) It speeds up the initial transcription process.
-

4. What is the default value for the `min_sentence_length` parameter?

- ☒ a) 5
 - ☒ b) 10
 - ☒ c) **20**
 - ☒ d) 50
-

5. If a sentence detected by the splitter is shorter than `min_sentence_length`, what might happen?

- ☒ a) It will be discarded.
 - ☒ b) It will cause an error.
 - ☒ c) **It might be grouped with adjacent sentences to form a larger chunk.**
 - ☒ d) It will be sent as a separate, very small chunk regardless.
-

6. The function returned by `get_sentence_based_splitter` takes which type of input?

- ☒ a) A list of sentences.
 - ☒ b) An audio buffer.
 - ☒ c) **A single str (the text to be split).**
 - ☒ d) A `TTSModelSettings` object.
-

7. What is the return type of the function returned by `get_sentence_based_splitter`?

- ☒ a) **list[str]**
- ☒ b) `AsyncIterator[str]`
- ☒ c) `str`
- ☒ d) `tuple[str, str]`

Note: This is actually context-dependent. Earlier versions returned `tuple[str, str]` per call, while in newer simplified implementations, a full list may be returned. Let me know your SDK version to adjust.

8. In the `tuple[str, str]` returned by the splitter function, what does the first `str` typically represent?

- ☒ a) The remaining, unprocessed text.
 - ☒ b) A random segment of the input.
 - ☒ c) **A complete chunk of text (one or more full sentences) ready for processing.**
 - ☒ d) An error message.
-

9. What does the second `str` in the returned `tuple[str, str]` typically represent?

- ☒ a) The first sentence found.
- ☒ b) **The remaining part of the text that couldn't form a complete sentence-based chunk yet.**
- ☒ c) A copy of the entire input text.

- ☒ d) A concatenated version of all sentences.
-

10. The `get_sentence_based_splitter` is part of the `src/agents/voice/utils.py` module, indicating it's a:

- ☒ a) Core pipeline component.
 - ☒ b) Model definition.
 - ☒ c) **Helper function for common tasks within voice agent development.**
 - ☒ d) Exception handler.
-

11. Could `get_sentence_based_splitter` be used outside of a `VoicePipeline` context?

- ☒ a) **Yes, it's a general utility for text processing.**
 - ☒ b) No, it's tightly coupled to the pipeline's internal state.
 - ☒ c) Only if you manually import all voice models.
 - ☒ d) Only for audio file processing.
-

12. If `min_sentence_length` is set to 0, what would be a likely behavior of the returned splitter?

- ☒ a) It would always return empty strings.
 - ☒ b) It would treat every character as a sentence.
 - ☒ c) **It would likely split on every sentence boundary, even for very short sentences.**
 - ☒ d) It would not split the text at all.
-

13. What problem does the `text_splitter` (where `get_sentence_based_splitter` is typically used) help solve for streaming TTS?

- ☒ a) Ensuring the correct language is used.
 - ☒ b) Reducing the CPU load of the TTS model.
 - ☒ c) **Preventing unnatural pauses or intonation issues in streamed audio caused by arbitrary text chunking.**
 - ☒ d) Authenticating with the TTS service.
-

14. Is `get_sentence_based_splitter` an asynchronous function?

- ☒ a) Yes, because it deals with streaming.
 - ☒ b) **No, the function it returns operates synchronously on strings.**
 - ☒ c) Only if `min_sentence_length` is very high.
 - ☒ d) Only when used within an `AsyncIterator`.
-

15. What type of splitting logic does this utility specifically employ?

- ☒ a) Character-based splitting.
- ☒ b) Word-based splitting.
- ☒ c) **Sentence-based splitting.**

- ☐ d) Paragraph-based splitting.
-

16. If the input text to the splitter function is "Hello.", and min_sentence_length=20, what might the output tuple be?

- ☐ a) ("Hello.", "")
 - ☒ b) ("", "Hello.")
 - ☐ c) ("Hello", ".")
 - ☐ d) ("Hello", "!")
-

17. The Callable type hint signifies that the returned object is:

- ☐ a) A class instance.
 - ☐ b) A variable.
 - ☒ c) **Something that can be called like a function.**
 - ☐ d) A data structure.
-

18. This utility is typically used as the value for which attribute of TTSMModelSettings?

- ☐ a) voice
 - ☐ b) instructions
 - ☒ c) **text_splitter**
 - ☐ d) buffer_size
-

19. get_sentence_based_splitter is an example of a:

- ☐ a) Class method.
 - ☐ b) Static method.
 - ☒ c) **Factory function.**
 - ☐ d) Constructor.
-

20. What is the primary goal of the get_sentence_based_splitter in improving the user experience of a voice agent?

- ☐ a) To make the agent respond faster.
 - ☐ b) To make the agent's voice sound more human-like.
 - ☒ c) **To ensure that the agent's spoken responses flow naturally and are easily understandable.**
 - ☐ d) To minimize the data transferred over the network.
-

OpenAI Agents SDK - OpenAIVoiceModelProvider MCQs

1. What is the primary function of OpenAIVoiceModelProvider?

- ☐ a) To manage audio input/output devices.

- ☐ b) To define a new type of OpenAI model.
 - ☒ c) **To provide access to OpenAI's Speech-to-Text and Text-to-Speech models.**
 - ☐ d) To handle network communication for voice agents.
-

2. OpenAIVoiceModelProvider inherits from which abstract base class?

- ☐ a) TTSMModel
 - ☐ b) STTModel
 - ☒ c) **VoiceModelProvider**
 - ☐ d) AsyncOpenAI
-

3. Which parameter in the constructor allows you to specify a custom OpenAI API endpoint?

- ☐ a) api_key
 - ☒ b) **base_url**
 - ☐ c) openai_client
 - ☐ d) organization
-

4. If the api_key parameter is None during OpenAIVoiceModelProvider initialization, where will it attempt to get the API key from?

- ☐ a) It will raise an immediate error.
 - ☒ b) **It will typically look for the OPENAI_API_KEY environment variable or other default locations.**
 - ☐ c) It will use a publicly available key.
 - ☐ d) It will prompt the user for the key.
-

5. The openai_client parameter in the constructor is provided for what purpose?

- ☐ a) To automatically create a new API key.
 - ☒ b) **To allow injection of an already-configured AsyncOpenAI client instance.**
 - ☐ c) To specify which OpenAI model to use globally.
 - ☐ d) To disable all tracing.
-

6. If openai_client is provided to the constructor, what happens to api_key and base_url passed to the same constructor?

- ☒ a) **They are ignored, as the provided openai_client takes precedence for configuration.**
 - ☐ b) They are used to reconfigure the provided openai_client.
 - ☐ c) They cause an error if both are present.
 - ☐ d) They are stored for future use but not applied immediately.
-

7. Which method of OpenAIVoiceModelProvider is responsible for providing an STT model?

- ☐ a) get_voice_model()

- ☐ b) create_transcriber()
 - ☒ c) get_stt_model()
 - ☐ d) get_text_model()
-

8. The get_tts_model() method returns an instance that conforms to which abstract class?

- ☐ a) STTModel
 - ☒ b) TTSTModel
 - ☐ c) VoiceModelProvider
 - ☐ d) StreamedTranscriptionSession
-

9. What is the typical OpenAI STT model name you would pass to get_stt_model()?

- ☐ a) "tts-1"
 - ☐ b) "gpt-4"
 - ☒ c) "whisper-1"
 - ☐ d) "davinci"
-

10. The organization and project parameters are primarily used for:

- ☐ a) Defining the agent's behavior.
 - ☐ b) Configuring audio input/output devices.
 - ☒ c) **Billing and resource management within OpenAI.**
 - ☐ d) Specifying the language of the voice model.
-

11. What kind of models does OpenAIVoiceModelProvider explicitly provide?

- ☐ a) Only Large Language Models (LLMs).
 - ☐ b) Only Image Generation Models.
 - ☒ c) **Speech-to-Text and Text-to-Speech models.**
 - ☐ d) Only Custom Trained Models.
-

12. If model_name is None when calling get_tts_model(), what will the provider likely do?

- ☐ a) Raise an error, as a model name is always required.
 - ☐ b) Return a random TTS model.
 - ☒ c) **Return a default TTS model (e.g., "tts-1").**
 - ☐ d) Revert to a local, non-OpenAI model.
-

13. The OpenAIVoiceModelProvider bridges the gap between the generic VoicePipeline framework and:

- ☐ a) Local file storage.
- ☐ b) User interface components.
- ☒ c) **OpenAI's specific API services for voice.**
- ☐ d) Other cloud providers.

14. Is `OpenAIVoiceModelProvider` a concrete class or an abstract class?

- ☐ a) Abstract, as it has abstract methods.
 - ☒ b) **Concrete, as it provides implementations for its base class's abstract methods.**
 - ☐ c) It's neither, it's a dataclass.
 - ☐ d) It depends on the Python version.
-

15. What is a key advantage of having a `VoiceModelProvider` abstraction like `OpenAIVoiceModelProvider`?

- ☐ a) It eliminates the need for API keys.
 - ☐ b) It forces all models to be from OpenAI.
 - ☒ c) **It allows the `VoicePipeline` to be model-agnostic, easily switching between different STT/TTS providers.**
 - ☐ d) It directly handles audio streaming to the user.
-

16. What type of client is typically instantiated internally by `OpenAIVoiceModelProvider` if `openai_client` is not provided?

- ☐ a) `OpenAIClient` (synchronous)
 - ☒ b) **`AsyncOpenAI`**
 - ☐ c) `requests.Session`
 - ☐ d) `httpx.Client`
-

17. The methods `get_stt_model` and `get_tts_model` both take `model_name: str | None`. What does this indicate?

- ☐ a) Only specific model names are allowed, no defaults.
 - ☐ b) The model name is always required.
 - ☒ c) **A specific model name can be requested, or a default will be used if `None` is provided.**
 - ☐ d) The model name is ignored.
-

18. In the context of a `VoicePipelineConfig`, an instance of `OpenAIVoiceModelProvider` would be assigned to which attribute?

- ☐ a) `stt_settings`
 - ☐ b) `tts_settings`
 - ☐ c) `workflow_name`
 - ☒ d) **`model_provider`**
-

19. `OpenAIVoiceModelProvider` is concerned with:

- ☐ a) The actual content of the user's speech.
- ☐ b) The conversational flow and agent's logic.
- ☒ c) **Providing the underlying AI services for voice processing.**
- ☐ d) Storing historical data of voice interactions.

20. What would be a reason to provide a custom AsyncOpenAI client via the constructor, rather than letting the provider create one?

- ☐ a) To bypass API key validation.
- ☐ b) To use a synchronous OpenAI client.
- ☒ c) **To apply shared configurations, custom headers, or manage connection pools externally for the OpenAI client.**
- ☐ d) To force the use of local models.

OpenAI Agents SDK - OpenAI STT MCQs

1. From which base class does OpenAISTTTranscriptionSession inherit?

- ☐ a) STTModel
- ☒ b) **StreamedTranscriptionSession**
- ☐ c) AudioInput
- ☐ d) AsyncOpenAI

2. What is the primary function of OpenAISTTTranscriptionSession?

- ☐ a) To synthesize speech from text
- ☒ b) **To manage a continuous, real-time speech transcription session with OpenAI's STT service**
- ☐ c) To transcribe a single, complete audio file
- ☐ d) To configure OpenAI API keys

3. OpenAISTTModel inherits from which abstract base class?

- ☐ a) TTSTModel
- ☒ b) **STTModel**
- ☐ c) VoiceModelProvider
- ☐ d) OpenAISTTTranscriptionSession

4. Which parameter is required in the OpenAISTTModel constructor to specify the particular OpenAI STT model to use?

- ☐ a) api_key
- ☒ b) **model**
- ☐ c) base_url
- ☐ d) settings

5. What type of client does OpenAISTTModel require in its constructor to make API calls to OpenAI?

- ☐ a) SyncOpenAI
- ☐ b) requests.Session

- ✓c) AsyncOpenAI
 - ✗d) HTTPClient
-

6. The transcribe method of OpenAISTTModel is designed for:

- ✗a) Streaming audio input in real-time
 - ✗b) Yielding transcription turns continuously
 - ✓c) **Transcribing static, complete AudioInput objects**
 - ✗d) Converting text to speech
-

7. What is the return type of the OpenAISTTModel.transcribe() method?

- ✗a) AsyncIterator[str]
 - ✗b) bytes
 - ✓c) **str**
 - ✗d) AudioInput
-

8. Which method of OpenAISTTModel is used to initiate a real-time, streamed transcription process?

- ✗a) transcribe()
 - ✗b) run_session()
 - ✓c) **create_session()**
 - ✗d) start_transcription()
-

9. When OpenAISTTModel.create_session() is called, what does it return?

- ✗a) The full transcription string
 - ✗b) An AudioInput object
 - ✓c) **A StreamedTranscriptionSession (specifically an OpenAISTTTranscriptionSession)**
 - ✗d) An AsyncOpenAI client
-

10. The input parameter for OpenAISTTModel.transcribe() is of type AudioInput. What does this imply about the audio data?

- ✗a) It's expected to be pushed incrementally
 - ✗b) It must come from a live microphone
 - ✓c) **It represents a complete, static audio segment**
 - ✗d) It's always in a compressed format
-

11. Which of the following is a common OpenAI STT model name used with OpenAISTTModel?

- ✗a) "tts-1"
- ✗b) "gpt-3.5-turbo"
- ✓c) **"whisper-1"**

- ☒ d) "dall-e-3"
-

12. The `trace_include_sensitive_data` and `trace_include_sensitive_audio_data` parameters in both `transcribe` and `create_session` methods are related to:

- ☒ a) Data compression
 - ☒ b) Audio processing quality
 - ☒ c) **Observability and tracing configuration**
 - ☒ d) API rate limits
-

13. The `OpenAISTTTranscriptionSession` class internally manages the communication to OpenAI's STT service, which often involves:

- ☒ a) One-time HTTP POST requests
 - ☒ b) UDP broadcast messages
 - ☒ c) **WebSocket or streaming HTTP connections**
 - ☒ d) FTP transfers
-

14. What is the role of `STTModelSettings` when passed to `OpenAISTTModel` methods?

- ☒ a) To define the `model_name`
 - ☒ b) To provide the `api_key`
 - ☒ c) **To configure specific transcription parameters like prompt or language**
 - ☒ d) To manage the `AsyncOpenAI` client instance
-

15. If you wanted to continuously send audio from a live microphone to OpenAI for transcription, you would first call `OpenAISTTModel.create_session()` and then interact with the returned:

- ☒ a) `OpenAISTTModel` object
 - ☒ b) `AudioInput` object
 - ☒ c) **`OpenAISTTTranscriptionSession` object**
 - ☒ d) `STTModelSettings` object
-

16. What does the `OpenAISTTTranscriptionSession.transcribe_turns()` method yield?

- ☒ a) Audio bytes
 - ☒ b) Error messages
 - ☒ c) **Text transcriptions, typically one "turn" at a time**
 - ☒ d) Configuration settings
-

17. If an `OpenAISTTModel` is initialized with `model="whisper-1"`, but you pass a different model name to its `transcribe()` or `create_session()` methods, what would likely happen?

- ☒ a) The method would ignore the new model name

- ☒ b) The constructor model is used by default, but a new model name passed to the method may override it if allowed
 - ☐ c) An error would be raised
 - ☐ d) It would dynamically switch to the new model for that call without warning
-

18. The `OpenAISttModel` is considered a "concrete" class because:

- ☐ a) It does not inherit from any other class
 - ☐ b) It is defined in a Python file
 - ☒ c) It provides actual implementations for the abstract methods of its base class `STTModel`
 - ☐ d) It only deals with static data
-

19. `OpenAITTTranscriptionSession.close()` is essential for:

- ☐ a) Starting a new transcription
 - ☐ b) Retrieving the final transcript
 - ☒ c) Properly cleaning up resources associated with the streamed transcription connection
 - ☐ d) Changing the STT model settings
-

20. Which class is responsible for encapsulating the audio data itself, that is passed to `OpenAISttModel.transcribe()`?

- ☐ a) `STTModelSettings`
 - ☐ b) `StreamedAudioInput`
 - ☒ c) **`AudioInput`**
 - ☐ d) `TTSTModel`
-

OpenAI Agents SDK - OpenAI TTS MCQs

1. From which base class does `OpenAITTSModel` inherit?

- ☐ a) `STTModel`
 - ☒ b) **`TTSTModel`**
 - ☐ c) `VoiceModelProvider`
 - ☐ d) `AsyncOpenAI`
-

2. What is the primary function of `OpenAITTSModel`?

- ☐ a) To transcribe spoken audio into text
- ☐ b) To manage real-time STT transcription sessions
- ☒ c) **To convert text into speech using OpenAI models**
- ☐ d) To provide configuration settings for voice models

3. Which parameter in the `OpenAITTSModel` constructor specifies the particular OpenAI TTS model to use?

- ☐ a) `api_key`
 - ☒ b) **`model`**
 - ☐ c) `base_url`
 - ☐ d) `settings`
-

4. What type of client does `OpenAITTSModel` require in its constructor to interact with the OpenAI API?

- ☐ a) `SyncOpenAI`
 - ☐ b) `requests.Session`
 - ☒ c) **`AsyncOpenAI`**
 - ☐ d) `HTTPClient`
-

5. The `run` method of `OpenAITTSModel` takes `text` and `settings` as parameters. What is the type of `settings`?

- ☐ a) `STTModelSettings`
 - ☒ b) **`TTSModelSettings`**
 - ☐ c) `VoicePipelineConfig`
 - ☐ d) `dict`
-

6. What is the return type of the `OpenAITTSModel.run()` method?

- ☐ a) `str`
 - ☐ b) `bytes`
 - ☒ c) **`AsyncIterator[bytes]`**
 - ☐ d) `float`
-

7. The `AsyncIterator[bytes]` return type of `run()` is crucial for:

- ☐ a) Reducing file size
 - ☐ b) Ensuring perfect fidelity
 - ☒ c) **Enabling real-time streaming and low-latency playback**
 - ☐ d) Offline batch processing
-

8. Which attribute from `TTSModelSettings` would `OpenAITTSModel` use to determine the voice's gender and tone?

- ☐ a) `speed`
 - ☐ b) `buffer_size`
 - ☒ c) **`voice`**
 - ☐ d) `instructions`
-

9. If `TTSModelSettings.speed` is set, how does `OpenAITTSModel` leverage this?

- ☐ a) Changes pitch
 - ☒ b) **Adjusts speaking rate**
 - ☐ c) Increases volume
 - ☐ d) Adds pauses
-

10. The `instructions` parameter in `TTSModelSettings` passed to `OpenAITTSModel.run()` is used for:

- ☐ a) Grammar correction
 - ☐ b) Language specification
 - ☒ c) **Guiding the model's tone or partial sentence handling**
 - ☐ d) API retry logic
-

11. What kind of audio format are the byte chunks typically in when yielded by `OpenAITTSModel.run()`?

- ☐ a) MP3
 - ☐ b) WAV
 - ☒ c) **PCM**
 - ☐ d) AAC
-

12. `OpenAITTSModel` is considered a "concrete" class because:

- ☐ a) Uses a specific model name
 - ☐ b) Is defined in a .py file
 - ☒ c) **Implements the abstract methods of `TTSModel`**
 - ☐ d) Relies on `AsyncOpenAI` client
-

13. Which class typically provides the `AsyncOpenAI` client instance to `OpenAITTSModel`'s constructor?

- ☐ a) `TTSModelSettings`
 - ☐ b) `VoicePipelineConfig`
 - ☒ c) **`OpenAIVoiceModelProvider`**
 - ☐ d) `StreamedAudioResult`
-

14. The `text_splitter` setting passed via `TTSModelSettings` to `OpenAITTSModel.run()` helps in:

- ☐ a) Volume control
 - ☐ b) Detecting silence
 - ☒ c) **Breaking input into chunks for natural streaming**
 - ☐ d) Translation
-

15. What happens if the model parameter passed to OpenAITTSModel's constructor is invalid?

- ☐ a) Defaults to "tts-1"
 - ☐ b) Uses STT instead
 - ☒ c) **Likely results in API call error**
 - ☐ d) Picks closest model
-

16. OpenAITTSModel directly interacts with:

- ☐ a) Microphone
 - ☐ b) VoiceWorkflowBase
 - ☒ c) **OpenAI's TTS API endpoint**
 - ☐ d) Local file storage
-

17. The transform_data attribute in TTSMModelSettings allows for:

- ☐ a) Change model at runtime
 - ☐ b) Convert PCM to MP3
 - ☒ c) **Apply function to raw audio before streaming**
 - ☐ d) Split into tracks
-

18. Primary benefit of AsyncIterator in run() for conversational AI:

- ☐ a) Offline use
 - ☐ b) Lower memory on OpenAI side
 - ☒ c) **Early playback while generating response**
 - ☐ d) Detailed logging
-

19. Which TTSVoice would you typically pass to TTSMModelSettings for OpenAITTSModel?

- ☐ a) "male" / "female"
 - ☐ b) "standard" / "neural"
 - ☒ c) **"alloy", "echo", "nova", etc.**
 - ☐ d) "english", "spanish"
-

20. OpenAITTSModel represents output generation in the voice pipeline, converting:

- ☐ a) Speech to text
 - ☒ b) **Text to speech**
 - ☐ c) Text to text
 - ☐ d) Speech to speech
-

Extensions

OpenAI Agents SDK - Handoff Filters MCQs

1. What is the primary function of `remove_all_tools`?

- ☐ a) To add new tools to a `HandoffInputData` object.
- ☐ b) To execute all tools defined in the pipeline.
- ☒ c) **To filter out all tool-related items from `HandoffInputData`.**
- ☐ d) To log all tool usage to a file.

2. `remove_all_tools` operates on which specific data type?

- ☐ a) `str`
- ☐ b) `dict`
- ☒ c) **`HandoffInputData`**
- ☐ d) `ToolCall`

3. Which of the following is explicitly filtered out by `remove_all_tools`?

- ☐ a) User messages
- ☐ b) Agent responses
- ☒ c) **Web search outputs**
- ☐ d) Conversation summaries

4. The function signature `-> HandoffInputData` indicates that `remove_all_tools` returns:

- ☐ a) A boolean value indicating success or failure.
- ☐ b) The original, unmodified `HandoffInputData`.
- ☒ c) **A modified `HandoffInputData` object with tool items removed.**
- ☐ d) A different data type altogether.

5. Why might you want to remove tool information during a handoff to a human agent?

- ☐ a) To increase the speed of the handoff.
- ☐ b) To reduce the memory footprint of the agent.
- ☒ c) **To present a cleaner, more relevant, and less technical view of the conversation.**
- ☐ d) To allow the human agent to re-run the tools.

6. Which of these tool-related items does `remove_all_tools` target?

- ☐ a) Only file search results.
 - ☐ b) Only function call inputs.
 - ☐ c) Only web search queries.
 - ☒ d) **File search, web search, and function calls (including output).**
-

7. A key reason for using handoff filters like `remove_all_tools` is to enhance:

- ☐ a) Agent processing speed.
 - ☐ b) Model accuracy.
 - ☒ c) **Security and privacy during data transfer.**
 - ☐ d) The number of tools available to the agent.
-

8. If a `HandoffInputData` contains a record of an agent making an `add_to_cart` function call and its successful output, what would `remove_all_tools` do?

- ☐ a) Leave the `add_to_cart` record untouched.
 - ☒ b) **Remove both the `add_to_cart` function call and its output.**
 - ☐ c) Only remove the `add_to_cart` function call, keeping the output.
 - ☐ d) Only remove the `add_to_cart` output, keeping the function call.
-

9. What is the main benefit of simplification that `remove_all_tools` offers for handoff data?

- ☐ a) It makes the data smaller in size.
 - ☐ b) It encrypts sensitive data.
 - ☒ c) **It makes the data more digestible and focused for the receiving entity.**
 - ☐ d) It adds metadata to the handoff.
-

10. In what scenario would `remove_all_tools` be most beneficial?

- ☐ a) When debugging an agent's tool usage internally.
 - ☒ b) **When transferring a conversation context from an AI agent to a human customer support agent.**
 - ☐ c) When training a new tool usage model.
 - ☐ d) When conducting performance benchmarks of tool execution.
-

11. This function is categorized under "Handoff filters" because it:

- ☐ a) Filters out handoff requests.
 - ☐ b) Filters which agent receives the handoff.
 - ☒ c) **Modifies the data before it is handed off.**
 - ☐ d) Filters data after it has been received.
-

12. Could `remove_all_tools` prevent a human agent from seeing a user's sensitive query that was part of a web search tool call?

- ☒ a) Yes, if the web search tool call and its input are part of the filtered items.
 - ☐ b) No, it only filters output, not input.
 - ☐ c) Only if the query was explicitly marked as sensitive.
 - ☐ d) It's unrelated to user queries.
-

13. If `HandoffInputData` represents the complete conversation history including tool calls, `remove_all_tools` aims to present a view of the conversation that is primarily focused on:

- ☐ a) Technical diagnostics.
 - ☐ b) Detailed tool execution steps.
 - ☒ c) The natural language dialogue and core intent.
 - ☐ d) External API response structures.
-

14. `remove_all_tools` helps prevent misinterpretation by the receiving entity because:

- ☐ a) It translates tool outputs into natural language.
 - ☒ b) Raw tool outputs can be confusing if the recipient doesn't have the context or capability to interpret them.
 - ☐ c) It ensures all tool outputs are fully explained.
 - ☐ d) It adds warnings to uninterpretable data.
-

15. Does `remove_all_tools` modify the agent's internal state or only the data being prepared for handoff?

- ☐ a) It modifies the agent's internal state.
 - ☒ b) It primarily modifies the data being prepared for handoff.
 - ☐ c) It modifies both.
 - ☐ d) It makes a read-only copy.
-

16. If an agent used a tool named `check_stock`, what would `remove_all_tools` filter out?

- ☐ a) Only the fact that `check_stock` was attempted.
 - ☐ b) Only the numerical output of `check_stock`.
 - ☒ c) Both the `check_stock` function call and its output.
 - ☐ d) Nothing, as `check_stock` is not explicitly mentioned in the documentation snippet.
-

17. The `src/agents/extensions/handoff_filters.py` path suggests that this function is part of:

- ☐ a) The core agent runtime.
 - ☐ b) The main VoicePipeline.
 - ☒ c) Optional extensions or utilities for agent handoff scenarios.
 - ☐ d) Low-level model definitions.
-

18. `remove_all_tools` contributes to providing a more concise representation of data during handoff by:

- ☐ a) Summarizing tool outputs into a single sentence.
 - ☐ b) Compressing the `HandoffInputData` object.
 - ☒ c) **Removing detailed, potentially irrelevant tool execution logs.**
 - ☐ d) Encrypting the entire handoff data.
-

19. Why would `HandoffInputData` likely include tool items in the first place before filtering?

- ☐ a) They are necessary for basic conversation flow.
 - ☒ b) **The `HandoffInputData` usually captures the full, rich internal state and history of the agent's processing.**
 - ☐ c) Tools are always public information.
 - ☐ d) It's a design error.
-

20. Which principle does `remove_all_tools` align with in data transfer scenarios?

- ☐ a) Principle of Maximum Information.
 - ☒ b) **Principle of Least Privilege/Need-to-know.**
 - ☐ c) Principle of Data Redundancy.
 - ☐ d) Principle of Universal Data Access.
-

OpenAI Agents SDK - Handoff Prompt MCQs

1. What is the main purpose of `RECOMMENDED_PROMPT_PREFIX`?

- ☐ a) To define the agent's personality.
 - ☒ b) **To provide system-level instructions about the multi-agent system and handoffs.**
 - ☐ c) To store user conversation history.
 - ☐ d) To list available tools for the agent.
-

2. `RECOMMENDED_PROMPT_PREFIX` is a:

- ☐ a) Function.
 - ☐ b) Class.
 - ☒ c) **String constant (module attribute).**
 - ☐ d) Data structure.
-

3. The instruction *"Transfers between agents are handled seamlessly in the background; do not mention or draw attention to these transfers in your conversation with the user"* primarily serves to improve:

- ☐ a) Agent processing speed.
- ☐ b) Model accuracy.

- ☒ c) Developer debugging experience.
 - ☒ d) **User experience and conversational flow.**
-

4. What type of function call does RECOMMENDED_PROMPT_PREFIX instruct agents to use for handoffs?

- ☒ a) call_external_api()
 - ☒ b) initiate_transfer()
 - ☒ c) handoff_conversation()
 - ☒ d) **transfer_to_<agent_name>**
-

5. What does the prompt_with_handoff_instructions function do?

- ☒ a) It executes a handoff.
 - ☒ b) **It prepends RECOMMENDED_PROMPT_PREFIX to a given agent prompt.**
 - ☒ c) It removes handoff instructions from a prompt.
 - ☒ d) It validates the correctness of a handoff prompt.
-

6. If original_prompt = "You are a sales agent.", what would prompt_with_handoff_instructions(original_prompt) return?

- ☒ a) original_prompt unchanged.
 - ☒ b) Only RECOMMENDED_PROMPT_PREFIX.
 - ☒ c) **RECOMMENDED_PROMPT_PREFIX followed by original_prompt.**
 - ☒ d) original_prompt followed by RECOMMENDED_PROMPT_PREFIX.
-

7. The phrase "# System context" at the beginning of RECOMMENDED_PROMPT_PREFIX is a hint for the LLM that these instructions are:

- ☒ a) From the user.
 - ☒ b) Optional suggestions.
 - ☒ c) **Core directives about its operational environment.**
 - ☒ d) Debugging information.
-

8. Why is it important for the agent to know that it's part of a "multi-agent system"?

- ☒ a) To allow it to communicate directly with other agents.
 - ☒ b) To enable it to train other agents.
 - ☒ c) **To provide context for its ability to hand off tasks to other specialized agents.**
 - ☒ d) To help it choose the right language for the user.
-

9. If an agent is not given the RECOMMENDED_PROMPT_PREFIX, and is expected to perform handoffs, what is a likely outcome?

- ☒ a) It will automatically figure out how to hand off.

- ☐ b) It will refuse to communicate with the user.
 - ☒ c) **It might not know how to call the `transfer_to_<agent_name>` function or might explicitly mention transfers to the user.**
 - ☐ d) It will only respond in short sentences.
-

10. The `prompt_with_handoff_instructions` function returns a `str`. This means the output is a:

- ☐ a) List of instructions.
 - ☐ b) Boolean value.
 - ☒ c) **Single, combined string.**
 - ☐ d) Dictionary of prompt components.
-

11. This handoff prompt mechanism helps in guiding the agent's:

- ☐ a) Memory management.
 - ☐ b) Audio processing.
 - ☒ c) **Tool selection and conversational etiquette during transfers.**
 - ☐ d) External API integration.
-

12. The `RECOMMENDED_PROMPT_PREFIX` specifies that handoffs are "seamlessly in the background" from whose perspective?

- ☐ a) The developer's.
 - ☐ b) The agent's.
 - ☒ c) **The user's.**
 - ☐ d) The system administrator's.
-

13. What is the type of `prompt` parameter in `prompt_with_handoff_instructions`?

- ☐ a) `list[str]`
 - ☒ b) **`str`**
 - ☐ c) `HandoffInputData`
 - ☐ d) `AgentConfig`
-

14. The `RECOMMENDED_PROMPT_PREFIX` defines two primary abstractions of the Agents SDK. What are they?

- ☐ a) Users and Tools.
 - ☐ b) Input and Output.
 - ☒ c) **Agents and Handoffs.**
 - ☐ d) Workflows and Models.
-

15. If `prompt_with_handoff_instructions` were not used, and you manually tried to add the prefix, what might be a potential issue?

- ☐ a) The prompt would become too short.
 - ☐ b) The `RECOMMENDED_PROMPT_PREFIX` might change unexpectedly.
 - ☒ c) **You might accidentally modify the original `RECOMMENDED_PROMPT_PREFIX` or miss future updates.**
 - ☐ d) The agent would ignore the instructions.
-

16. The statement "An agent encompasses instructions and tools" from the prefix means:

- ☐ a) Agents can only understand instructions.
 - ☐ b) Agents are purely rule-based.
 - ☒ c) **An agent has a set of instructions guiding its behavior and access to external functionalities.**
 - ☐ d) Agents are limited to internal operations only.
-

17. The `src/agents/extensions/handoff_prompt.py` path indicates that these are:

- ☐ a) Mandatory core components for any agent.
 - ☐ b) Tools for external system integration.
 - ☒ c) **Optional extensions to help manage agent handoffs.**
 - ☐ d) Low-level hardware drivers.
-

18. What is the benefit of making `transfer_to_<agent_name>` a "general" naming convention?

- ☐ a) It prevents name collisions.
 - ☐ b) It allows for dynamic agent creation.
 - ☒ c) **It provides a standardized and predictable way for the LLM to learn and invoke handoff tools.**
 - ☐ d) It makes the system more secure.
-

19. `RECOMMENDED_PROMPT_PREFIX` could be considered a form of:

- ☐ a) Output formatting.
 - ☐ b) User input validation.
 - ☒ c) **System message or meta-prompting.**
 - ☐ d) Error handling.
-

20. When an agent calls a `transfer_to_<agent_name>` function, who handles the actual "transfer" in the background according to the prompt?

- ☐ a) The user.
- ☐ b) The receiving agent.
- ☒ c) **The Agents SDK itself (implied by "Transfers between agents are handled seamlessly").**
- ☐ d) An external API.

OpenAI Agents SDK - LitellmModel MCQs

1. What is the primary purpose of the LitellmModel class?

- ☐ a) To integrate only OpenAI models into the Agents SDK.
- ☐ b) To manage local LLM deployments.
- ☒ c) **To enable the Agents SDK to use any LLM supported by the LiteLLM library.**
- ☐ d) To provide a new type of conversational agent.

2. LitellmModel inherits from which base class?

- ☐ a) Agent
- ☐ b) VoiceModelProvider
- ☒ c) **Model**
- ☐ d) AsyncOpenAI

3. What problem does LiteLLM aim to solve for developers working with multiple LLM providers?

- ☐ a) Reducing the size of LLM models.
- ☐ b) Improving the training speed of LLMs.
- ☒ c) **Providing a unified interface to various LLM APIs, simplifying integration.**
- ☐ d) Automating the creation of new LLM models.

4. Which of the following LLM providers can LiteLLM (and thus LitellmModel) potentially access?

- ☐ a) Only OpenAI and Anthropic.
- ☐ b) Only Google Gemini.
- ☐ c) Only Mistral.
- ☒ d) **OpenAI, Anthropic, Gemini, Mistral, and many others.**

5. What is a significant benefit of using LitellmModel in terms of model selection?

- ☐ a) It forces the use of a single, fixed LLM.
- ☒ b) **It allows easy switching between different LLM providers without changing core agent code.**
- ☐ c) It automatically selects the highest-cost model.
- ☐ d) It only supports open-source models.

6. If you want to optimize costs by using the cheapest available LLM, how might LitellmModel (via LiteLLM) help?

- ☐ a) By compressing prompts.
- ☐ b) By pre-caching all possible responses.

- ☒ c) By supporting routing requests to the most cost-effective provider.
 - ☐ d) By reducing the number of API calls made.
-

7. **LitellmModel acts as a(n) _____ within the Agents SDK to connect to external LLMs.**

- ☐ a) Executor
 - ☐ b) Logger
 - ☒ c) **Adapter/Bridge**
 - ☐ d) Validator
-

8. **What does "Model Agnosticism" mean in the context of LitellmModel?**

- ☐ a) The model is unaware of the prompt.
 - ☐ b) The model doesn't require an API key.
 - ☒ c) **The agent code does not need to be tightly coupled to a specific LLM provider's API.**
 - ☐ d) The model cannot be fine-tuned.
-

9. **If an agent in the SDK is configured to use a LitellmModel, and the primary LLM provider fails, what capability might LiteLLM offer?**

- ☐ a) It will immediately raise an error.
 - ☐ b) It will halt the agent's execution.
 - ☒ c) **It can be configured for automatic fallbacks to alternative models/providers.**
 - ☐ d) It will attempt to fix the failing provider.
-

10. **By integrating LiteLLM, LitellmModel contributes to:**

- ☐ a) Increased complexity in LLM integration.
 - ☐ b) Limited access to new LLMs.
 - ☒ c) **Simplified and unified API interaction for various LLMs.**
 - ☐ d) Slower response times from LLMs.
-

11. **Where is the source code for LitellmModel located?**

- ☒ a) **src/agents/core/models/litellm_model.py**
 - ☐ b) src/agents/main/litellm_model.py
 - ☐ c) src/agents/extensions/models/litellm_model.py
 - ☐ d) src/litellm/model.py
-

12. **The LitellmModel class allows developers to achieve which of the following without modifying core agent logic?**

- ☐ a) Changing the agent's personality.
- ☒ b) **Swapping the underlying LLM provider (e.g., from OpenAI to Anthropic).**
- ☐ c) Developing new tools.

- ☒ d) Deploying the agent to a new server.
-

13. Which of these is not a direct benefit of using `LitellmModel` (and `LiteLLM`)?

- ☒ a) Simplified integration with multiple LLM APIs.
- ☒ b) Potential for cost optimization.
- ☒ c) Training new, custom LLM models from scratch.
- ☒ d) Improved reliability through fallbacks.

✓ Trick question—**option C is NOT a benefit**, while the others **are**.

14. When an Agent uses a `LitellmModel`, the agent communicates with:

- ☒ a) Directly with the OpenAI API.
 - ☒ b) Directly with the Anthropic API.
 - ☒ c) The `LitellmModel` instance, which then uses `LiteLLM` to route the request.
 - ☒ d) A local `.json` file containing LLM responses.
-

15. The `LitellmModel` likely implements abstract methods from its `Model` base class such as:

- ☒ a) `get_api_key()`
 - ☒ b) `configure_database()`
 - ☒ c) `run_llm()` or `generate_response()` (inferred)
 - ☒ d) `install_dependencies()`
-

16. If a new LLM provider emerges, what is the most likely way it would be supported by the Agents SDK via `LitellmModel`?

- ☒ a) The `LitellmModel` class would need to be completely rewritten.
 - ☒ b) A new specific Agent subclass for that provider would be created.
 - ☒ c) **LiteLLM would add support for it, and `LitellmModel` would automatically gain access.**
 - ☒ d) The Agents SDK would need a new `VoiceModelProvider`.
-

17. `LitellmModel` is categorized under `src/agents/extensions` because it:

- ☒ a) Is a core, mandatory part of every agent.
 - ☒ b) Is an experimental feature that is not yet stable.
 - ☒ c) **Provides optional, extended functionality for integrating diverse LLMs.**
 - ☒ d) Is only for internal developer use.
-

18. What kind of API calls does `LiteLLM` typically unify?

- ☒ a) File system calls.
- ☒ b) Database queries.

- ☒ c) **Large Language Model (LLM) API calls.**
 - ☒ d) Operating system calls.
-

19. A developer looking to leverage a wide array of commercial and open-source LLMs in their agent with minimal code changes would find which class most beneficial?

- ☒ a) OpenAITTSModel
 - ☒ b) OpenAISTTModel
 - ☒ c) **LitellmModel**
 - ☒ d) VoiceModelProvider
-

20. The integration of LitellmModel into the Agents SDK demonstrates a design principle of:

- ☒ a) Tightly coupling to specific services.
 - ☒ b) **Modularity and extensibility.**
 - ☒ c) Limiting external dependencies.
 - ☒ d) Prioritizing performance over flexibility.
-